

Building Generative Adversarial Networks in PyTorch



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Understand Generative Adversarial Networks (GANs)

Candidate generation and evaluation

Role of generator

Role of discriminator

Generate handwritten digits using GANs trained on the MNIST dataset

Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs)

A new type of ML system (invented in 2014) that rely on two neural networks contesting in a game. Can be used to generate realistic images and videos of virtually anything.

Generative Adversarial Networks (GANs)

A new type of ML system (invented in 2014) that rely **on two neural networks contesting in a game**. Can be used to generate realistic images and videos of virtually anything.

Generative Adversarial Networks (GANs)

A new type of ML system (invented in 2014) that rely on two neural networks contesting in a game. Can be used to **generate realistic images and videos** of virtually anything.

Applications of GANs

AI-generated art

**Human image
synthesis**


**3D models from 2D
images**

**Improve
astronomical image**

**Traditional
classification**

**Traditional
regression**

GANs



Considered part of unsupervised learning

Developed in 2014 by Ian Goodfellow et al

Two distinct contesting neural networks

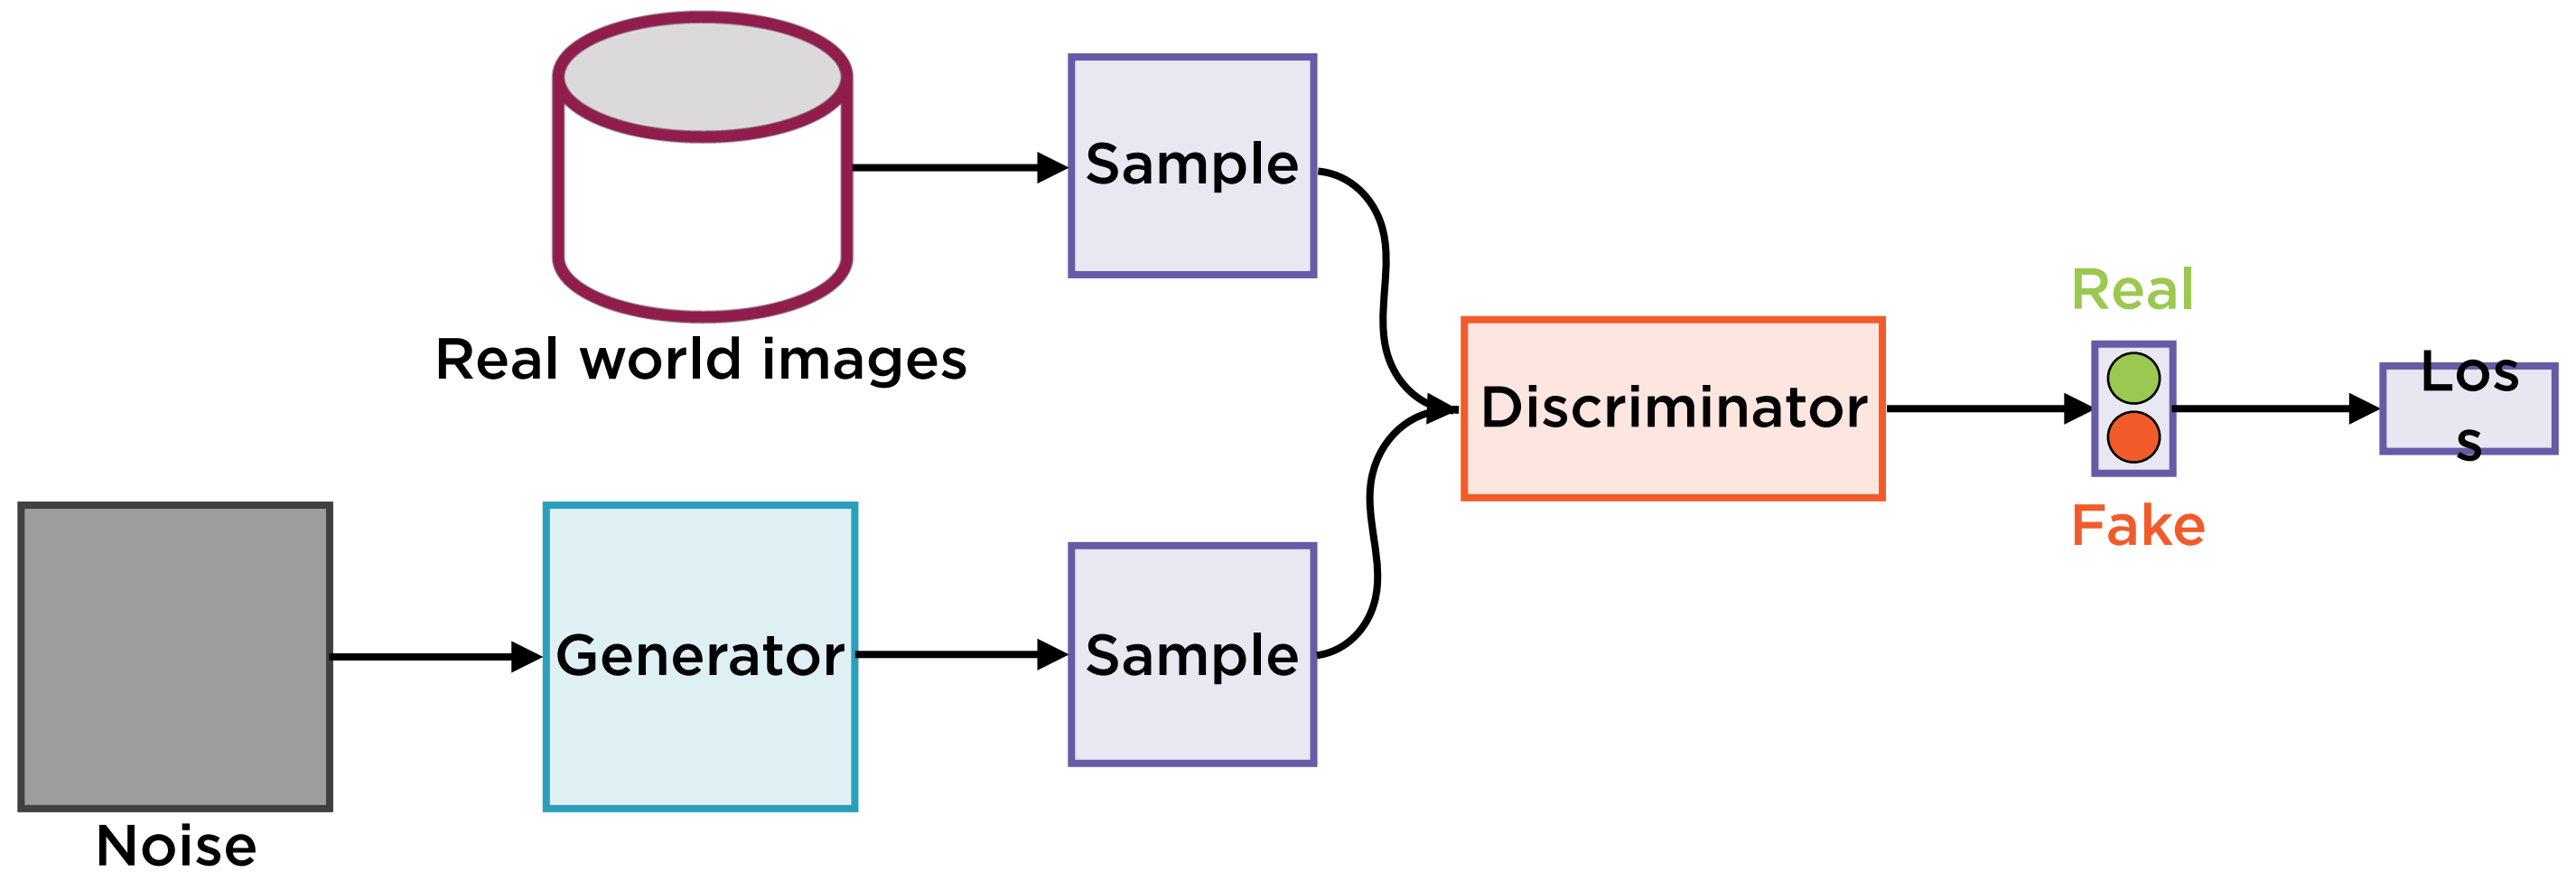
- Generative network
- Discriminative network

Two Neural Networks

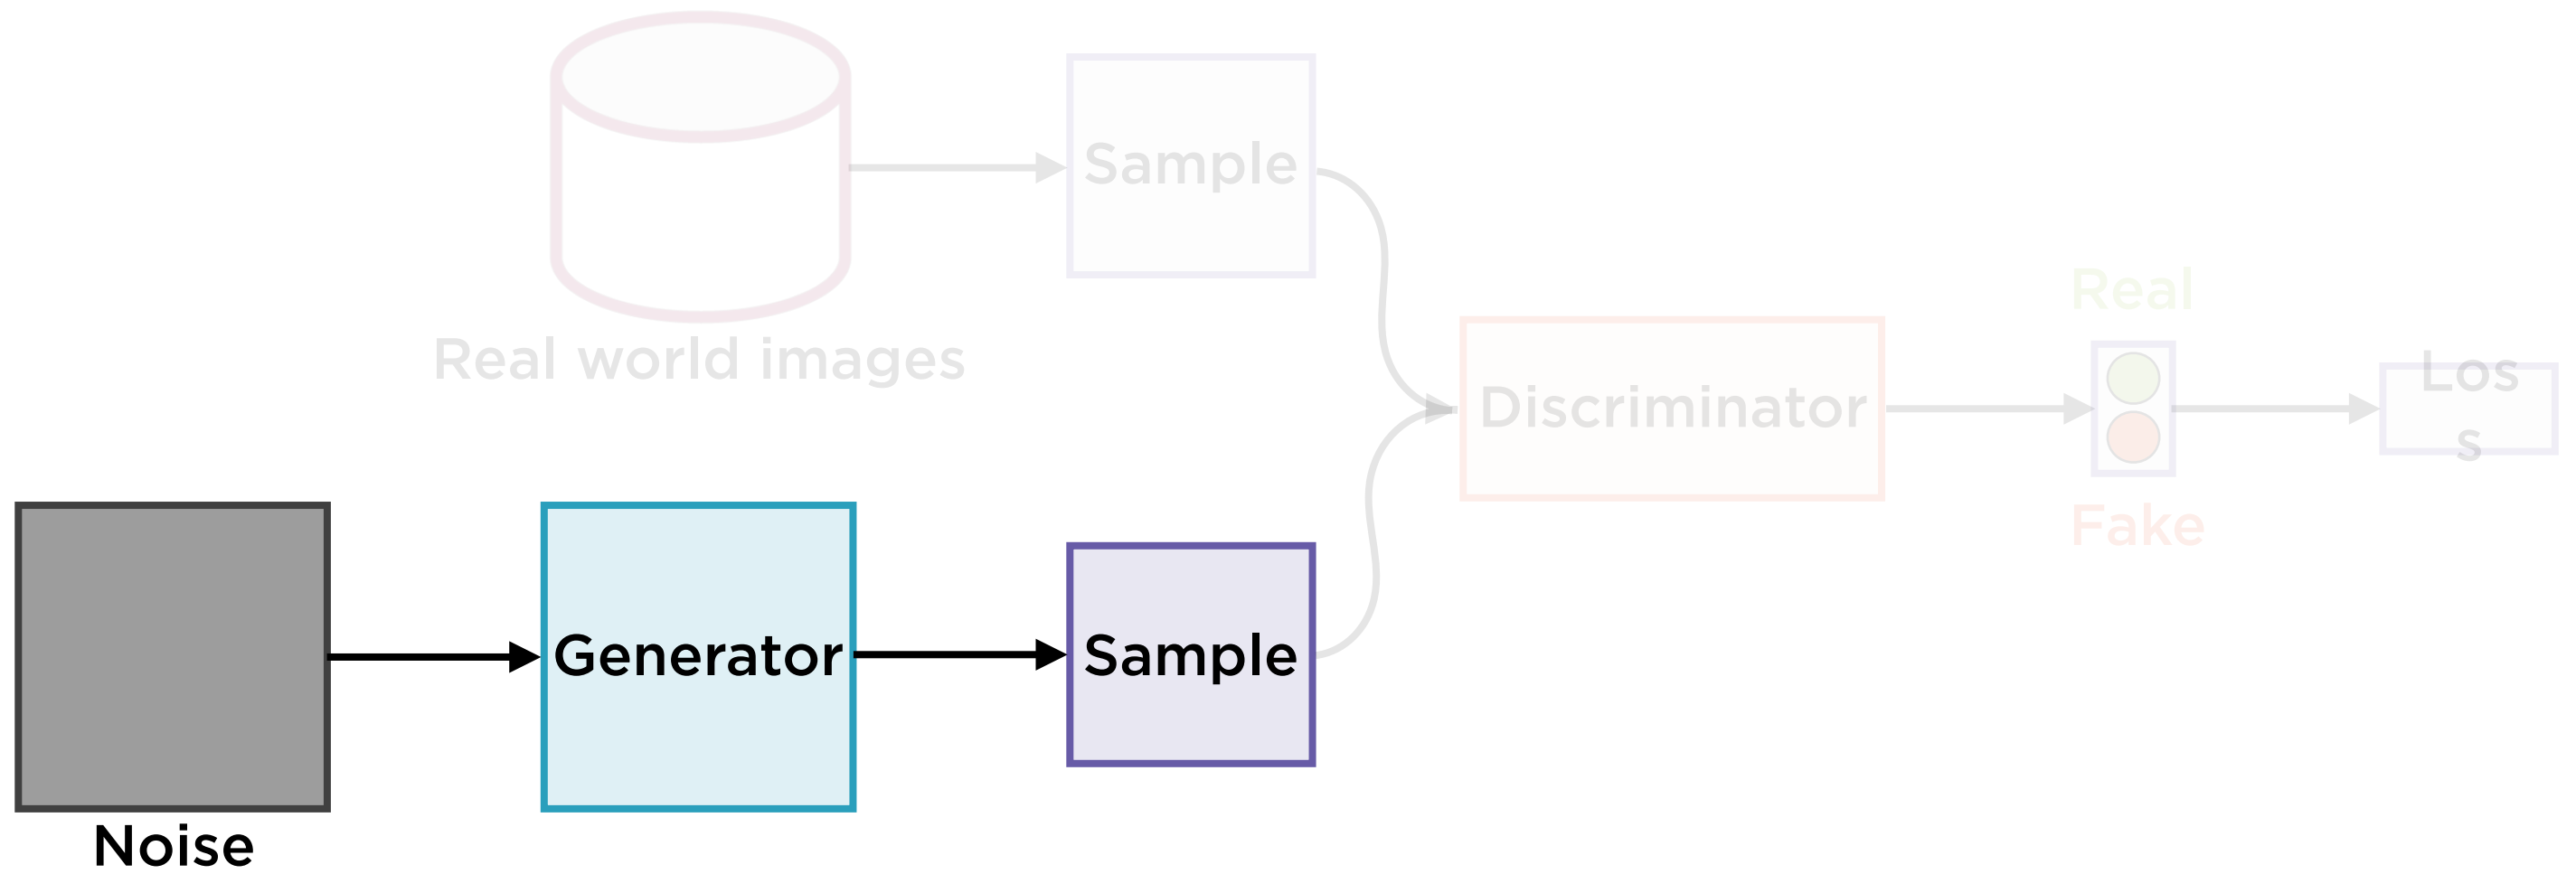
Generative Network
generates candidates

Discriminative Network
evaluates candidates

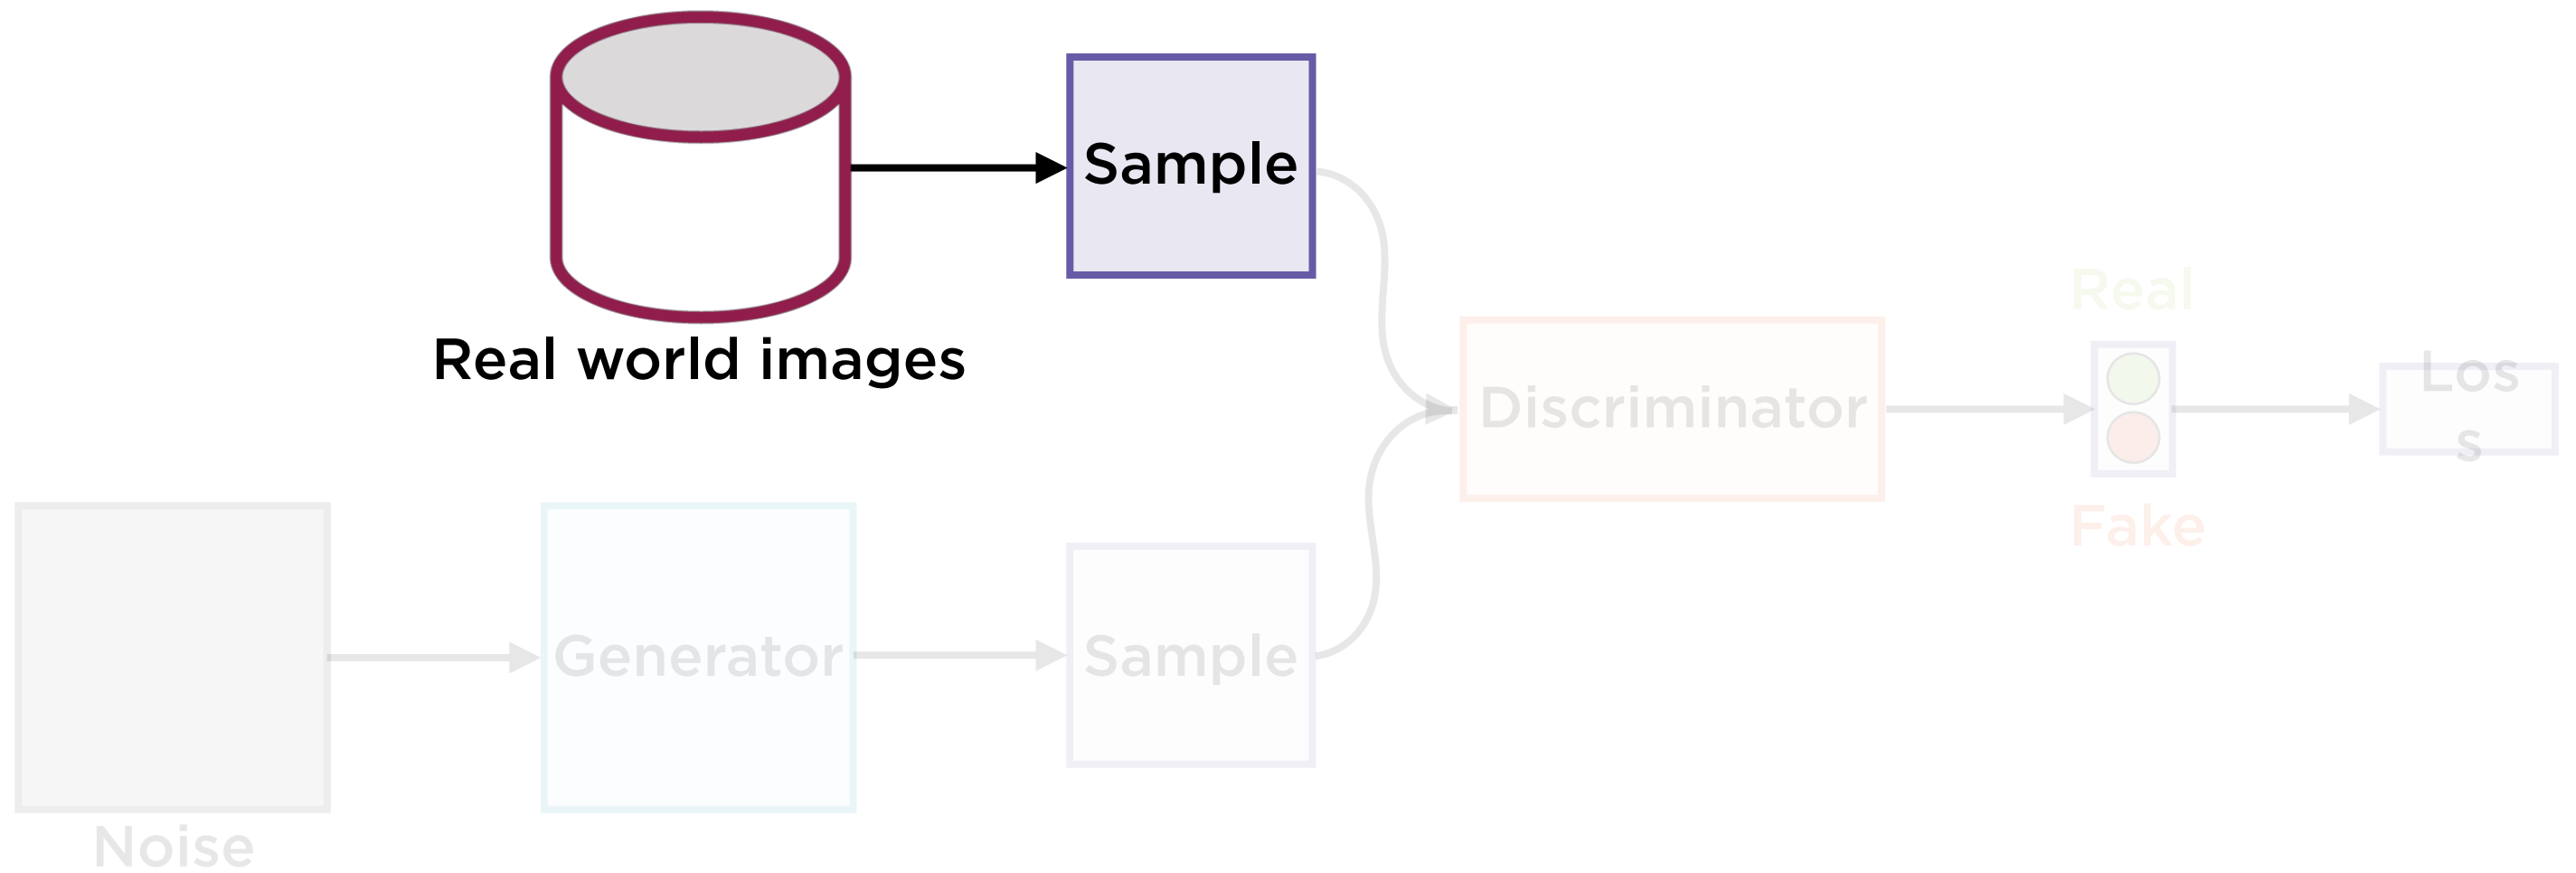
GANs



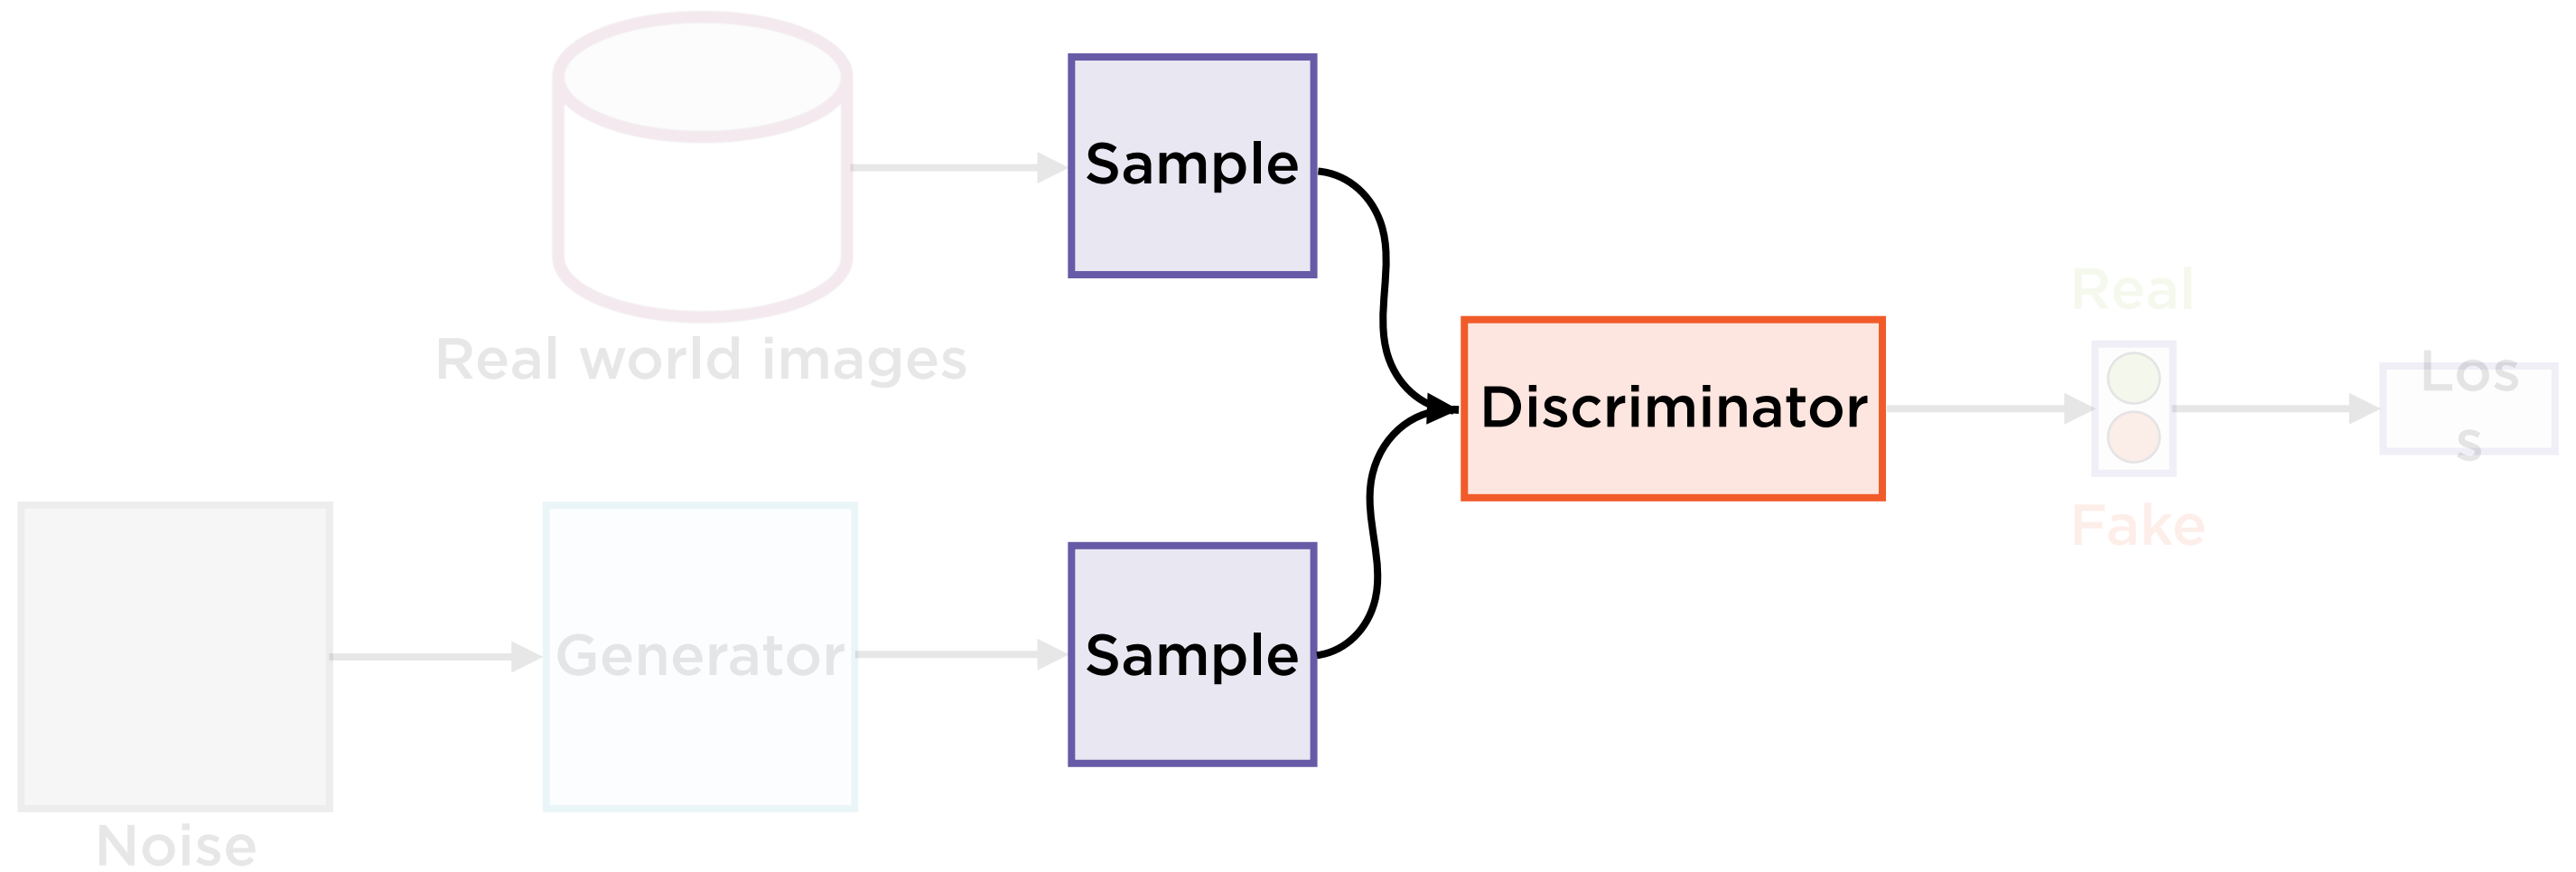
GANs



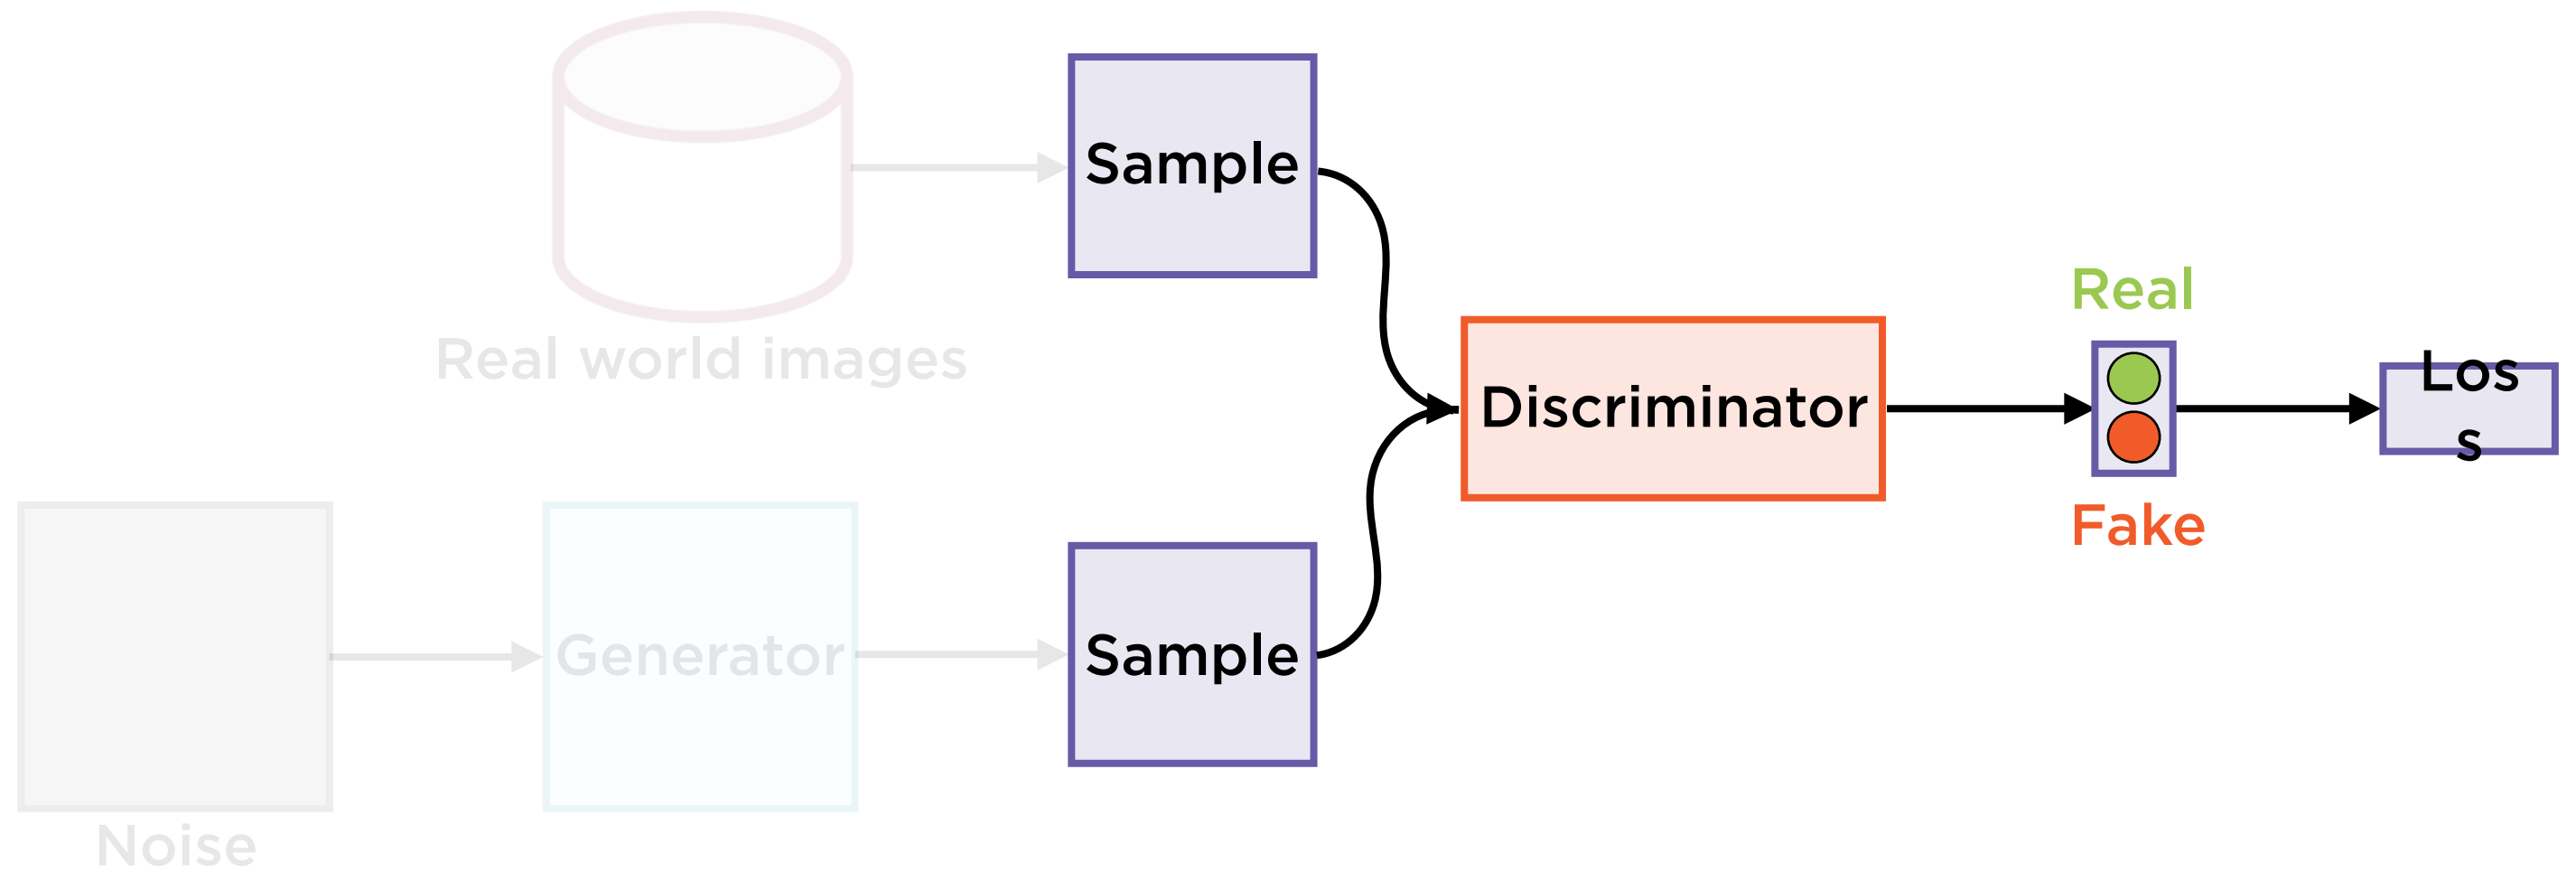
GANs



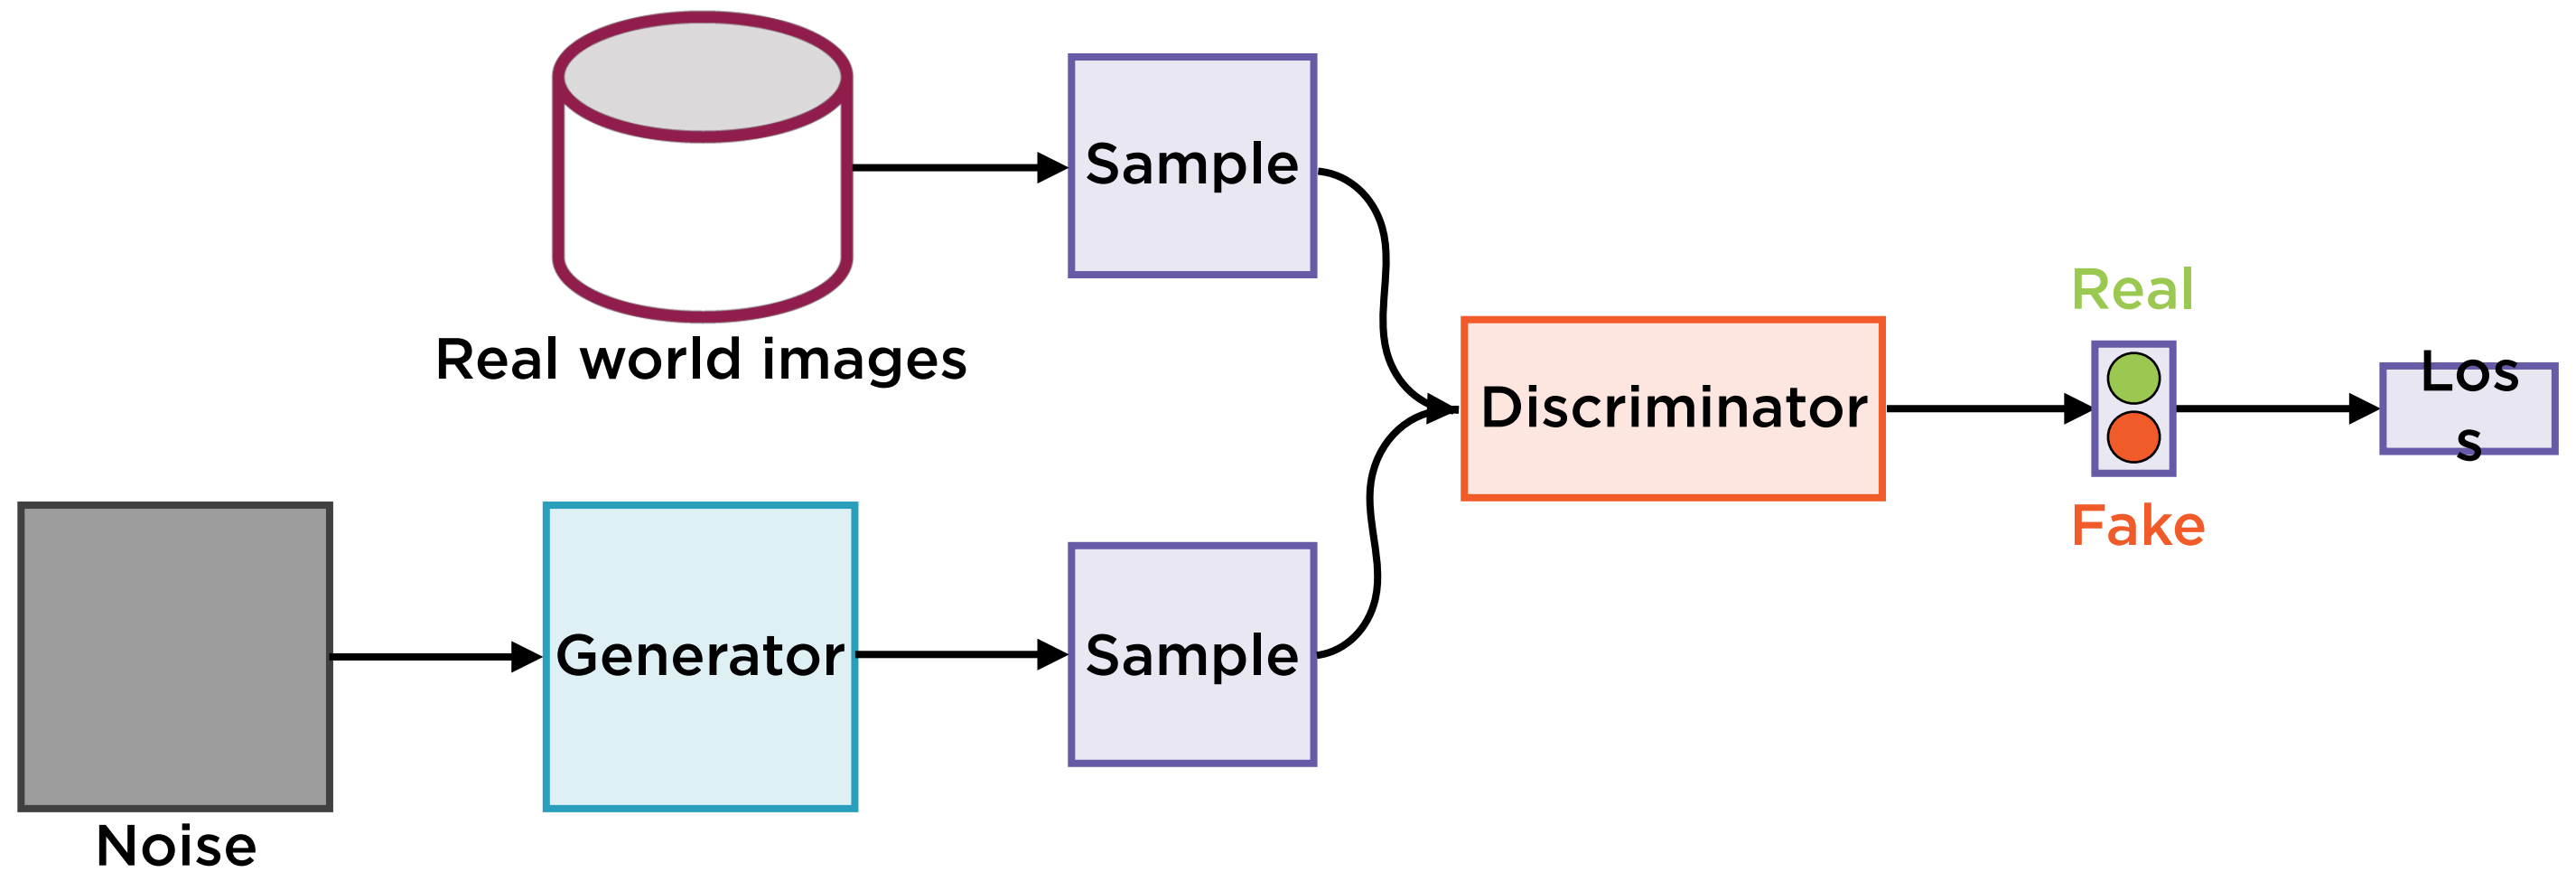
GANs




GANs



GANs



Generator




Generates data as realistically as possible

Trained to generate data similar to corpus

Seeks to fool discriminator

Discriminator




Generates probability that data is genuine

Classifies output of generator

Just like traditional classifier

Noise in GANs



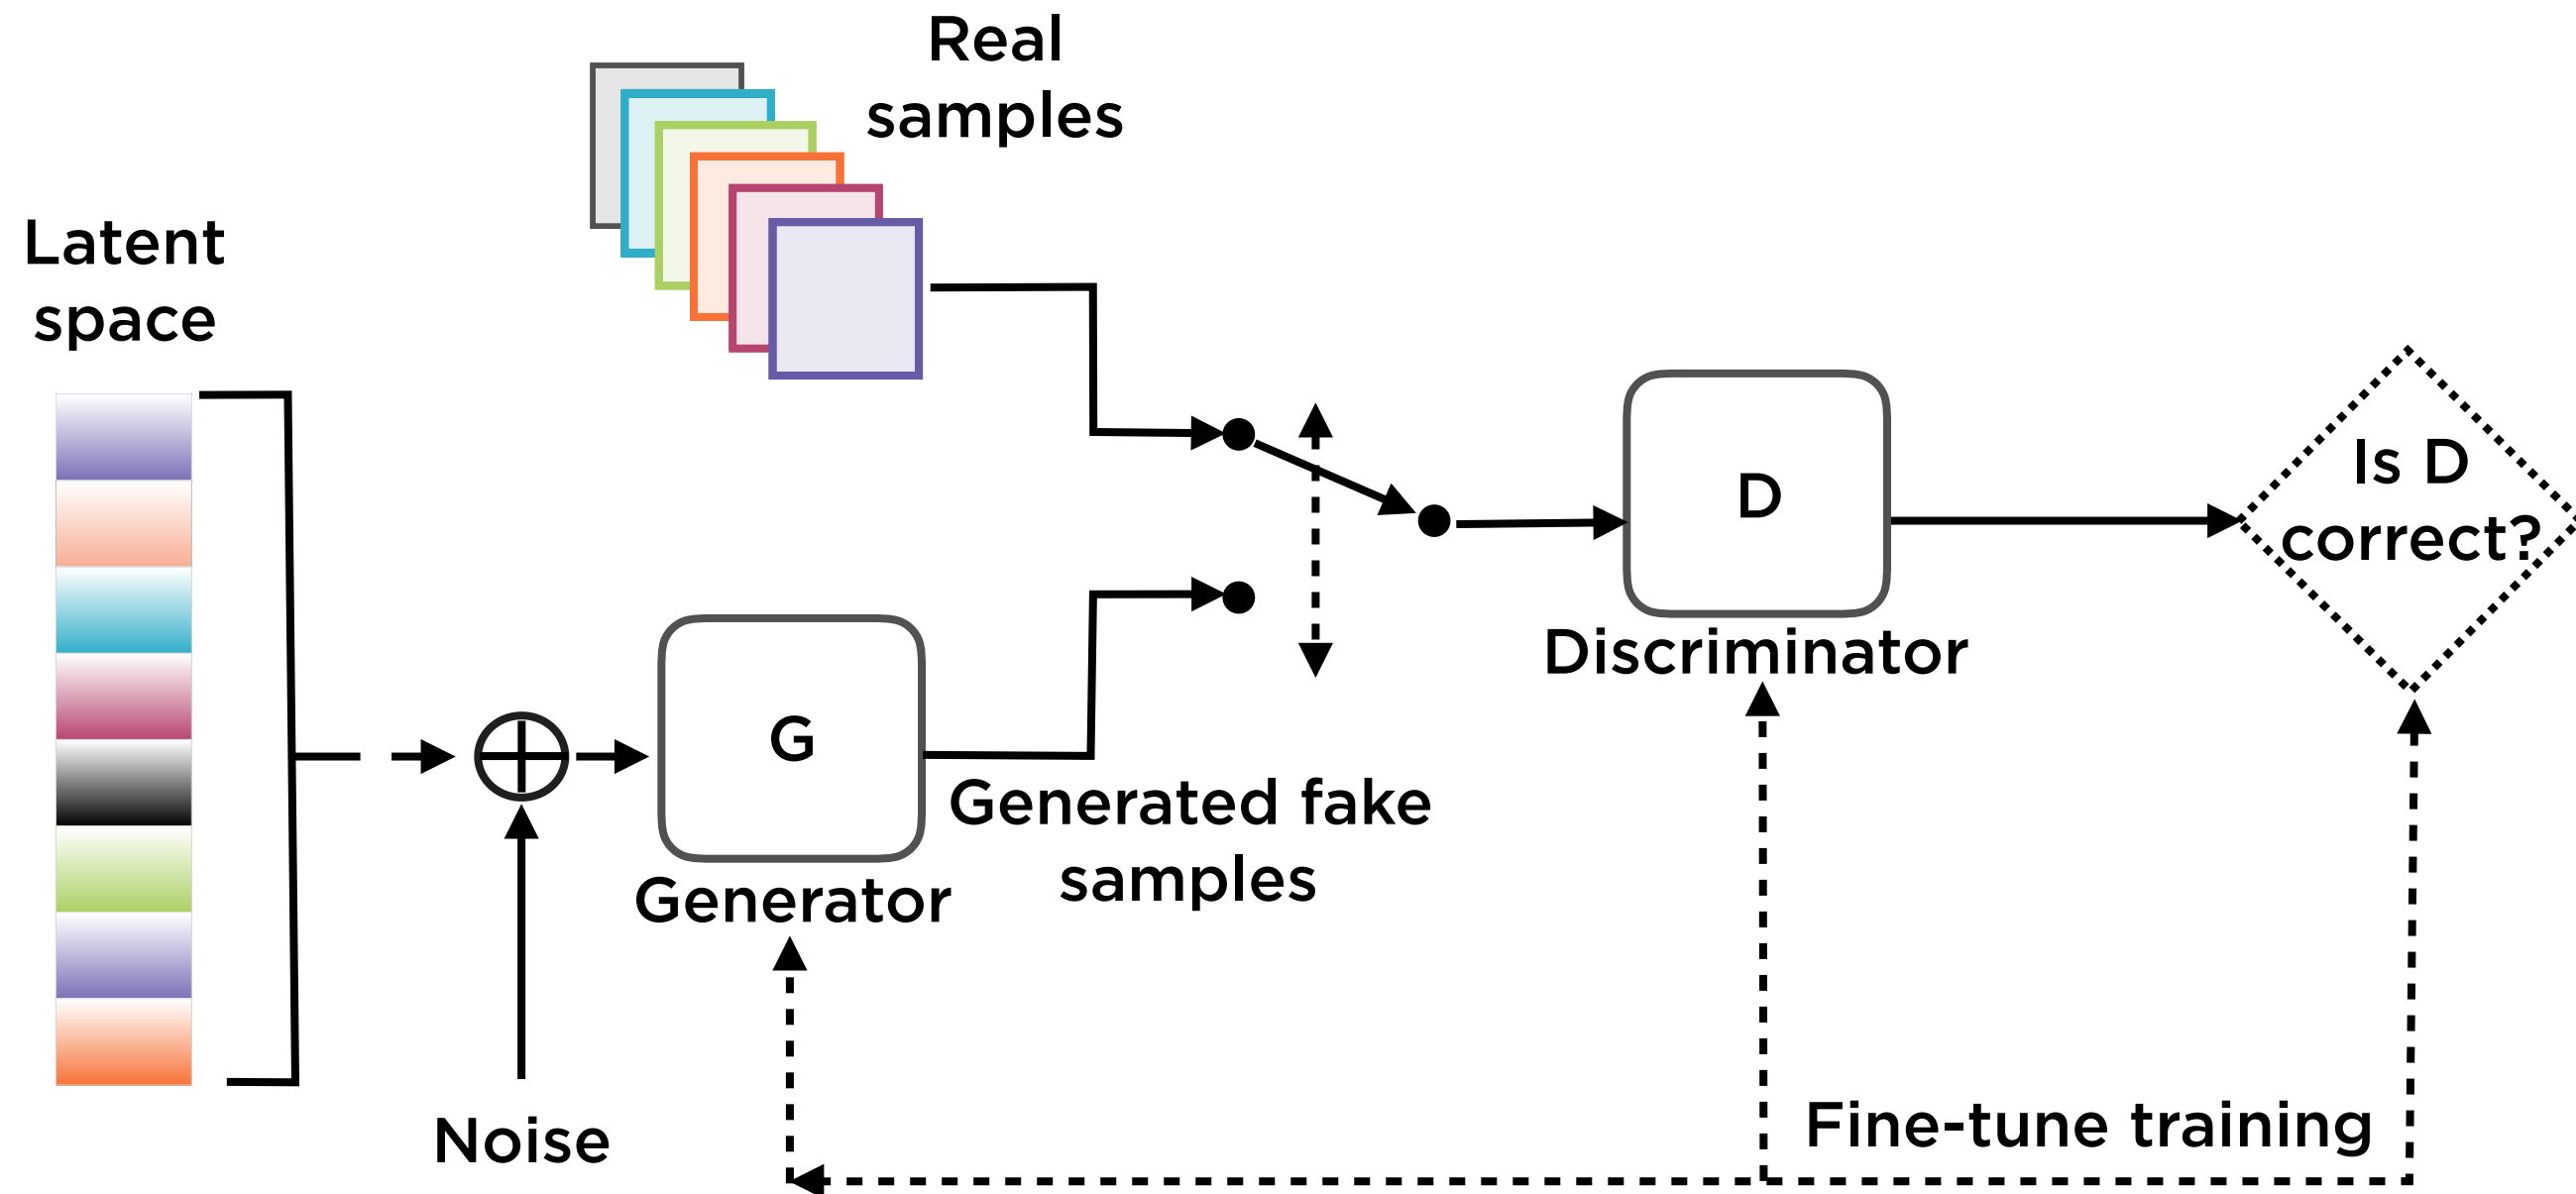
Requires function that generates noise

Create corpus of

- Real data points
- Noise function

Training a GAN

GANs



Training a GAN



Start with corpus of real points as well as noise

Train discriminator to tell them apart

Generate new noise points

Train generator to produce data that fools the discriminator

Repeat using optimizer


Discriminator



Maximizes probability of real data being classified as real

Minimizes probability of fake data being classified as real

Generator



**Maximizes probability of fake data
being classified as real**

Loss Functions




Need optimizers for both networks

Loss function used is Binary Cross-Entropy (BCE) Loss

Used to **heavily penalize** incorrect classifications

Generator and Discriminator




Adversaries during training

At some point generator will generate realistic data


Consistently fool the discriminator

Leaky ReLU Activation

A Neural Network




Unresponsive Neurons



What if the weights of the connections do not change in response to changing input?


Unresponsive Neurons




Neurons may be dead

Activation Functions


ReLU




logit



tanh




step




Various choices of activation functions exist and drive the design of your neural network

Activation Functions




Consider an S-shaped (sigmoid) activation function

Activation Functions




This is the **active** or **responsive** region of the function

Saturating Activation Functions




The activation function saturates at either end

Saturating Activation Functions



If a neuron operates within these **saturation regions** throughout training it might become unresponsive

Dying Neurons




**Neuron might become unresponsive -
output won't change as input changes**

**If this continues throughout training,
neuron is “dead”**


**Saturation of neuron occurs at both ends
of S-curve, for instance**

Saturating Activation Functions



Logit Activation


Saturates for very large and very small values of input



ReLU Activation

Saturates for very small (negative values) of inputs

Leaky ReLU Activation




Mitigates dying-neuron problem of ReLU

For positive values, output same as input (like ReLU)

For negative values, output is non-zero, but close to zero

Leaky ReLU found to out-perform strict ReLU

Leaky ReLU Activation



$$\text{LeakyRelu}_{\alpha}(x) = \max(x, \alpha x)$$

α is a hyper parameter

Typically 0.01

Small non-zero value of α ensures that neuron does not “die”

Demo

**Using Generative Adversarial Networks
(GANs) to generate handwritten digits**

Summary

Understand Generative Adversarial Networks (GANs)


Candidate generation and evaluation

Role of generator

Role of discriminator

Generate handwritten digits using GANs trained on the MNIST dataset

Related Courses



Building Features from Image Data

Image Classification with PyTorch

**Expediting Deep Learning with Transfer
Learning: PyTorch Playbook**