

Natural Language Processing with PyTorch

IMPLEMENTING RECURRENT NEURAL NETWORKS (RNNS)
IN PYTORCH



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Modifying neurons to endow them with state and memory

Understand Recurrent Neural Networks (RNNs)

Mitigate problems of vanishing and exploding gradients in training RNNs

Working with long-memory recurrent cells

Use RNNs in language modeling

Prerequisites and Course Outline

Prerequisites

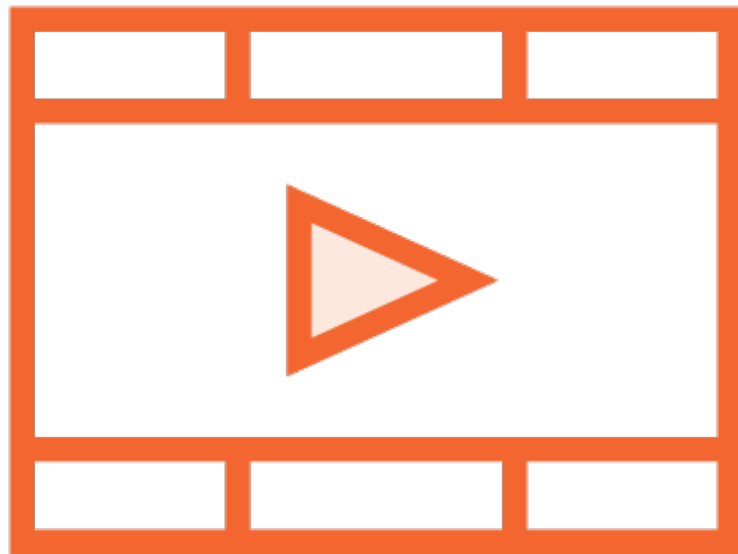


Comfortable programming in Python

Good understanding of neural networks

Used PyTorch to build and train neural networks

Prerequisite Courses



Foundations of PyTorch

Building Your First PyTorch Solution

Image Classification with PyTorch

Course Outline



Recurrent neural networks (RNNs)

Binary classification using words

Multi-class classification using characters

Sentiment analysis using word embeddings

Sequence-to-sequence models

RNNs and Natural Language Processing

$$y = f(x)$$

Machine Learning

Machine learning algorithms seek to “learn” the function f that links the features and the labels

$$y = Wx + b$$

$$f(x) = Wx + b$$

Linear regression specifies, up-front, that the function f is linear

```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

$f(x) = \text{doSomethingReallyComplicated}(x)$

ML algorithms such as neural network can “learn” (reverse-engineer) pretty much anything given the right training data

Sometimes **time** relationships in
data have special meaning

$$y_t = f(x_t, y_{t-1})$$

Learning the Past

Relationships where past values of the effect variable drive current values are called auto-regressive

$$y_t = f(x_t, y_{t-1})$$

Learning the Past

The output at one time instance depends on the current input at that time instance

$$\mathbf{y}_t = f(\mathbf{x}_t, \mathbf{y}_{t-1})$$

Learning the Past

And on the output from the previous time instance

Feed-forward networks cannot
learn from the past

Recurrent neural networks can

Text Is Sequential Data



Predict the next word in a sequence (autocomplete)

“The tallest building in the world is ...”



Language translations

“how are you” -> “Comment allez-vous”




Text classification, sentiment analysis, natural language processing

“This is not the worst restaurant not by a long way”


RNNs are great at learning
sequential data

Recurrent Neurons


Simplest Feed-forward Neuron



Simplest Recurrent Neuron



Recurrent Neuron




y_t = Output at time t


Depends upon

- y_{t-1} = Output at time t - 1
- x_t = New inputs available only at time t


Unrolling Through Time




Unrolling Through Time




Unrolling Through Time




Unrolling Through Time




Unrolling Through Time



Output of a Layer Fed to Next Layer



Recurrent Neuron



Regular neuron: input is feature vector,
output is scalar


$$Y = Wx + b$$

Recurrent neuron: **output is vector too**

Input: $[X_0, X_1, \dots X_t]$

Output: $[Y_0, Y_1, \dots Y_t]$

Memory and State




Recurrent neurons remember the past

They possess 'memory'

The stored state could be **more complex than simply y_{t-1}**

The internal state is represented by **h_t**


Recurrent Neuron



Now, each neuron has two weight vectors

W_x, W_y


Recurrent Neuron



Now, each neuron has two weight vectors

w_x , w_y


Recurrent Neuron



Now, each neuron has two weight vectors

W_x , W_y

Recurrent Neuron




Output of neuron as a whole is given as

$$y_t = \Phi(X_t W_x + y_{t-1} W_y + b)$$

(Φ is the activation function)


Training a Recurrent Neural Network

Layer of Recurrent Neurons



A layer of neurons forms an RNN cell - basic cell, LSTM cell, GRU cell (more on these later)


Layer of Recurrent Neurons



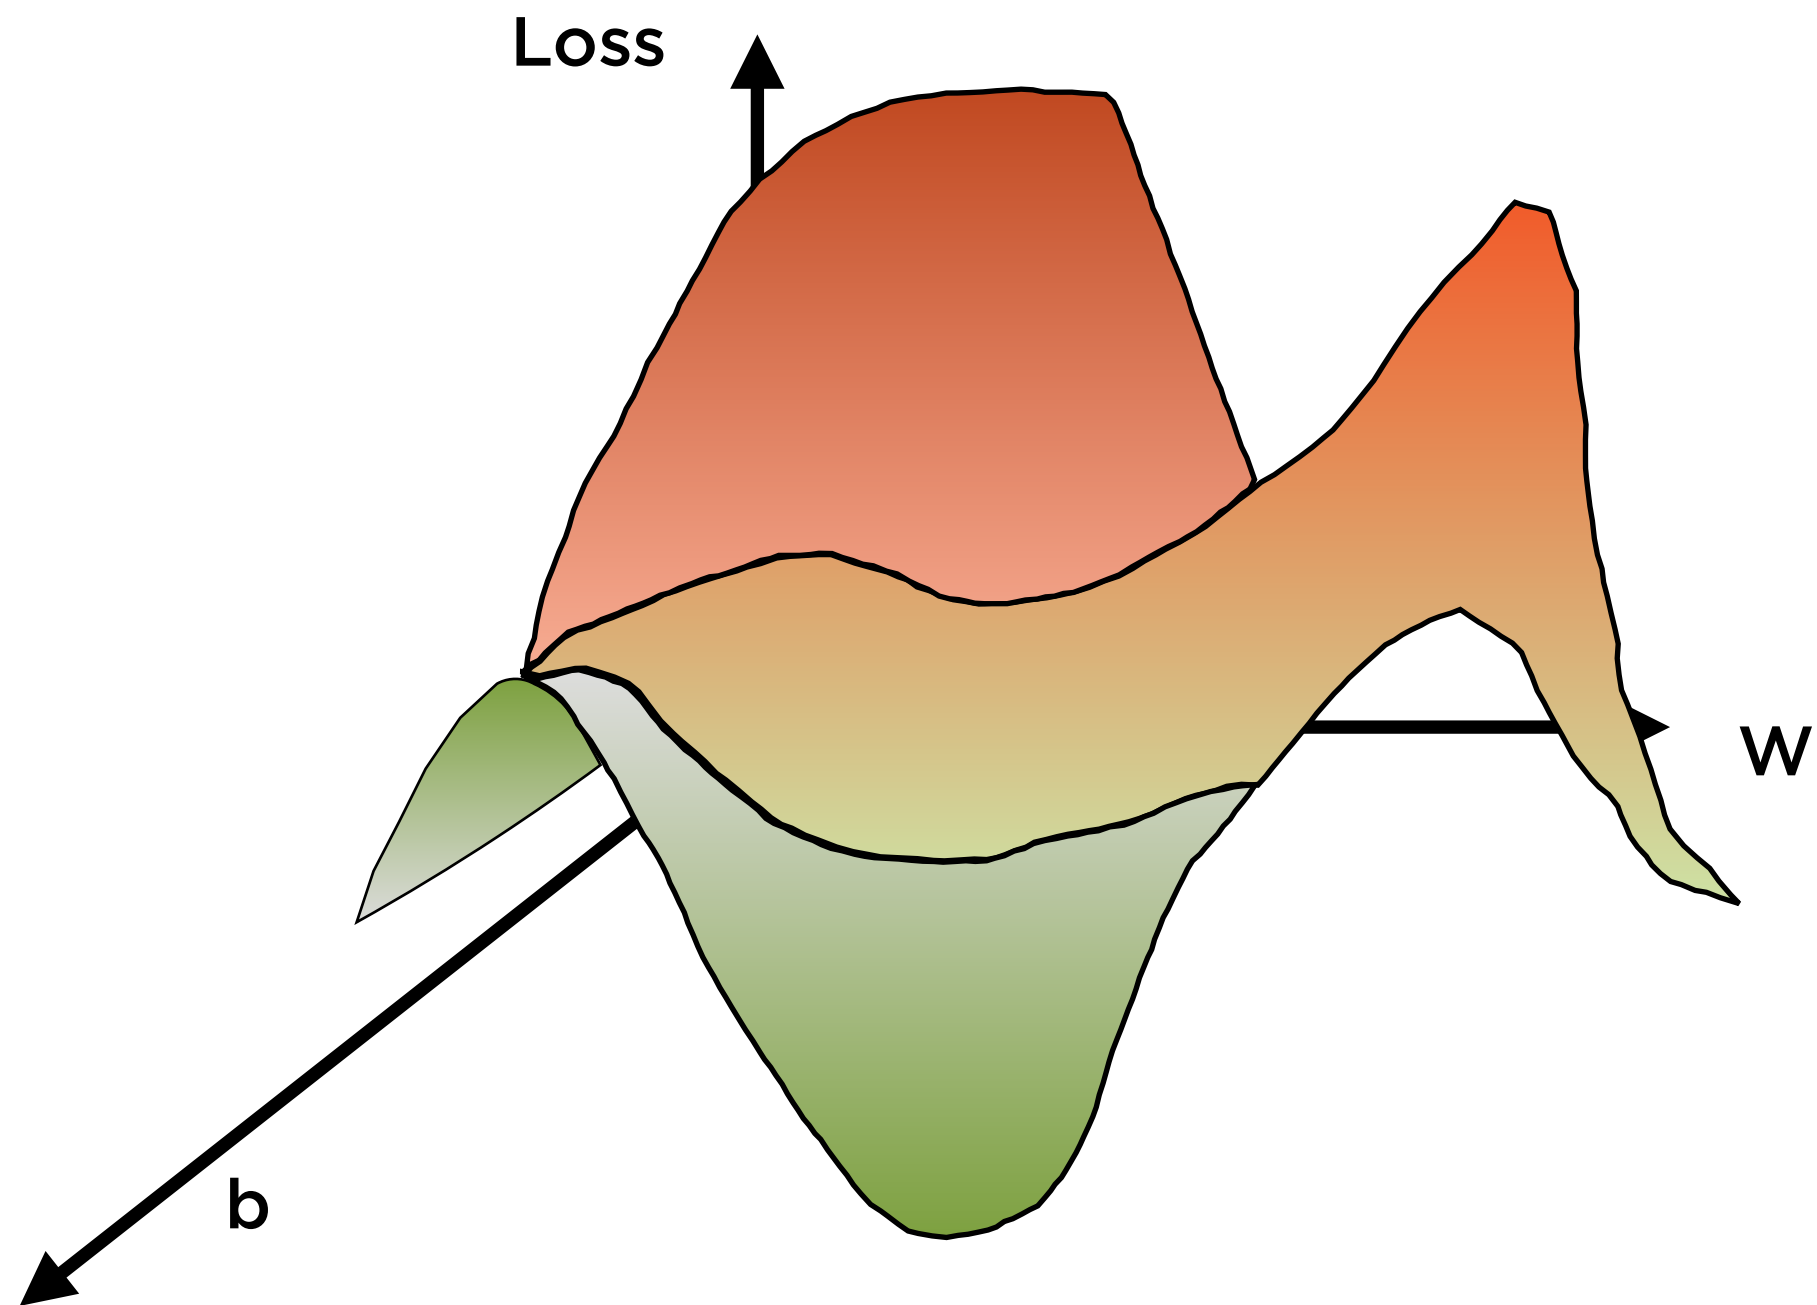
The cells unrolled through time form the layers of the
neural network

The actual training of a neural network happens via Gradient Descent Optimization


Minimizing Loss



Minimizing Loss




Minimizing Loss




Gradient Descent


Converging on the
“best” value using an
optimization algorithm




Back Propagation Through Time




Back Propagation Through Time




Back Propagation Through Time




Back Propagation Through Time



Back Propagation Through Time




Back Propagation Through Time



Back propagation allows the weights and biases of the neurons to **converge** to their final values

Back Propagation




This is an iterative process

Fails either if:

- gradients don't change at all
- gradients change too fast

BPTT




BPTT is the **tweaked** version of Back Propagation for RNN training

Need as many layers as past time periods

Vanishing and Exploding Gradients

Layer of Recurrent Neurons




$t=0$

$t=1$


time

Each of the RNN layers is a
cell made up of neurons,
unrolled through time

Back Propagation Through Time (BPTT)



Back Propagation




This is an iterative process

Fails either if:

- gradients don't change at all
- gradients change too fast

BPTT



BPTT is the **tweaked** version of Back Propagation for RNN training

Need as many layers as past time periods

$$y_t = f(x_t, y_{t-1}, y_{t-2})$$

Learning the (Recent) Past

Unrolling the RNN through time helps learn the past

$$y_t = f(x_t, y_{t-1}, y_{t-2} \dots, y_{t-1000})$$

Learning the Distant Past

The unrolled RNN will be very, very deep - many layers to train, the gradient has to be propagated a long way

$$y_t = f(x_t, y_{t-1}, y_{t-2} \dots, y_{t-1000})$$

Learning the Distant Past


The unrolled RNN will be very, very deep - many layers to train, the gradient has to be propagated a long way

Recurrent neural networks may be
unrolled **very far back in time**

They're prone to the **vanishing** and
exploding gradients issue


Vanishing Gradient Problem

Gradient becomes zero
and stops changing




Exploding Gradient Problem

Gradient changes abruptly and “explodes”




Vanishing and Exploding Gradients



Back propagation fails if:

- gradients are **vanishing**
- gradients are **exploding**

Training RNNs



Training RNNs poses some specific challenges

Coping with Vanishing/Exploding Gradients


Proper initialization

**Non-saturating activation
function**

Batch normalization

Gradient clipping


BPTT



If output relies on distant past
Vanishing/exploding gradients
very likely

One option - **truncated BPTT**

Truncated BPTT




Simply truncate input sequence

E.g. predict stock movement tomorrow:

- Use only last week's data
- Do not use data for last year at all
- Daily data for last week, monthly before that

Truncated BPTT




Truncated BPTT can kill prediction performance

What if stock move tomorrow depends on stock move on last quarter-ending date?


Use **long-memory** cells to store
additional state in neuron

Long-memory Cells


Simplest Recurrent Neuron




Simplest Recurrent Neuron



Long Memory Recurrent Neuron



Long Memory RNNs




Increase the amount of state in neuron
Effect is to increase memory of neuron

Could explicitly add:

- long-term state (c)
- short-term state (h)

Long memory neurons have several advantages over basic RNNs

Long Memory RNNs



Advantages in Training

Faster training, nicer gradients



Advantages in Prediction

No need to truncate BPTT

Long-term Dependencies in Text

The sky is **cloudy**, it looks like



The gap between the relevant information needed to
predict the next word is small

Long-term Dependencies in Text

I've lived in **France** a long time. I first went there as a tourist, then applied for a job and I've been here ever since. I now speak fluent



The context for which language to predict here is much farther back in the sentence

Long-term Dependencies in Text

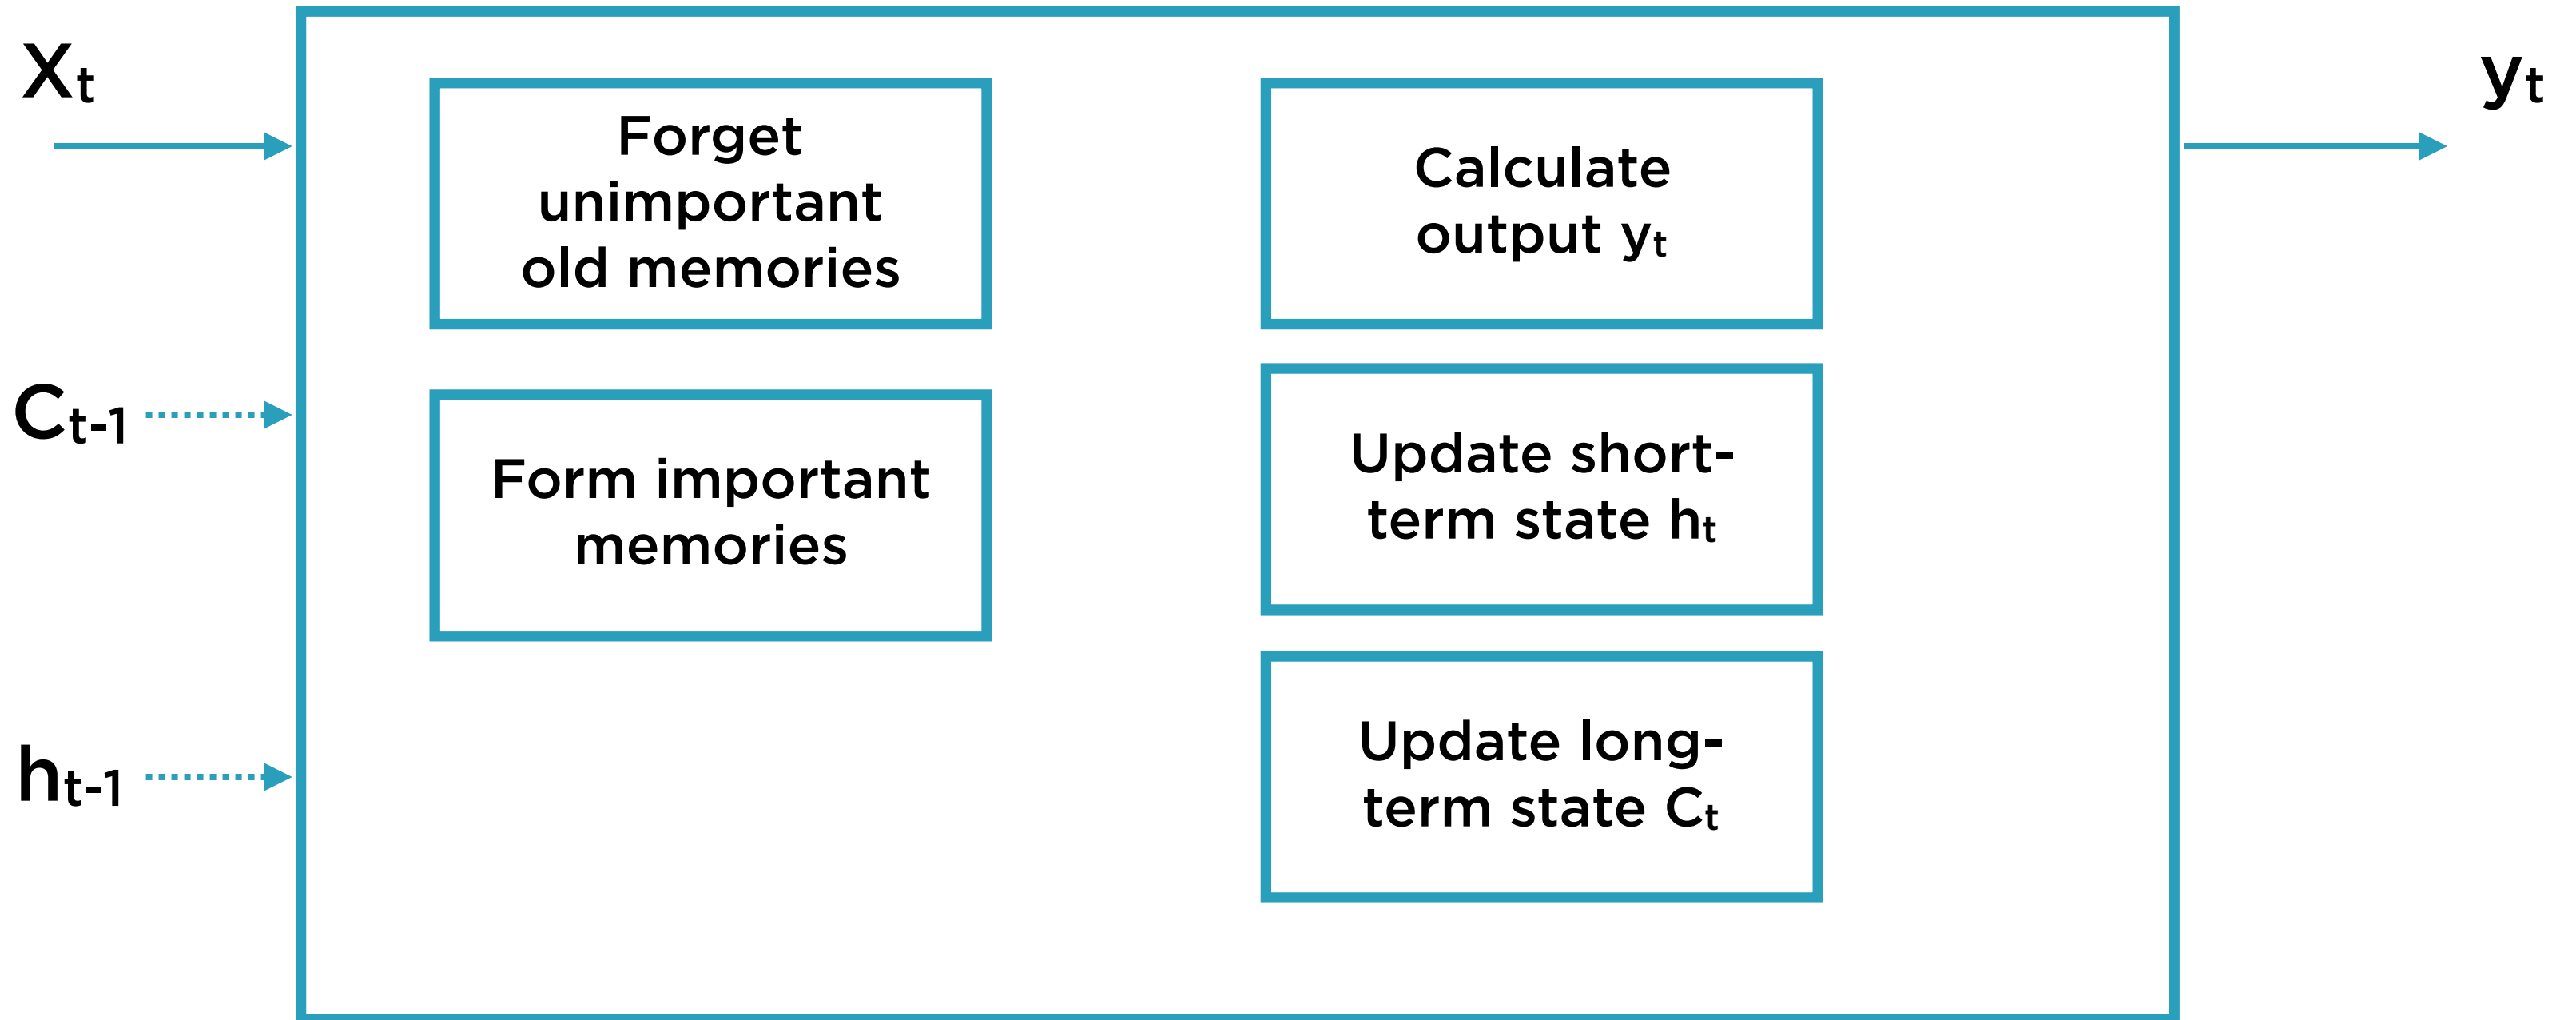
I've lived in **France** a long time. I first went there as a tourist, then applied for a job and I've been here ever since. I now speak fluent



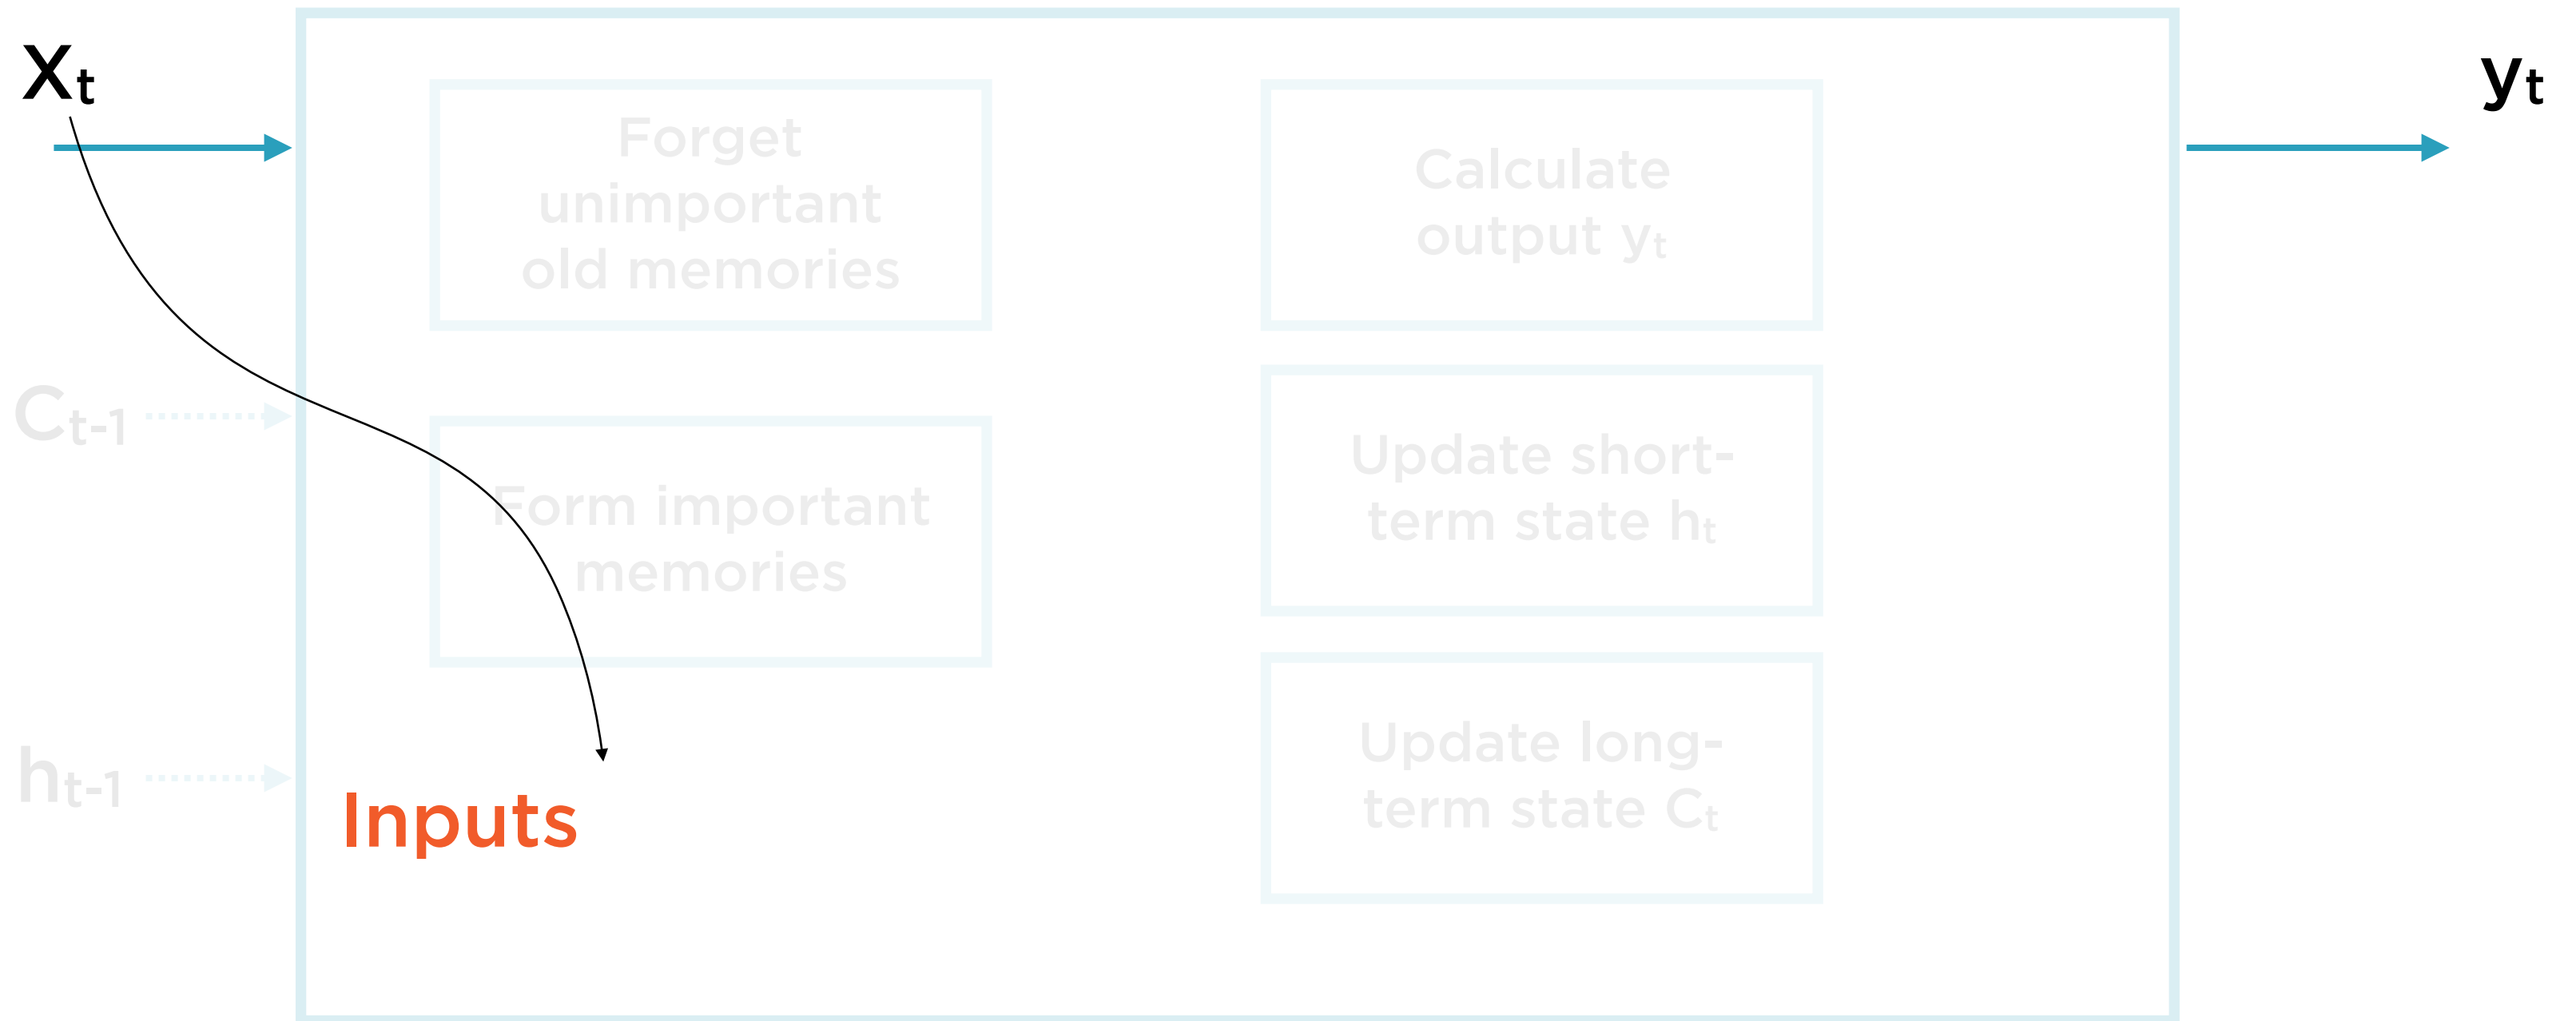
As this gap grows, RNNs become unable to learn to connect the information

Long/Short-Term Memory Cell
(**LSTM**) - a popularly used long
memory cell in RNNs

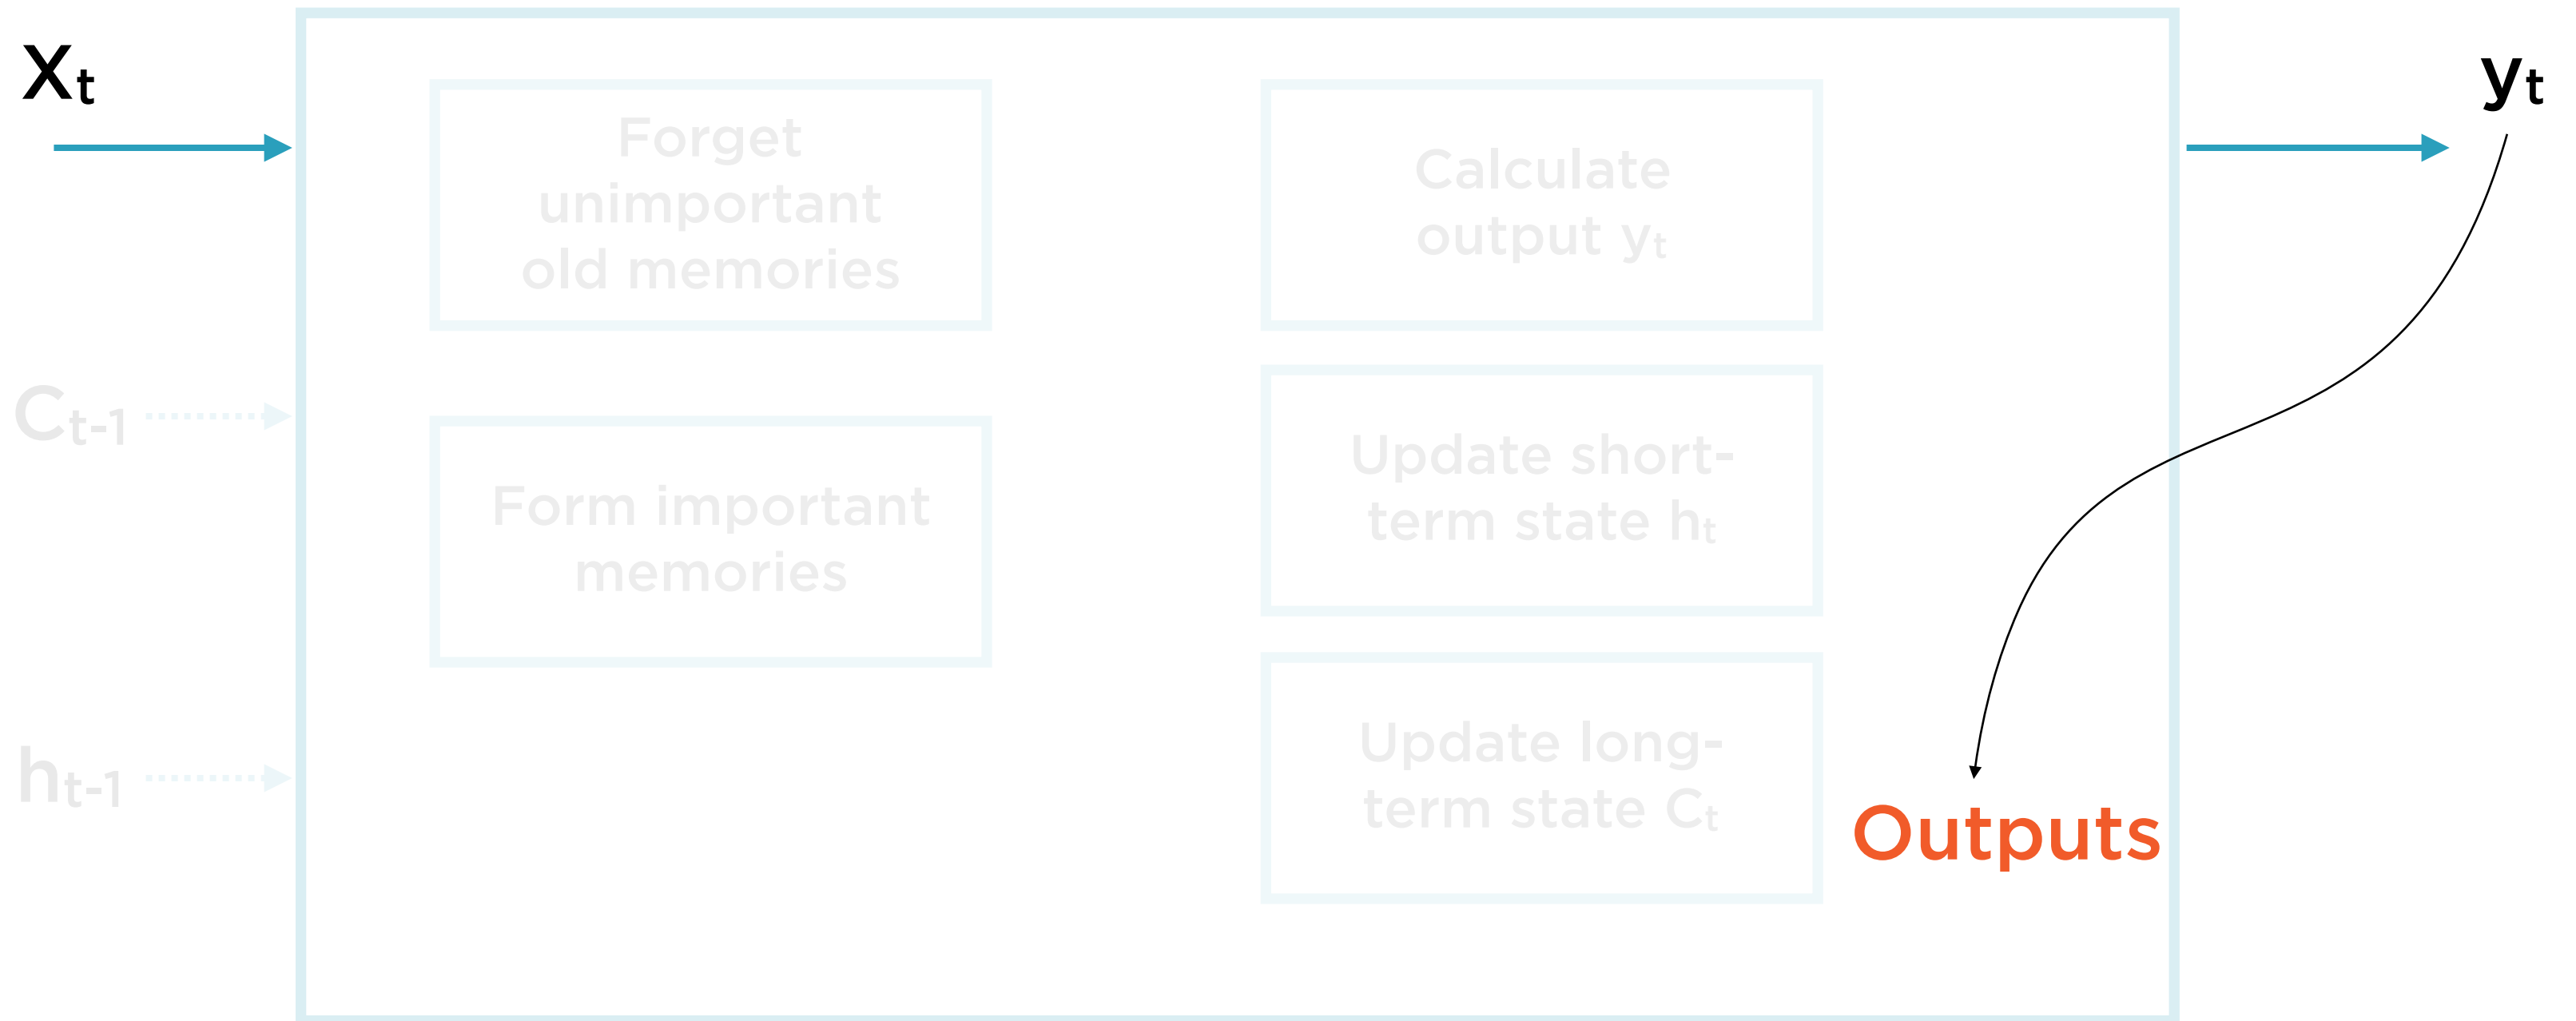
LSTM



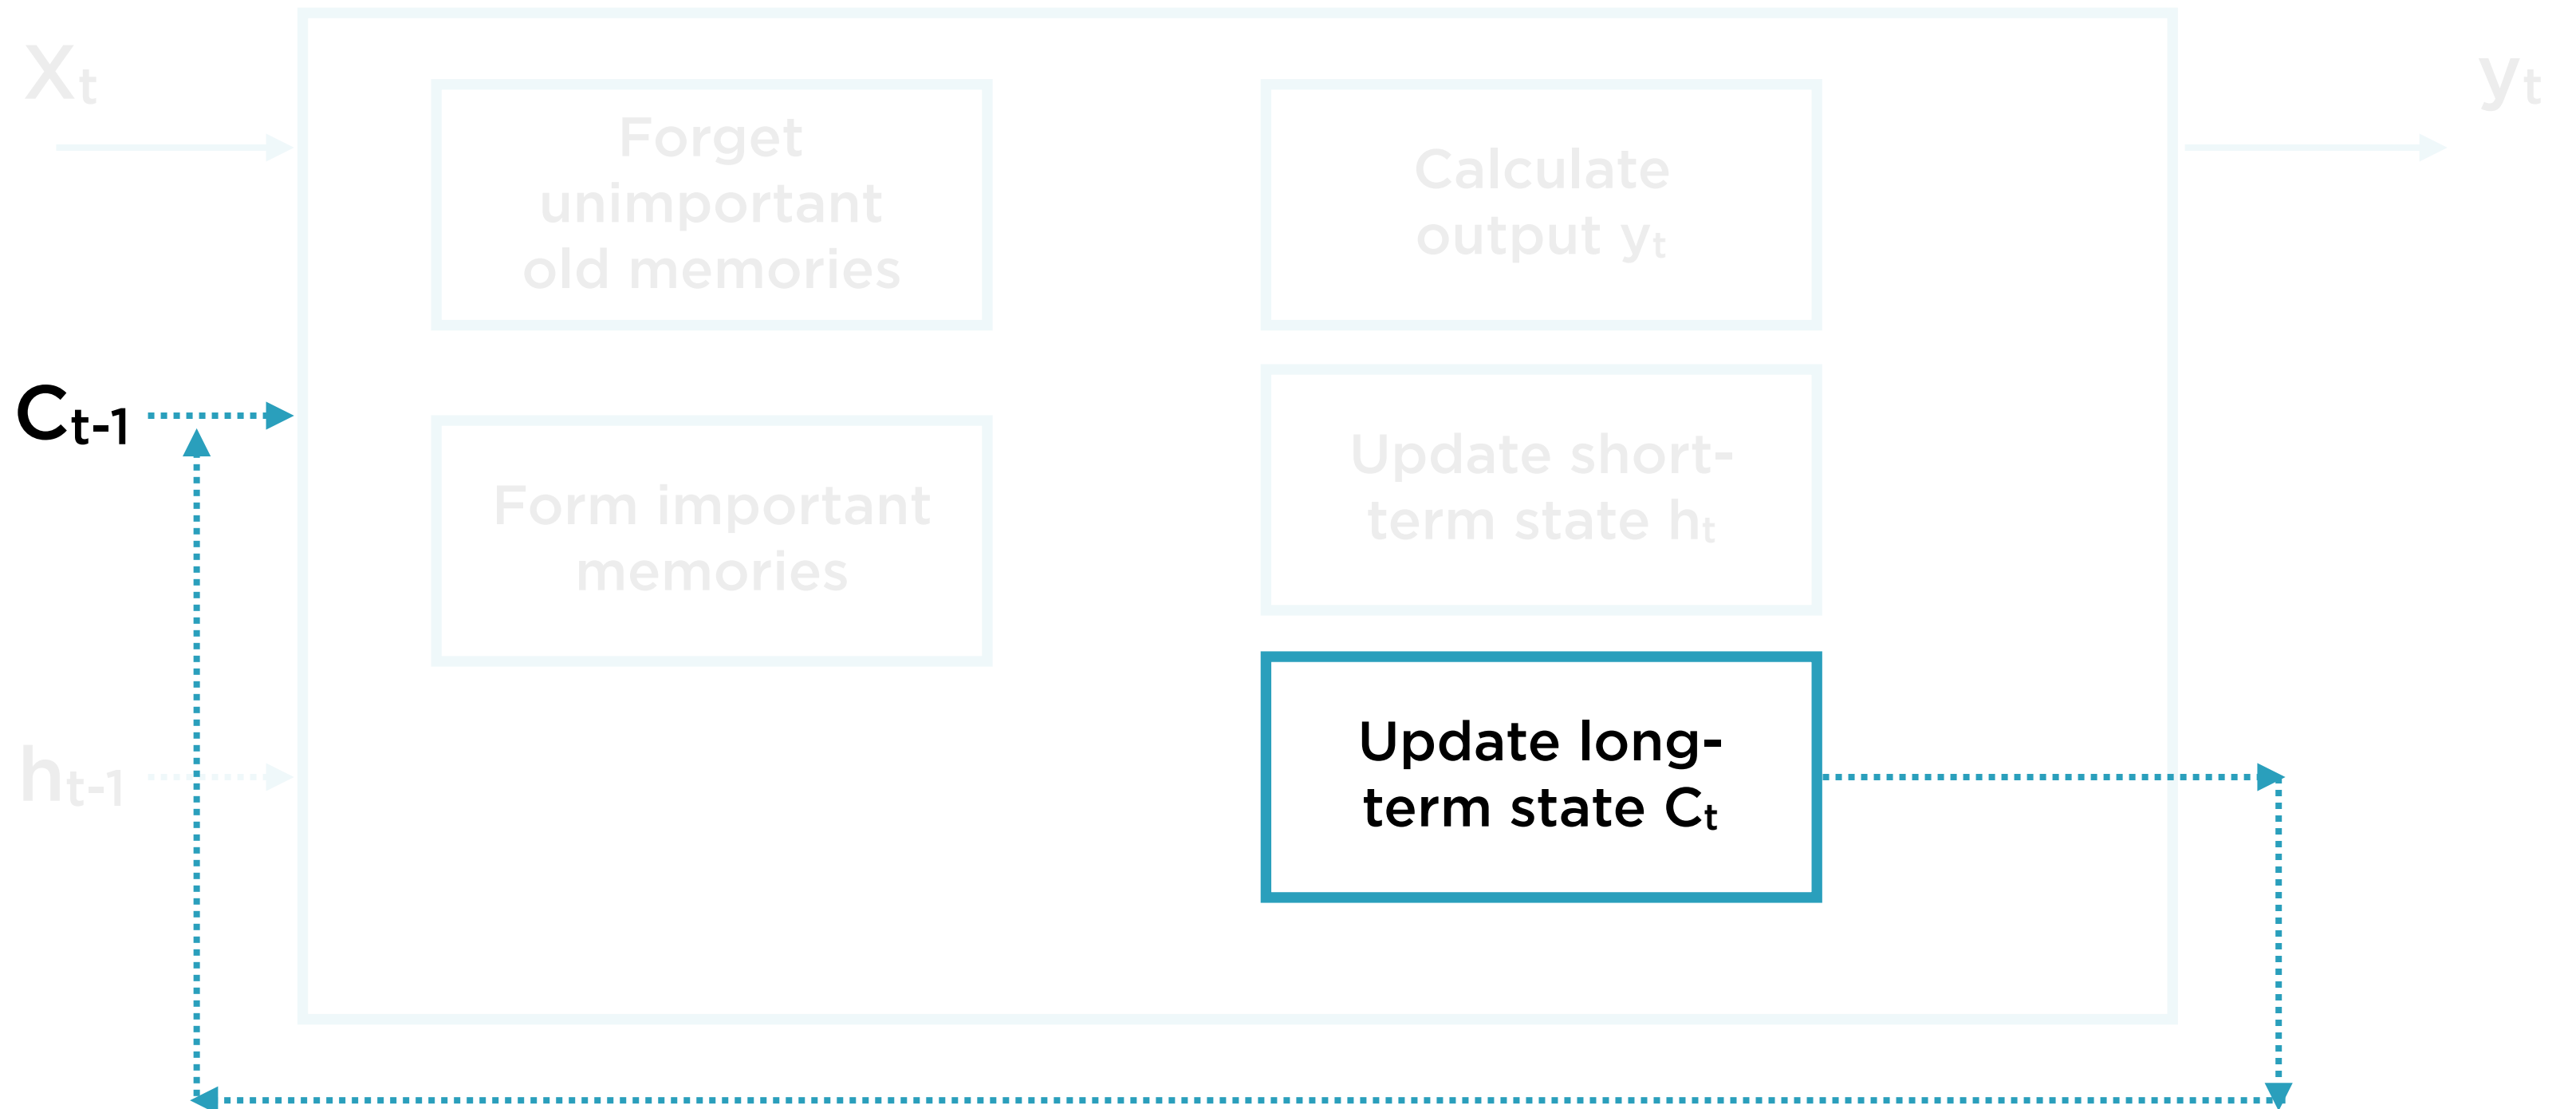
LSTM



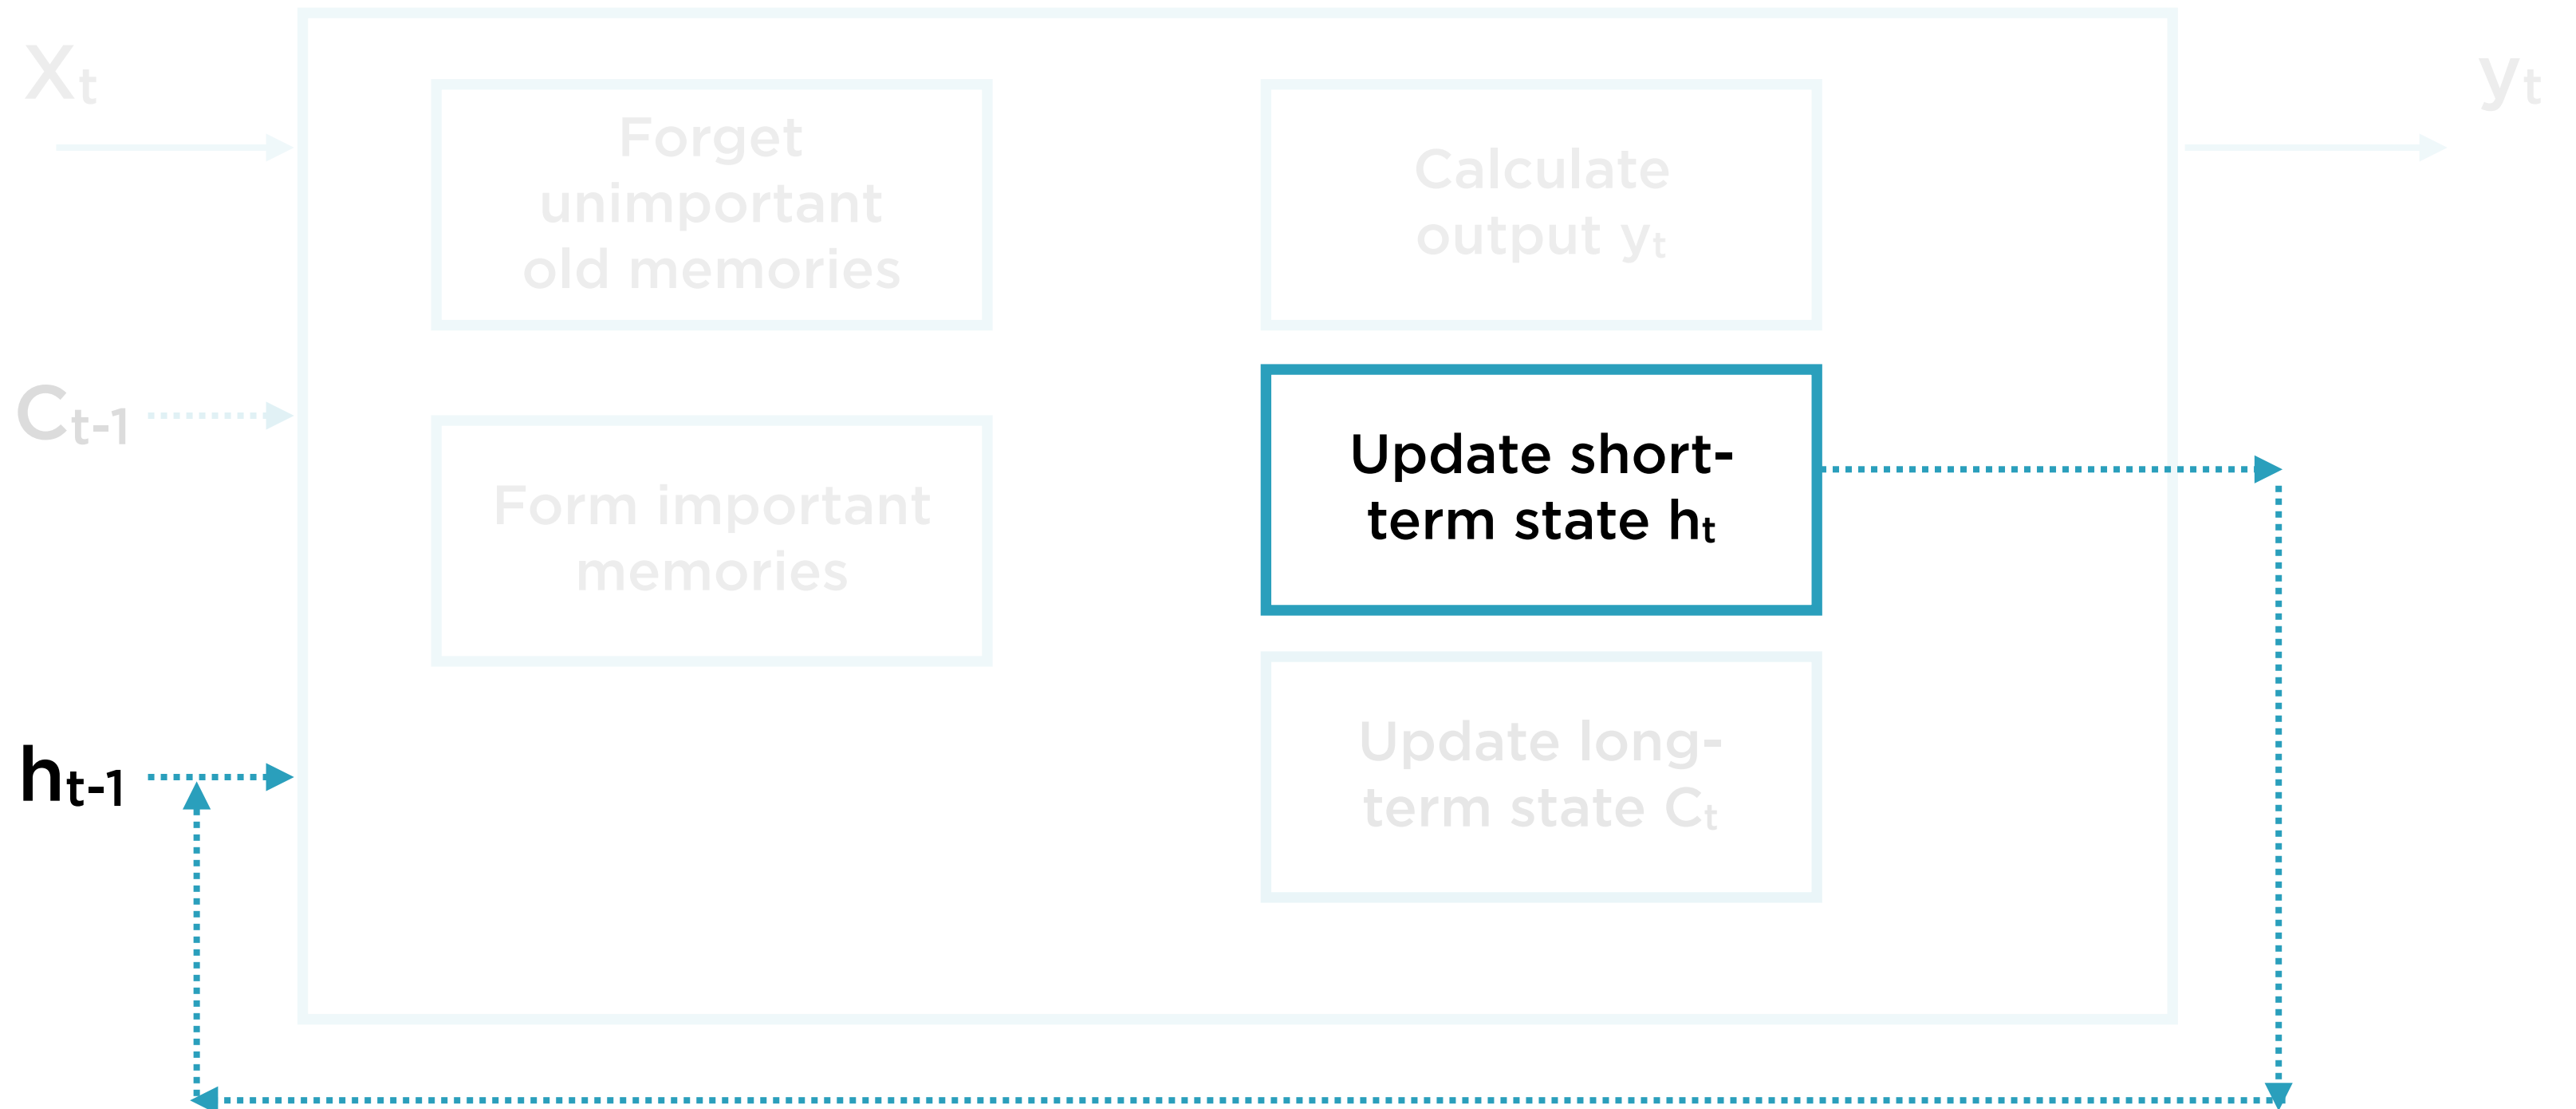
LSTM



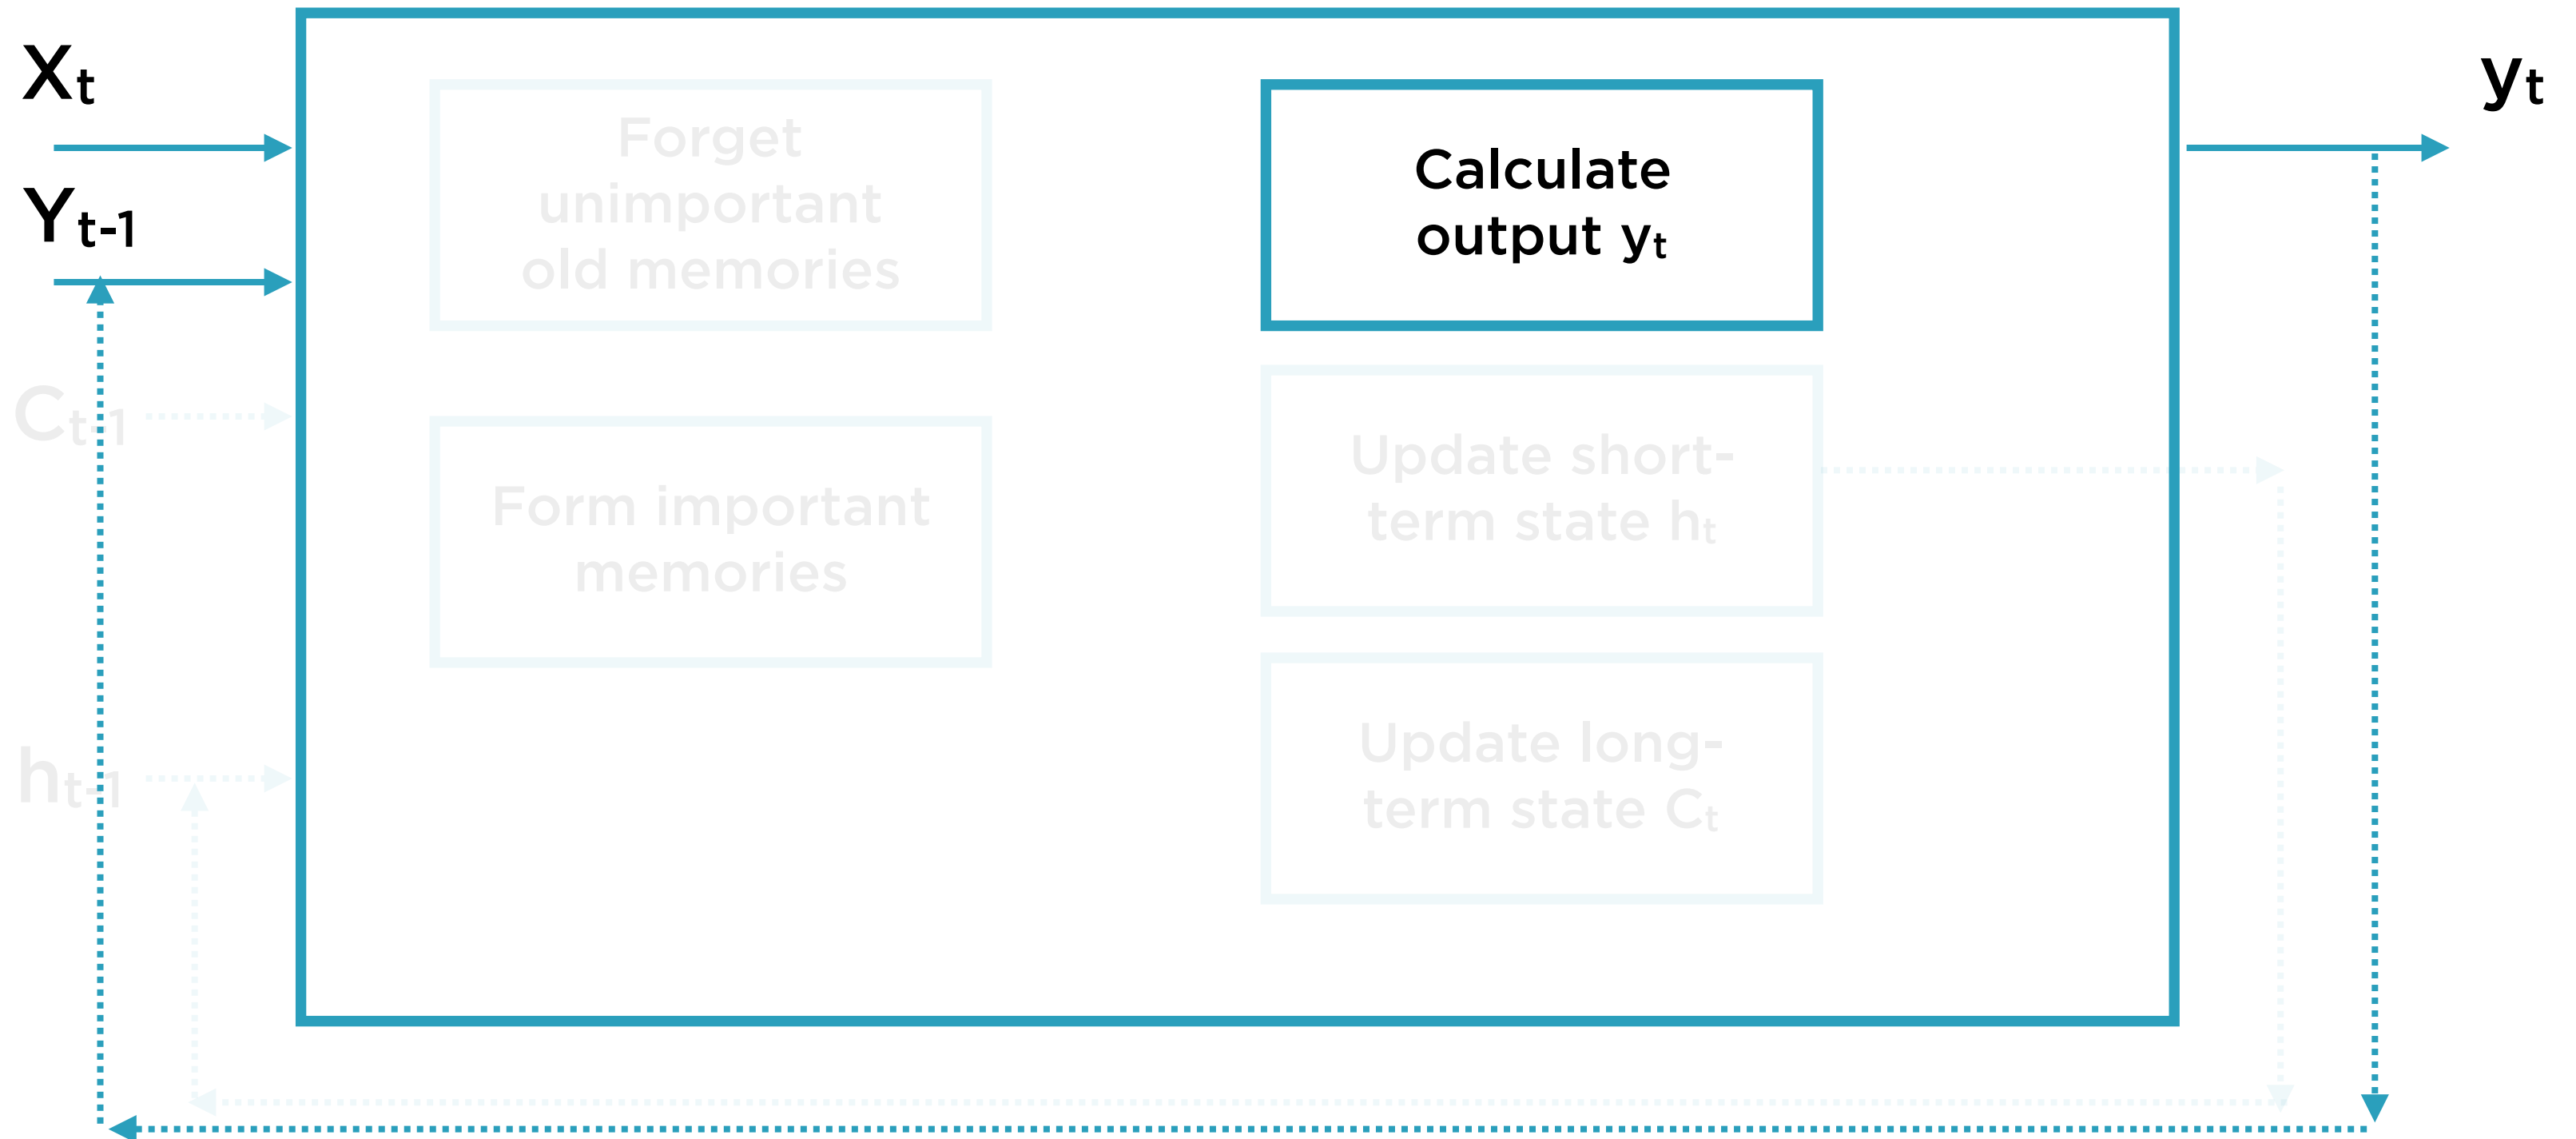
LSTM



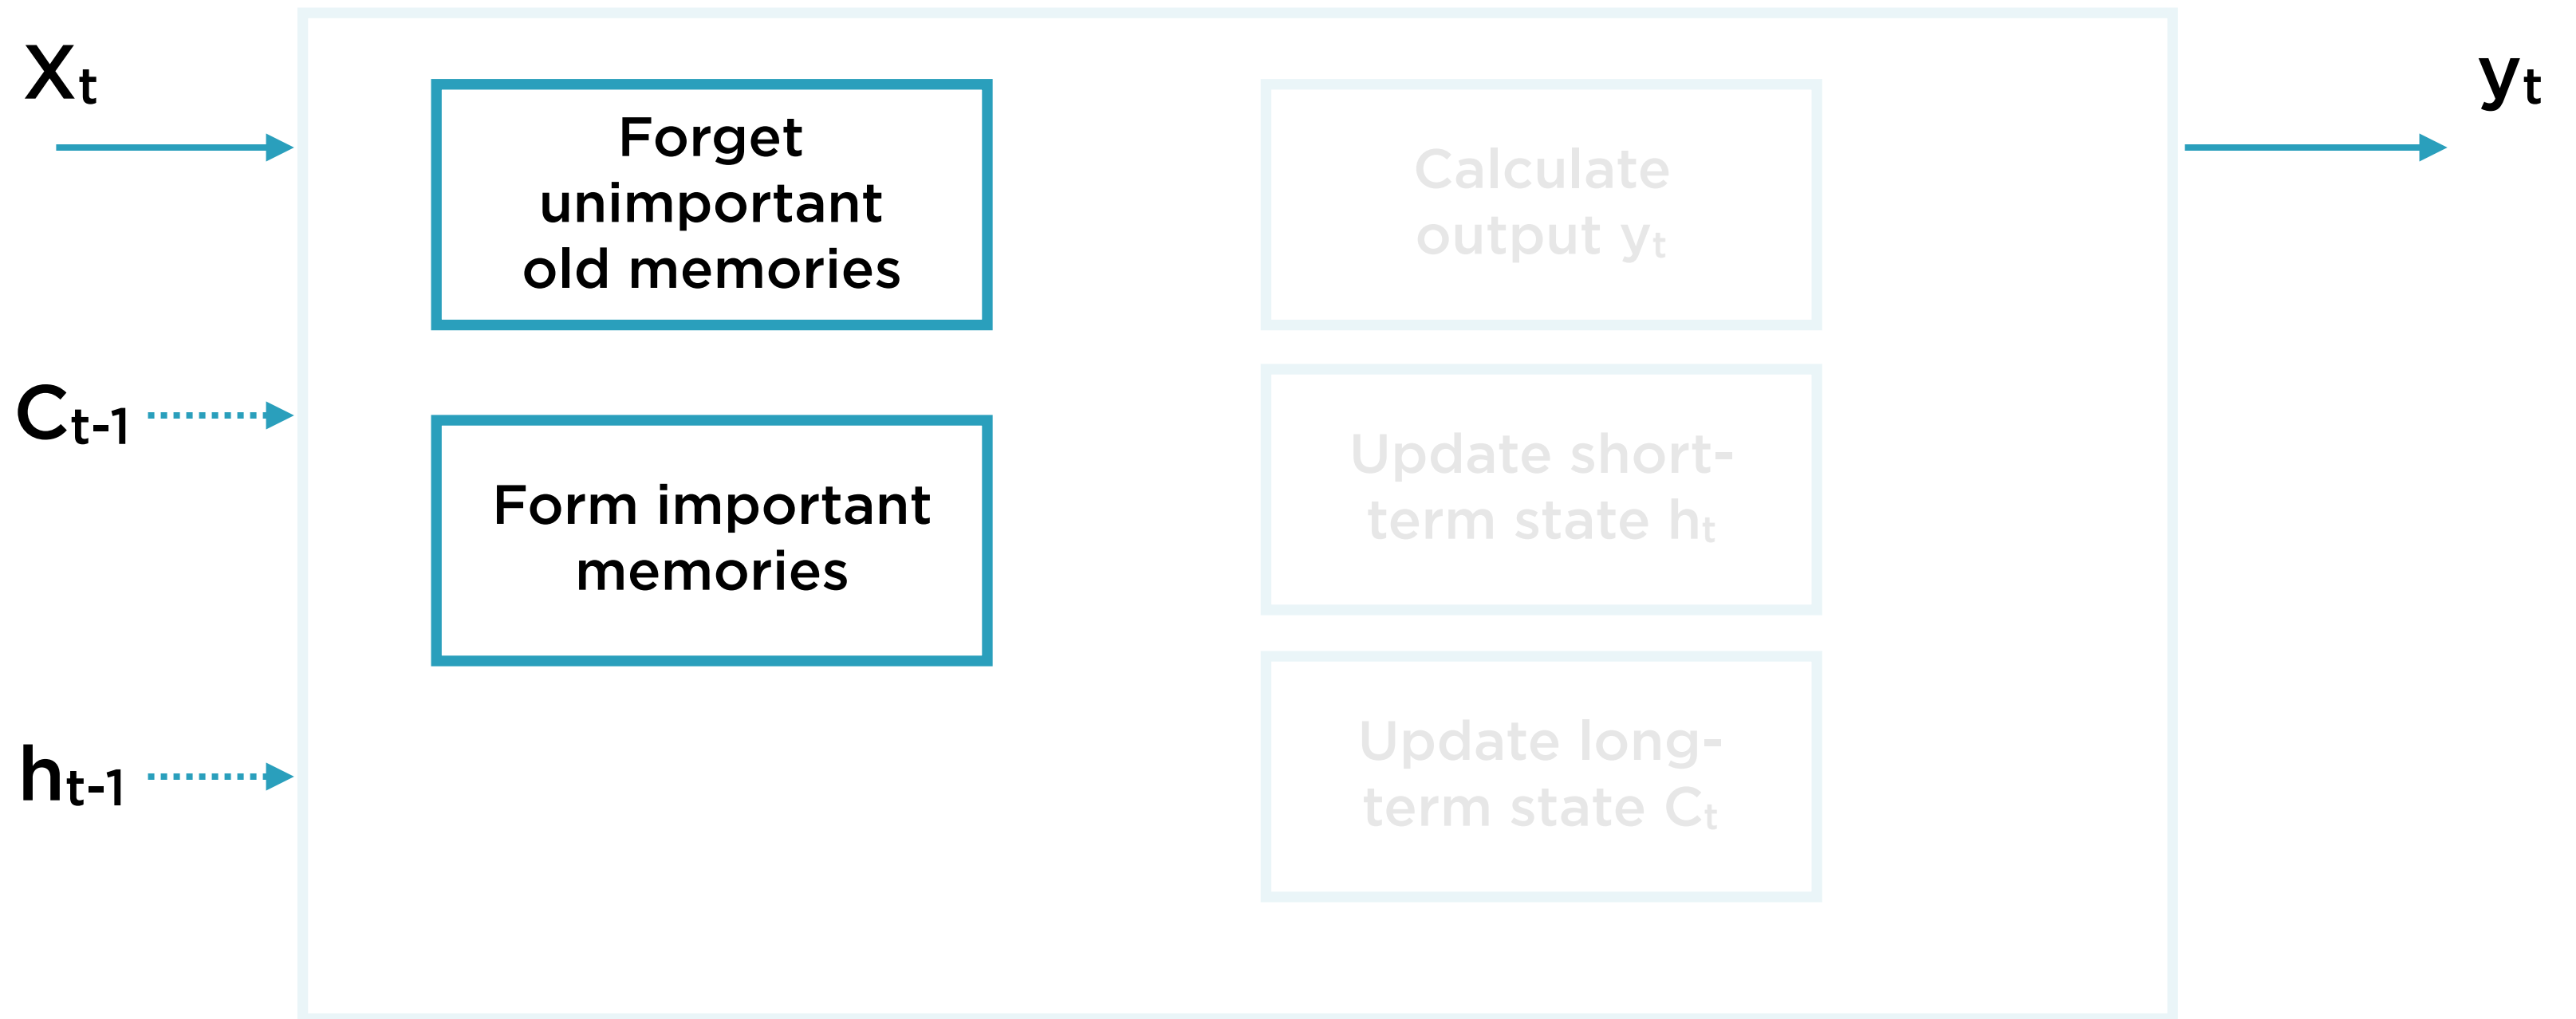
LSTM



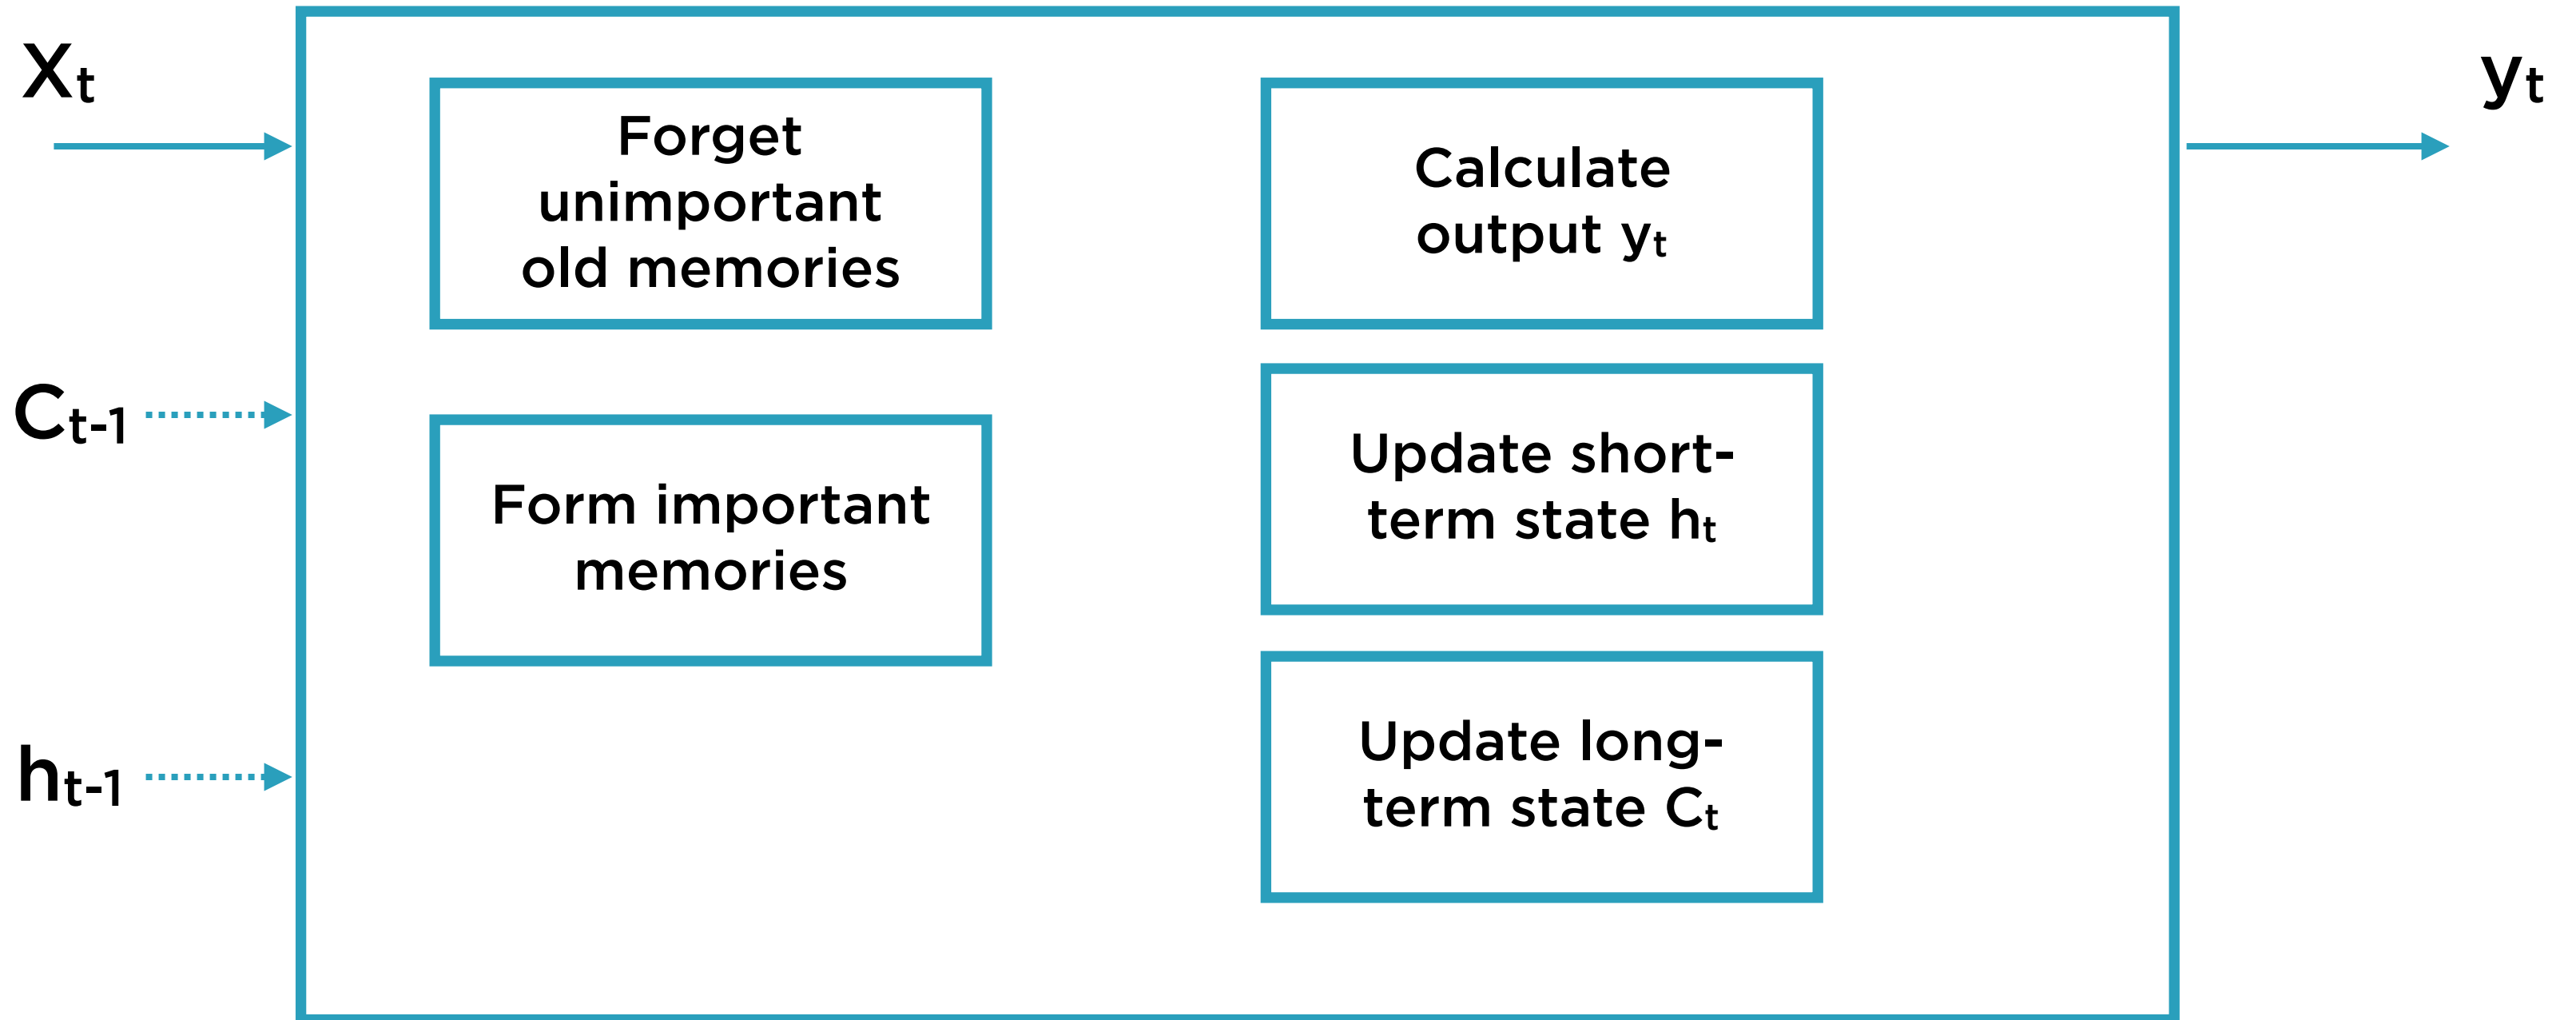
Basic RNN Cell



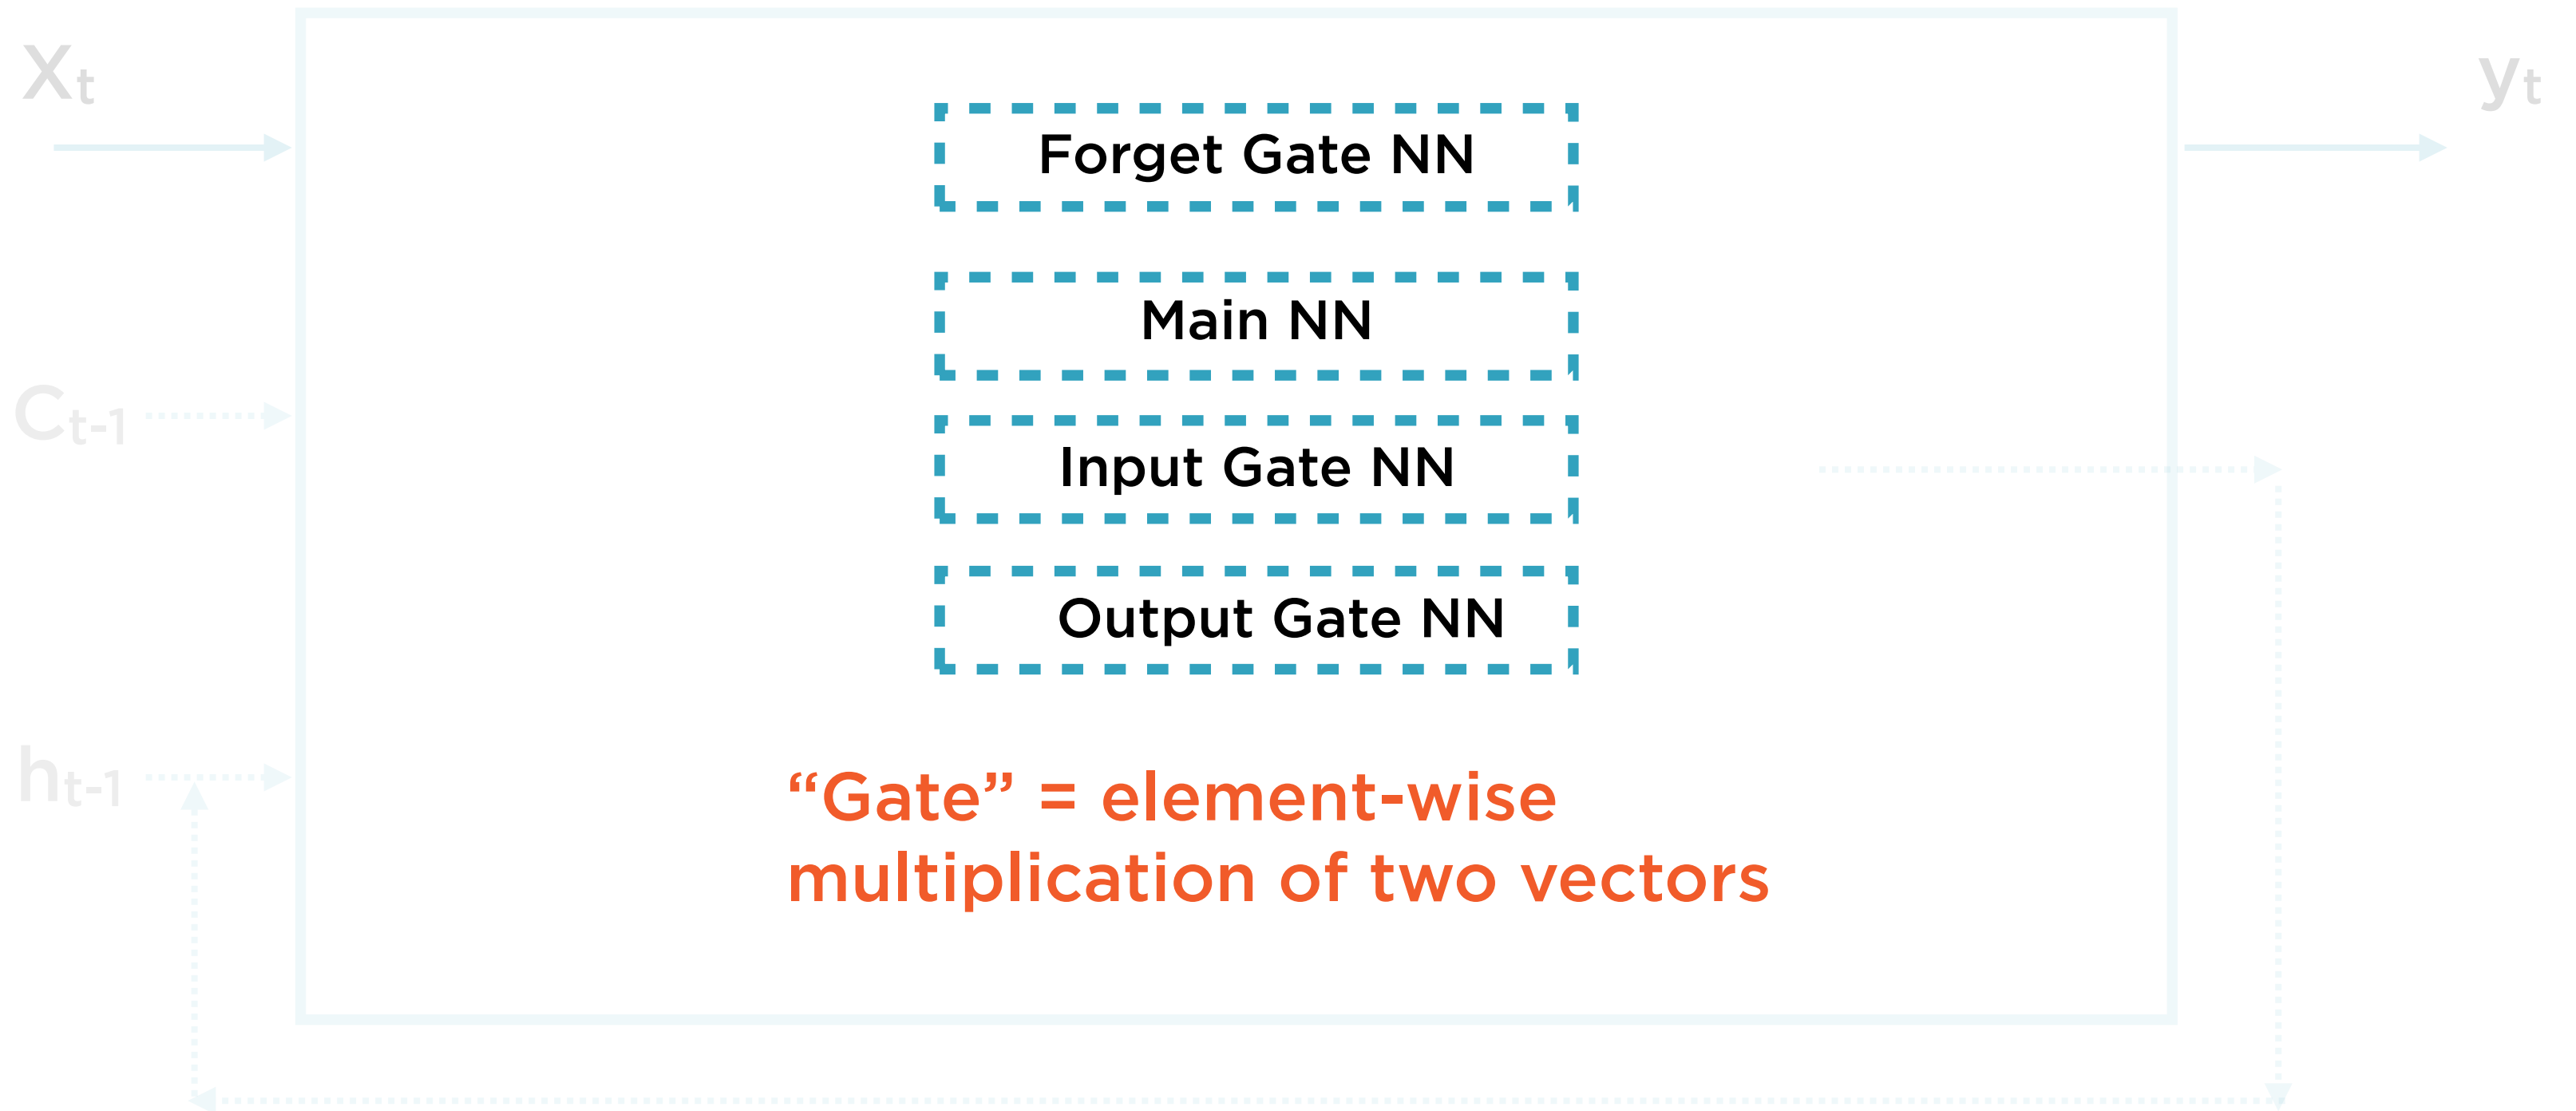
LSTM



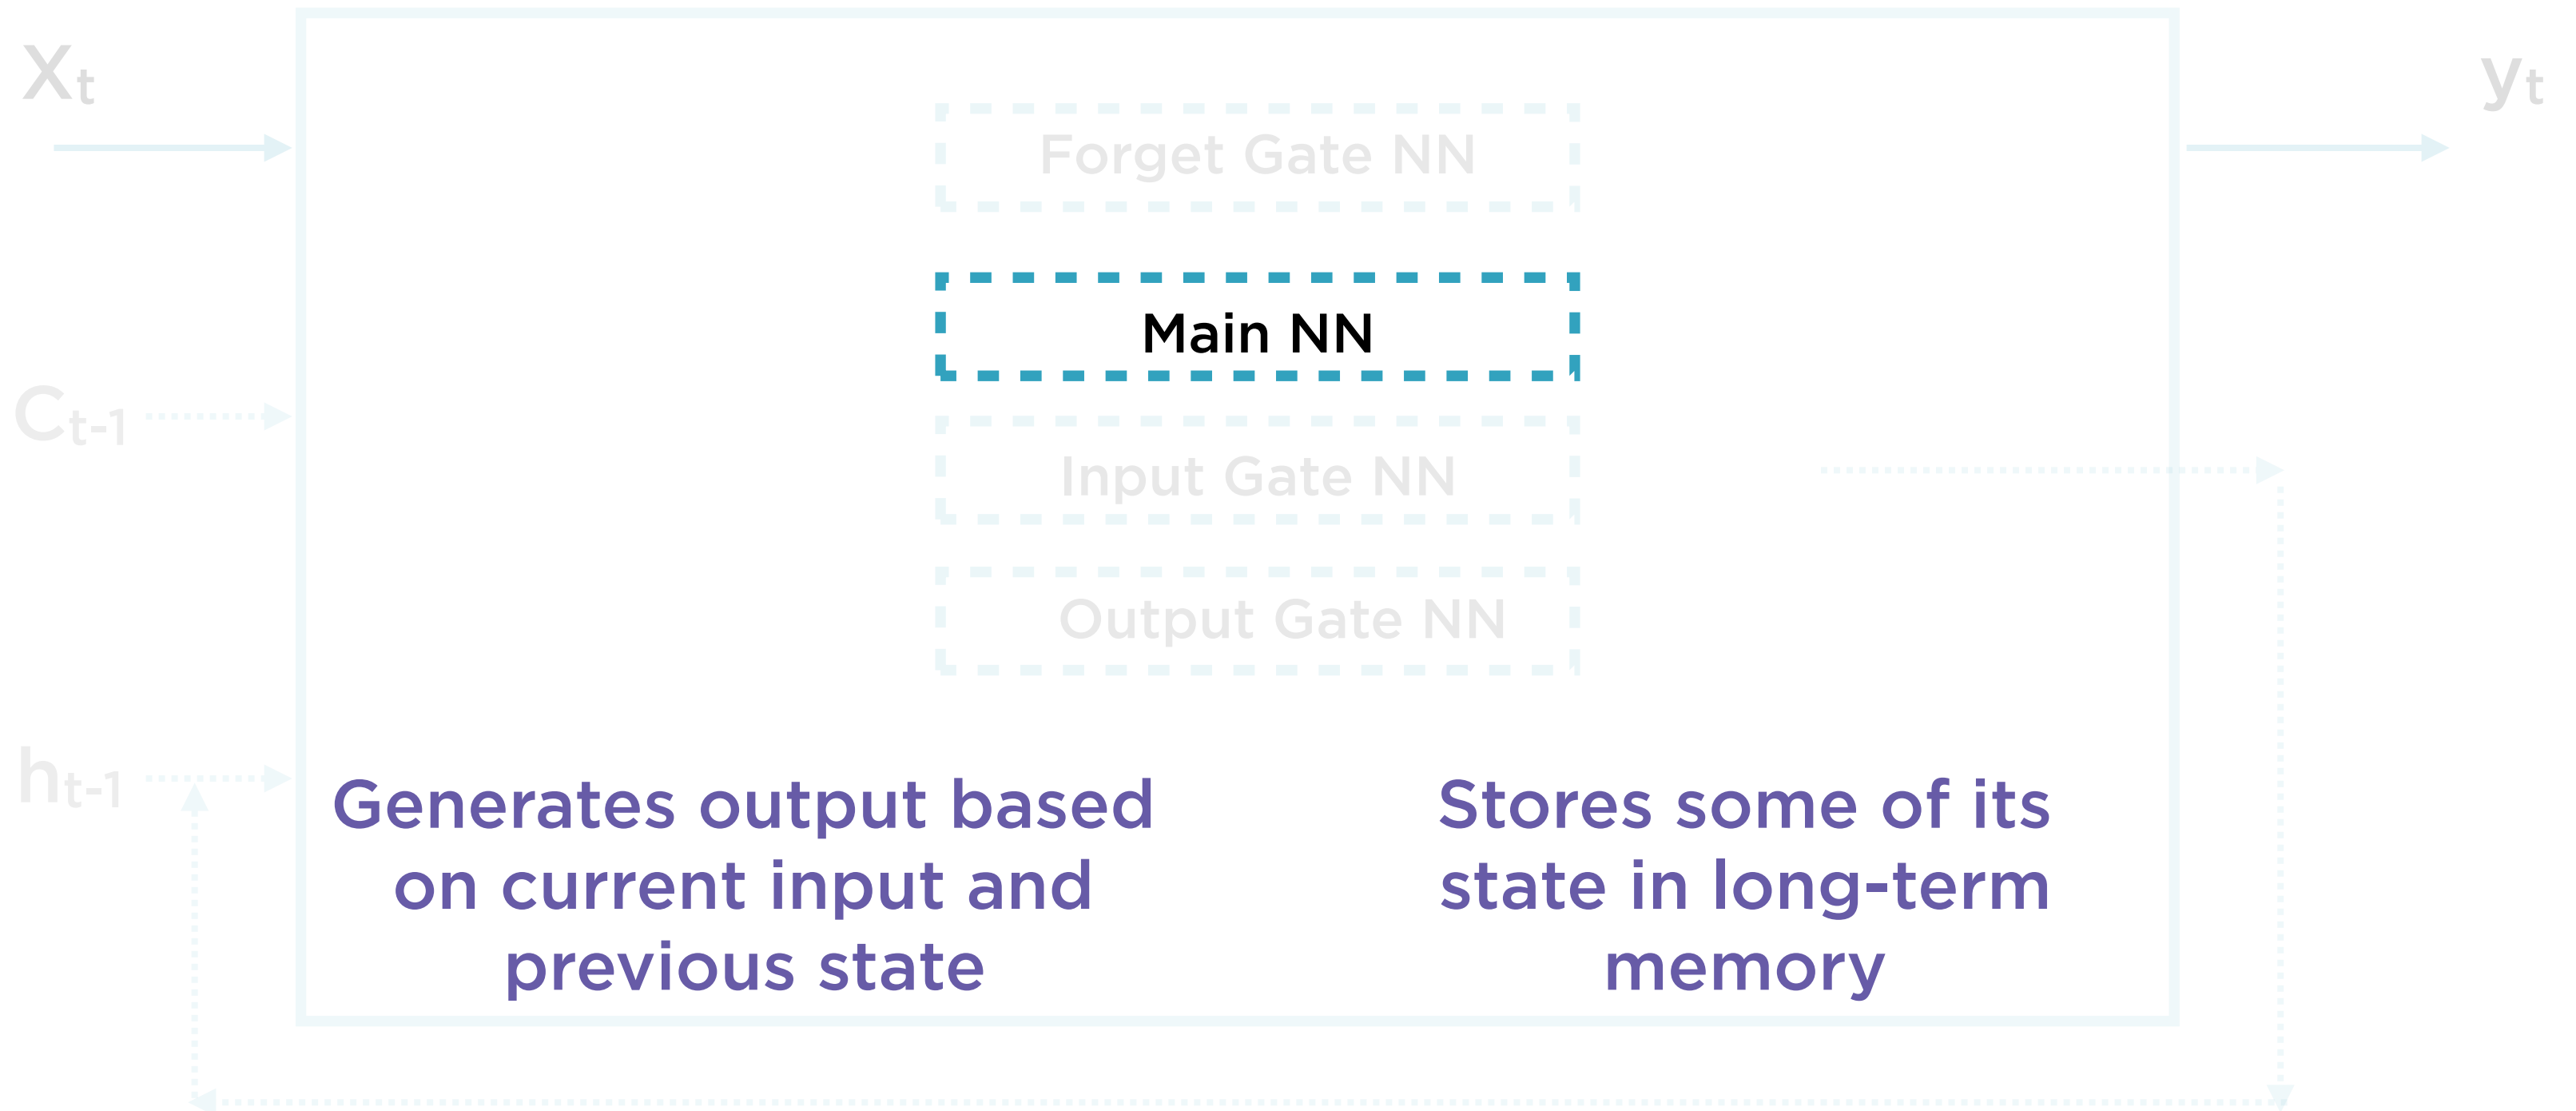
LSTM



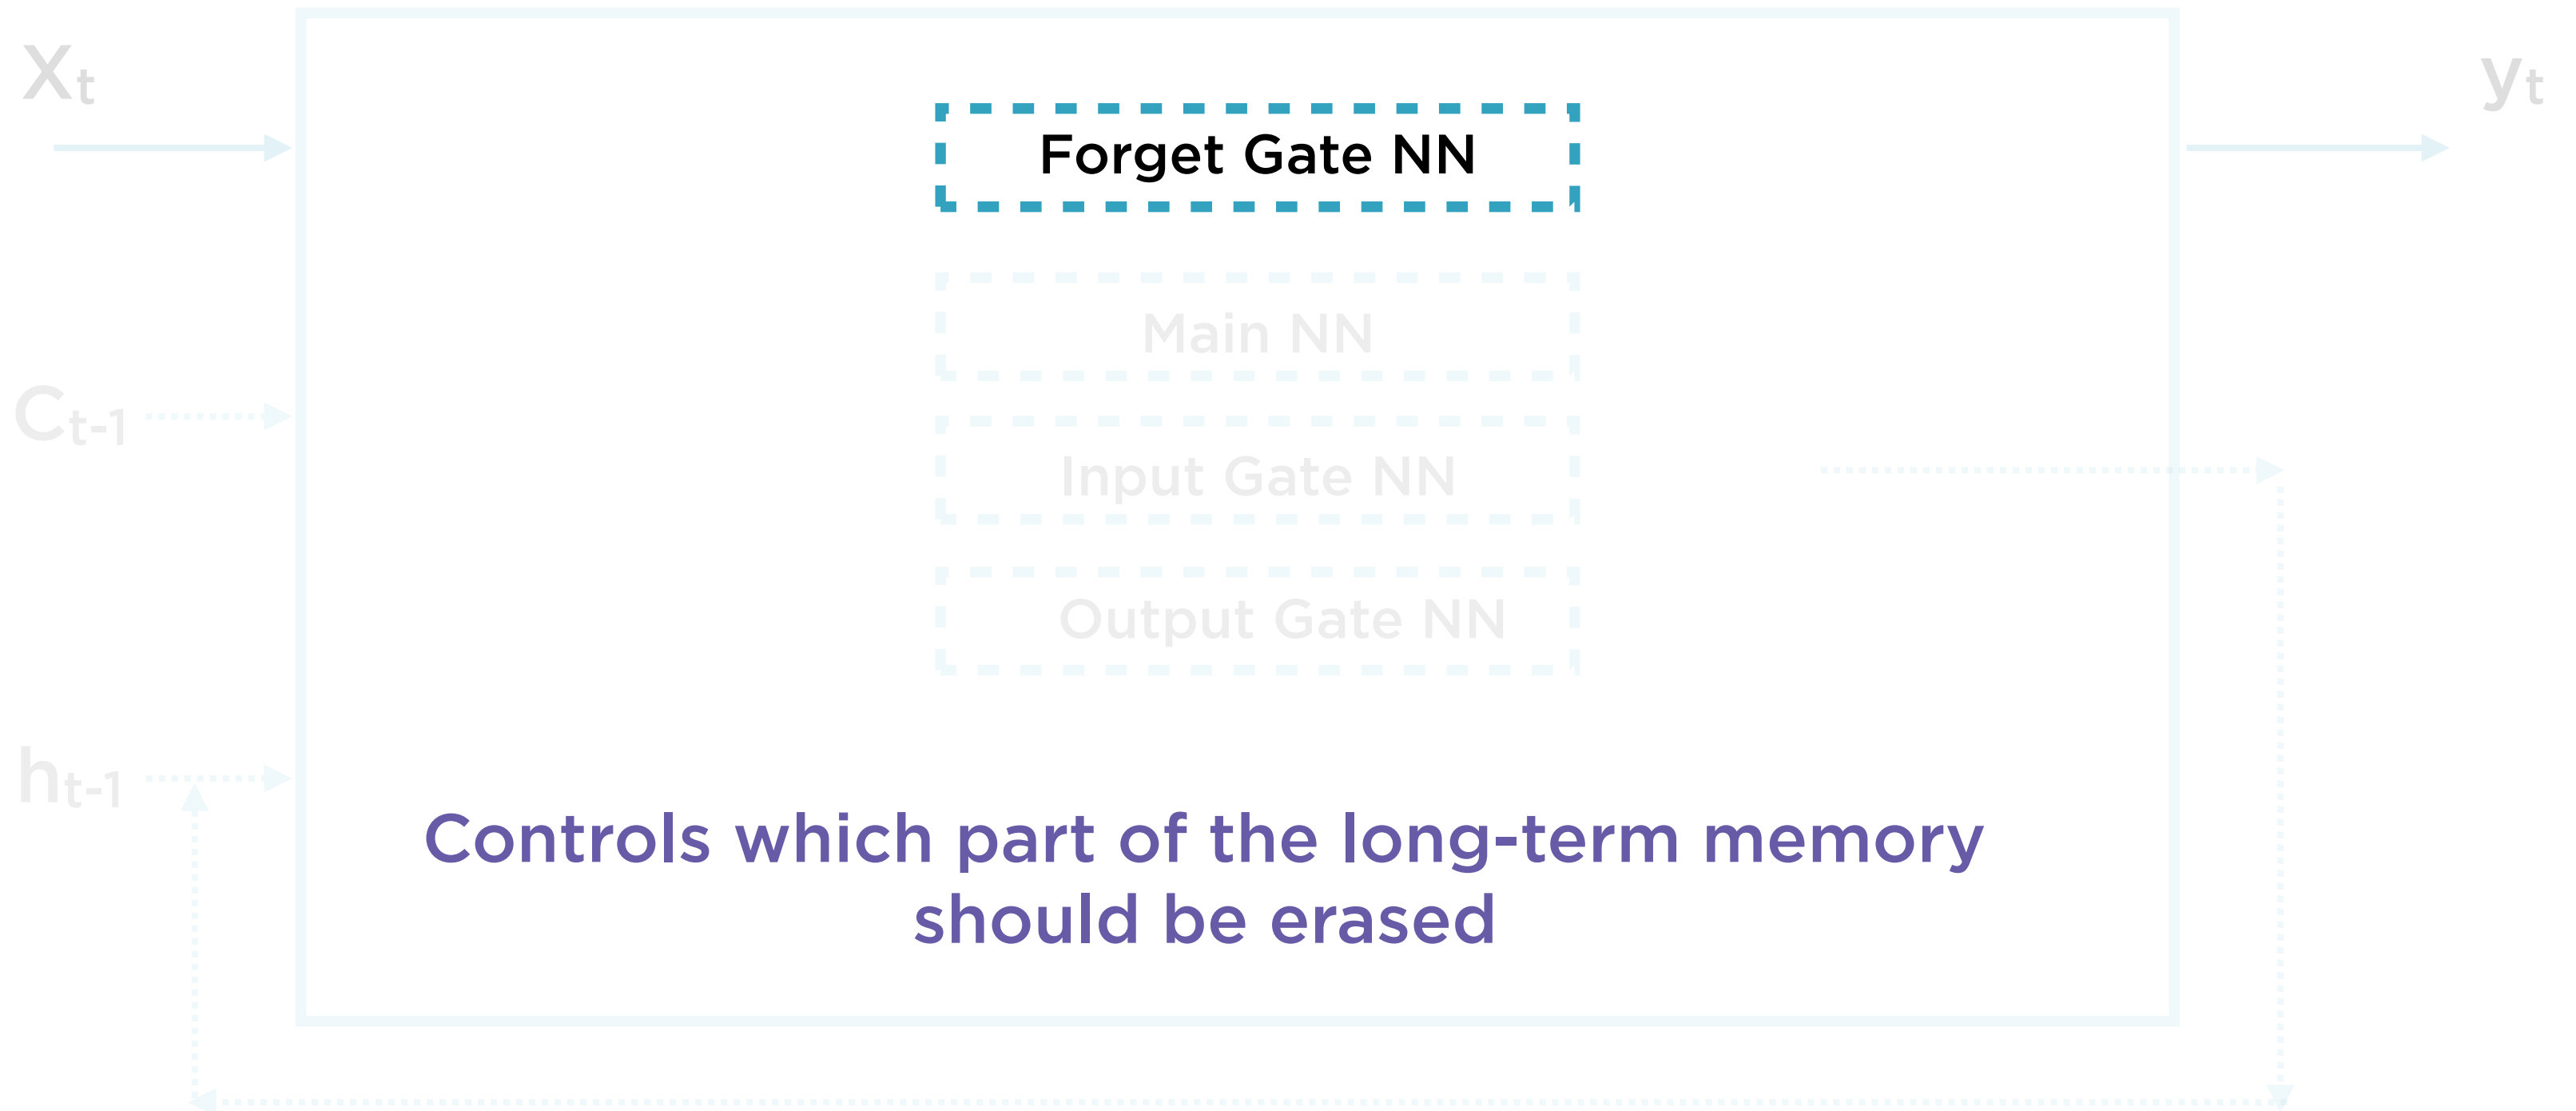
LSTM



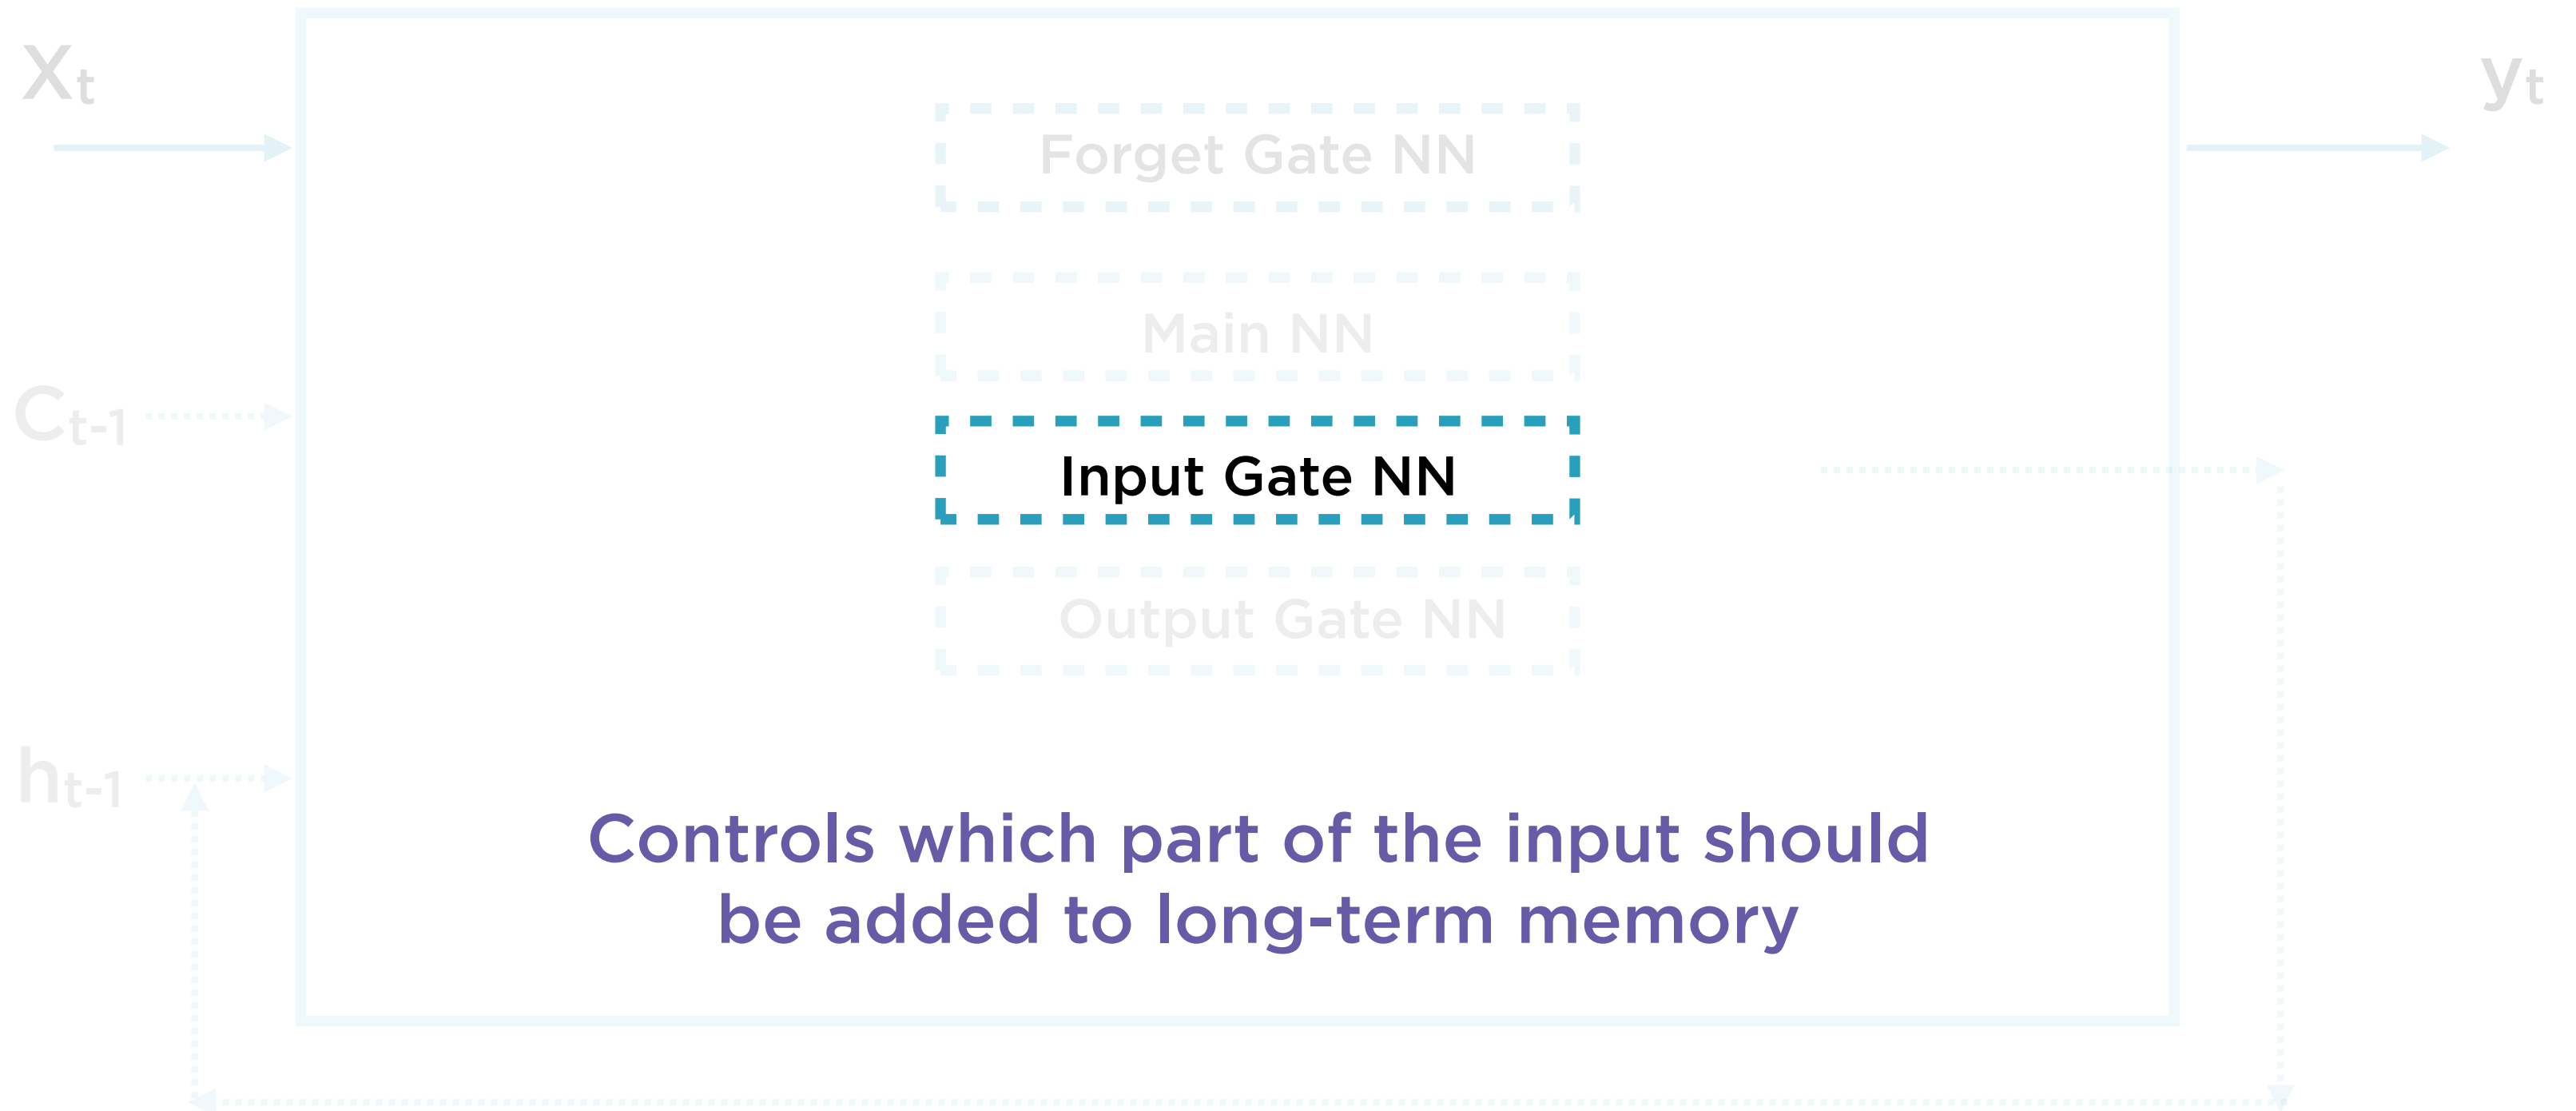
LSTM



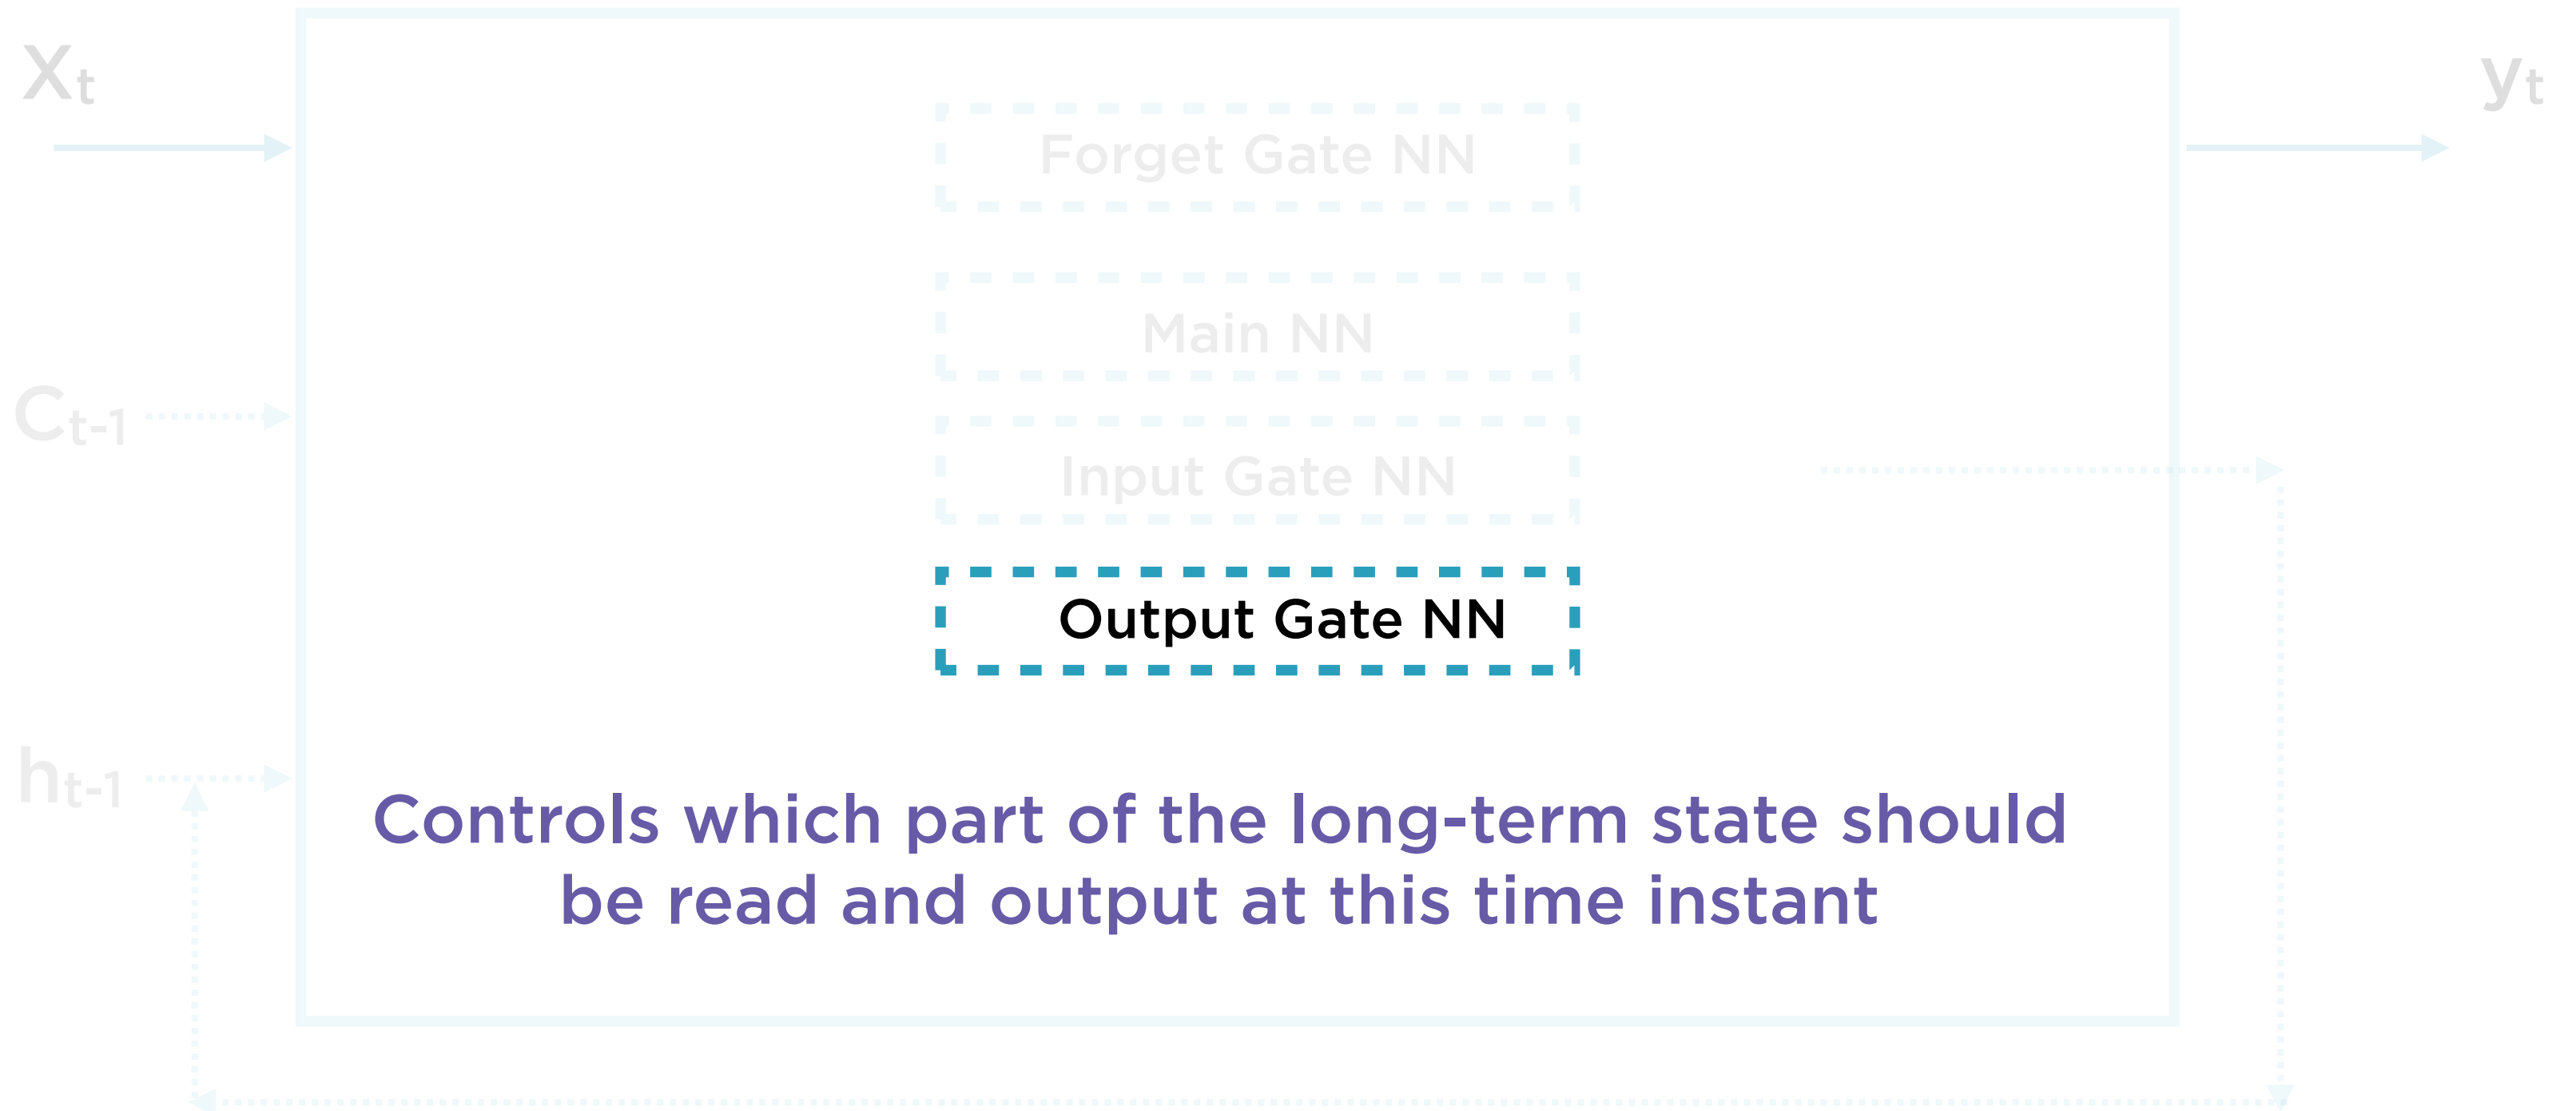
LSTM




LSTM



LSTM



LSTM Cells




Functionally like basic RNN cell

Performance far better

Amazing success at long-term patterns

Long text sequences, time series

Variants



Peephole connections: LSTM cells that store state for more than 1 period

Gated Recurrent Unit (GRU): Simplified LSTM with better performance

- Only 1 state vector
- Fewer internal gates and NNs

Summary

Modifying neurons to endow them with state and memory

Understand Recurrent Neural Networks

Mitigate problems of vanishing and exploding gradients in training RNNs

LSTM and GRU neurons in RNNs

Use RNNs in language modeling