

Performing Sentiment Analysis Using Word Embeddings



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Representing textual data as features for machine learning

Count, frequency, and probability-based embeddings

GloVe and word2vec for pre-trained word embeddings

Performing sentiment analysis using pre-trained word embeddings

Bi-directional RNNs for sentiment analysis

Numeric Representations of Text

`d = "This is not the worst restaurant in the metropolis,
not by a long way"`

Document as Word Sequence

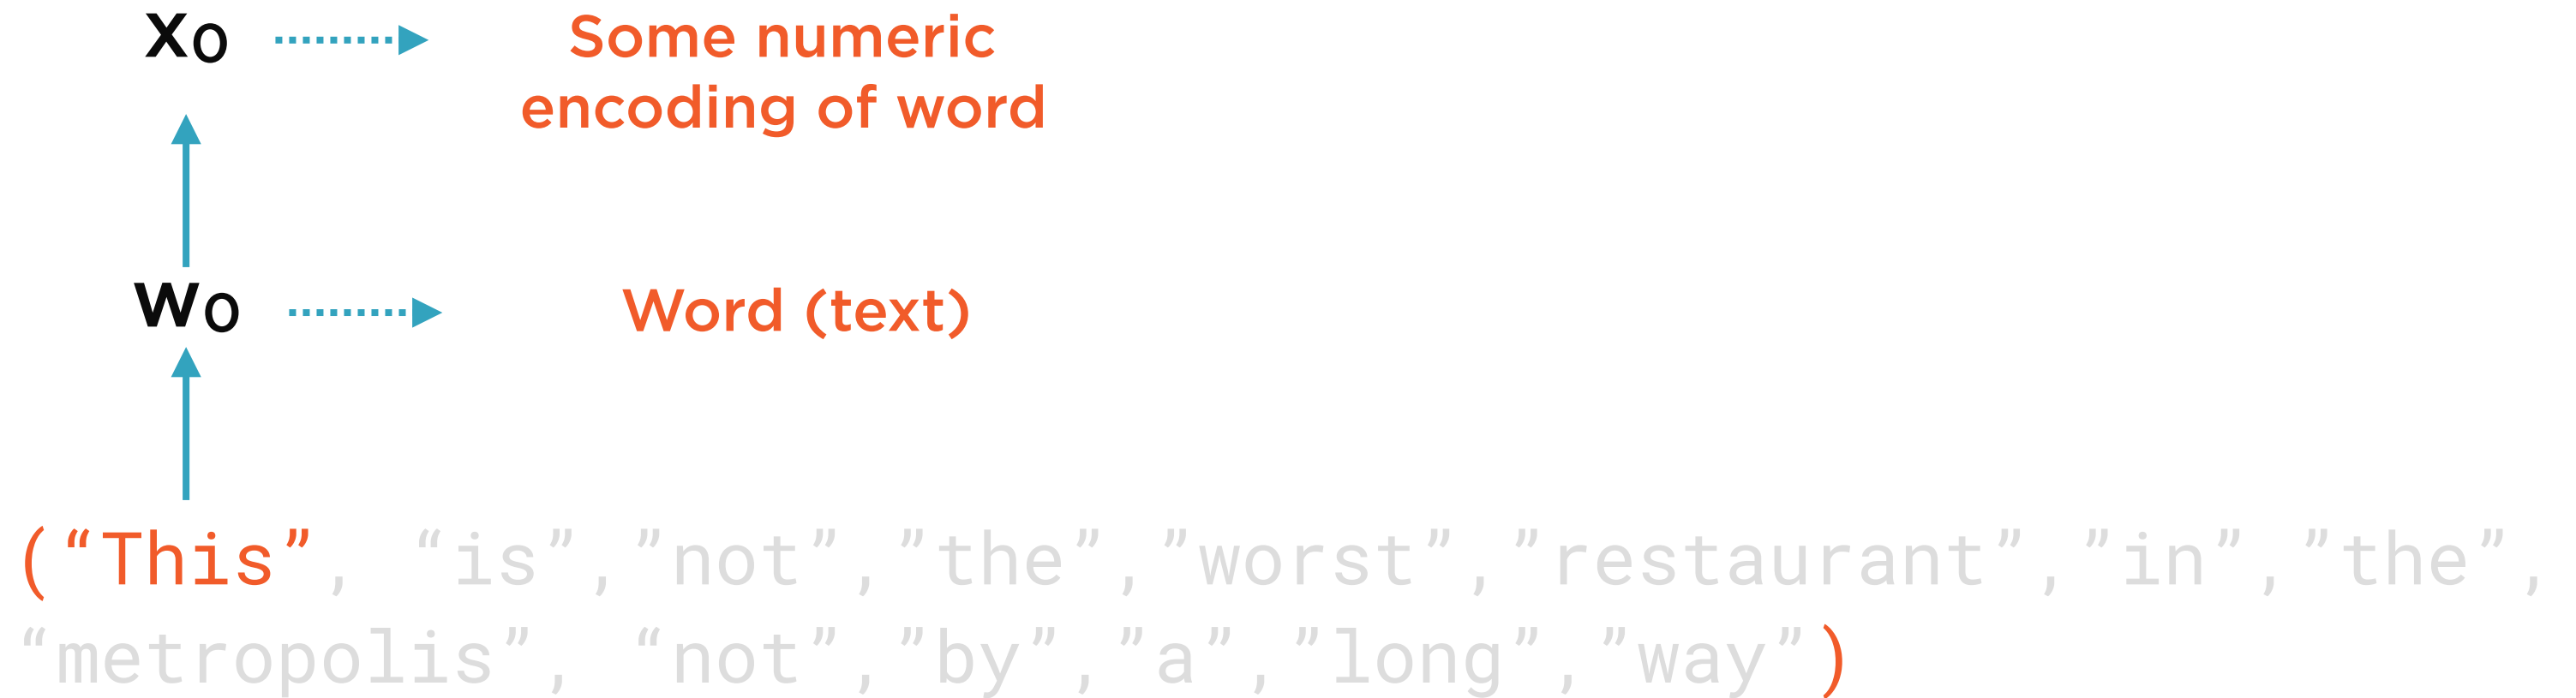
Model a document as an ordered sequence of words

`d = "This is not the worst restaurant in the metropolis,
not by a long way"`

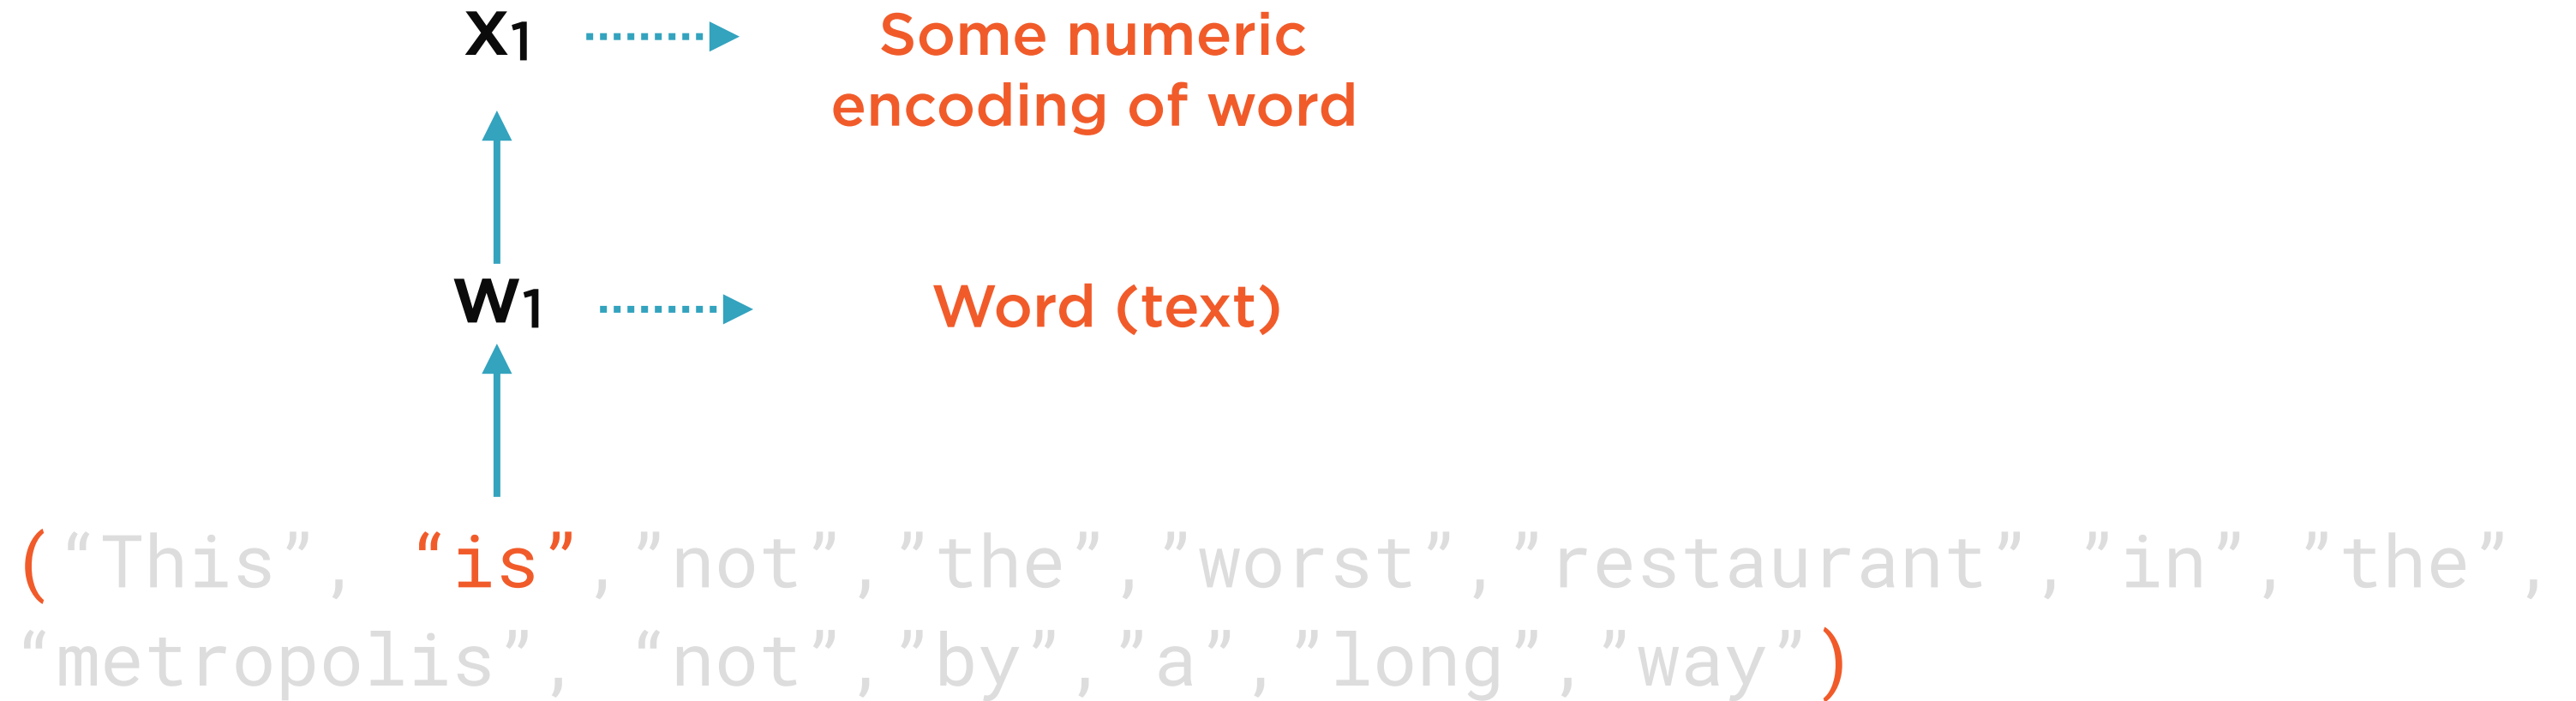
`("This", "is", "not", "the", "worst", "restaurant", "in", "the",
"metropolis", "not", "by", "a", "long", "way")`

Document as Word Sequence

Tokenize document into individual words



Represent Each Word as a Number



Represent Each Word as a Number



Represent Each Word as a Number

$$d = [x_0, x_1, \dots x_n]$$

Document as Tensor

Represent each word as numeric data, aggregate into tensor

Numeric Representations of Text

One-hot

Frequency-based

Prediction-based

Numeric Representations of Text



One-hot

Frequency-based

Prediction-based

Represent each word in text by its
presence or absence

Numeric Representations of Text

One-hot

Frequency-based

Prediction-based

Frequency-based Embeddings

Count

TF-IDF

Co-occurrence

Frequency-based Embeddings

Count

TF-IDF

Co-occurrence

Capture how often a word occurs in a document i.e. the **counts** or the **frequency**

Frequency-based Embeddings

Count

TF-IDF

Co-occurrence

Captures how often a word
occurs in a **document** as well as
the **entire corpus**

Tf-Idf



Frequently in a single document

Might be important



Frequently in the corpus

**Probably a common word like
“a”, “an”, “the”**

Frequency-based Embeddings

Count

TF-IDF

Co-occurrence

Similar words will occur
together and will have similar
context

Context Window

A window centered around a word, which includes a certain number of neighboring words

Co-occurrence

The number of times two words w_1 and w_2 have occurred together in a context window

Word Embeddings

One-hot

Frequency-based

Prediction-based

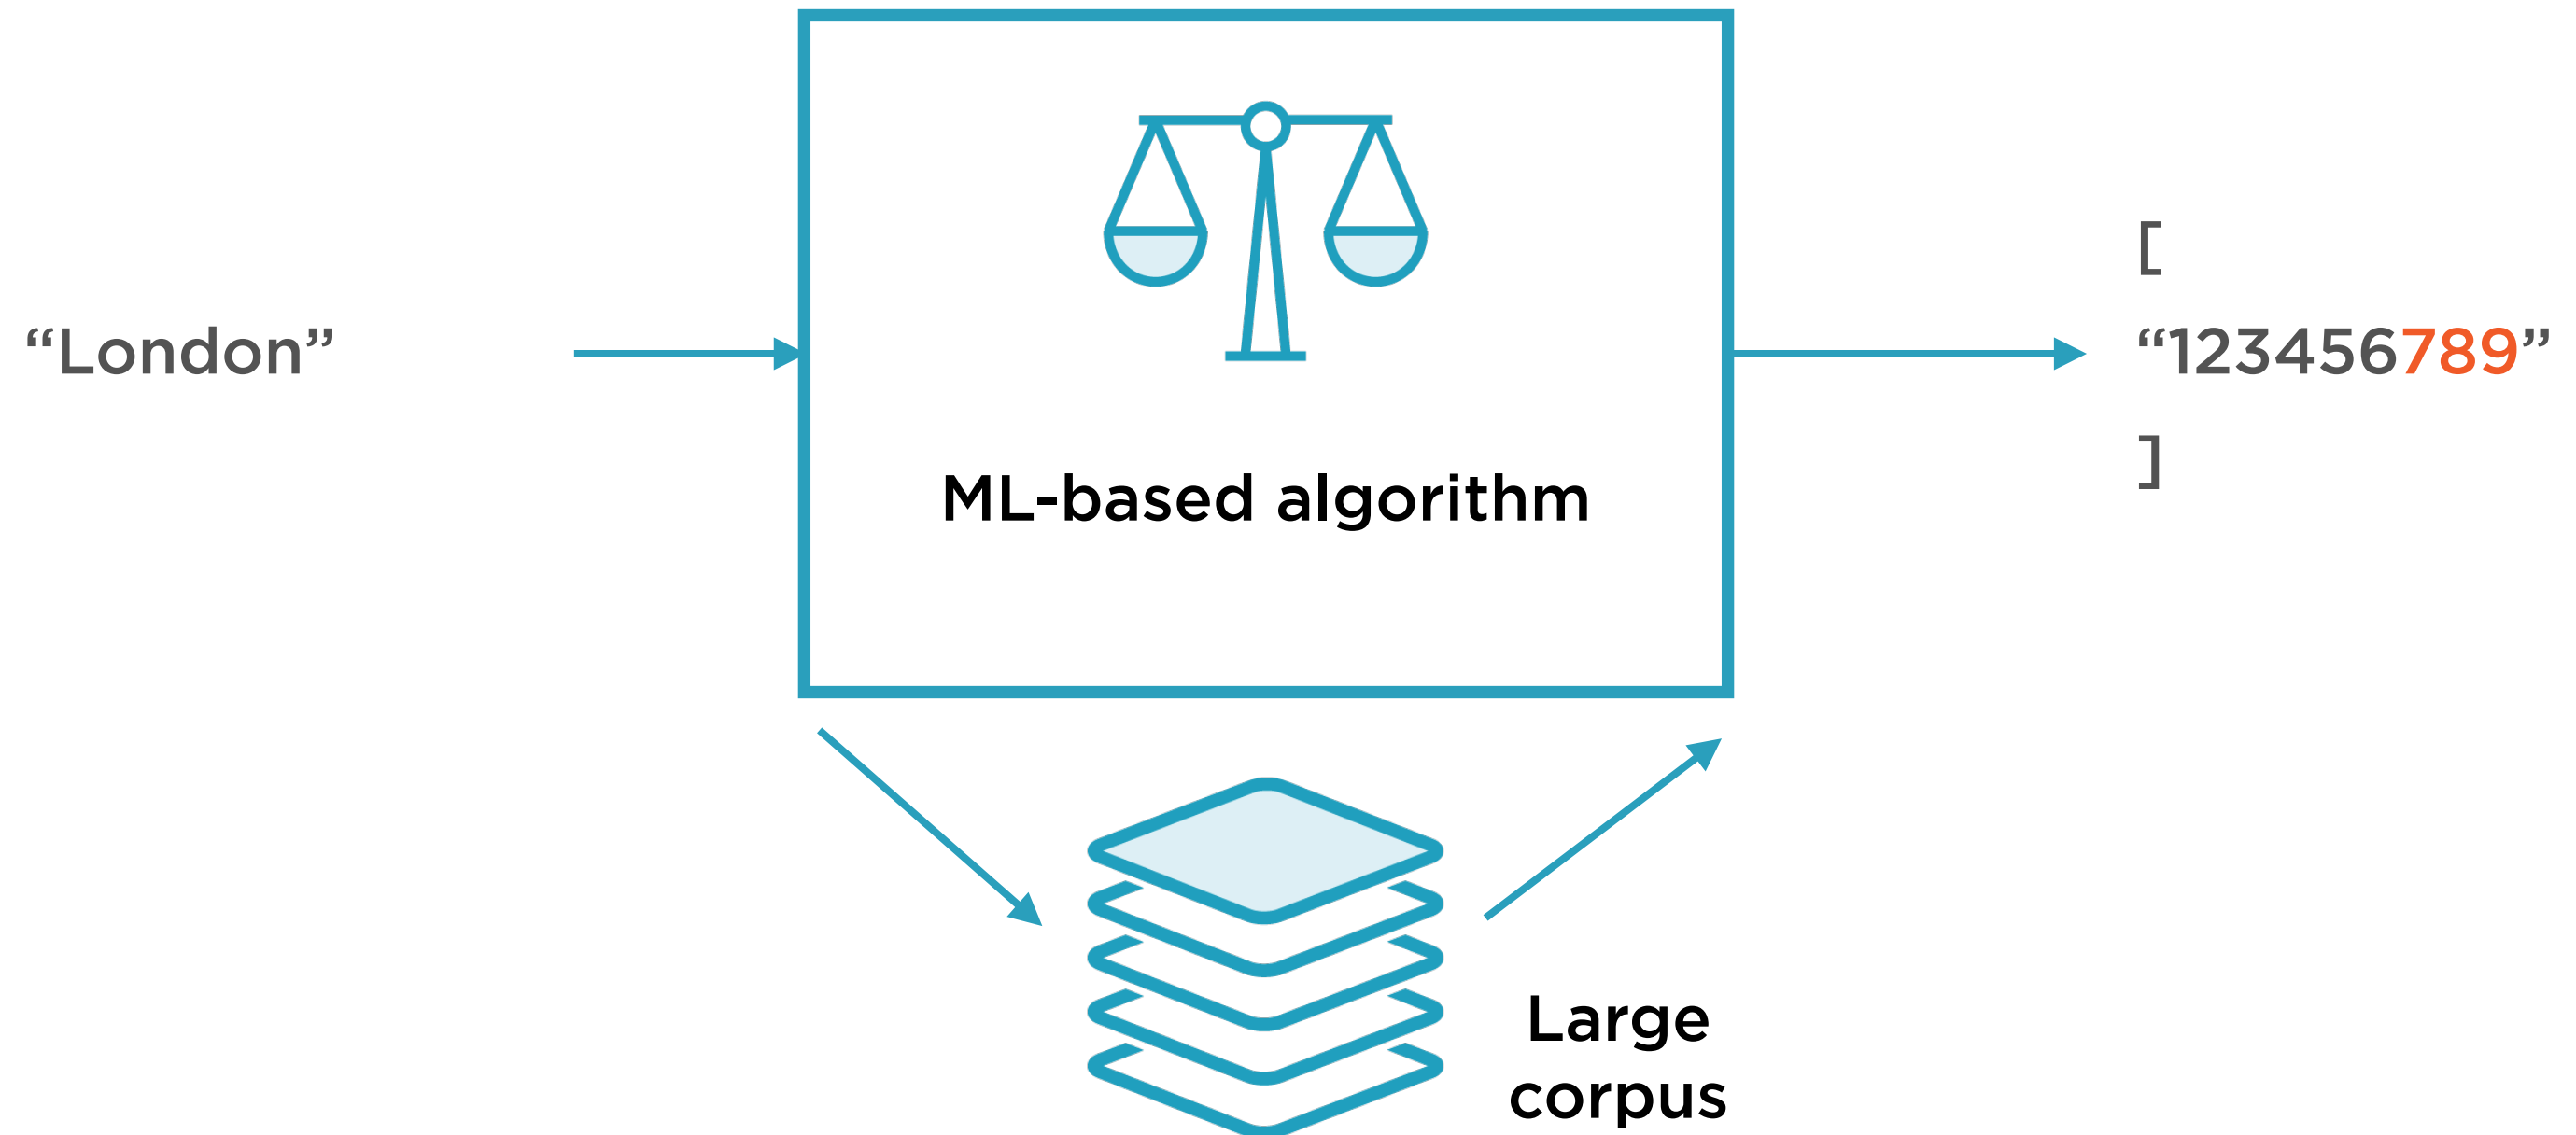


Predictions-based embeddings

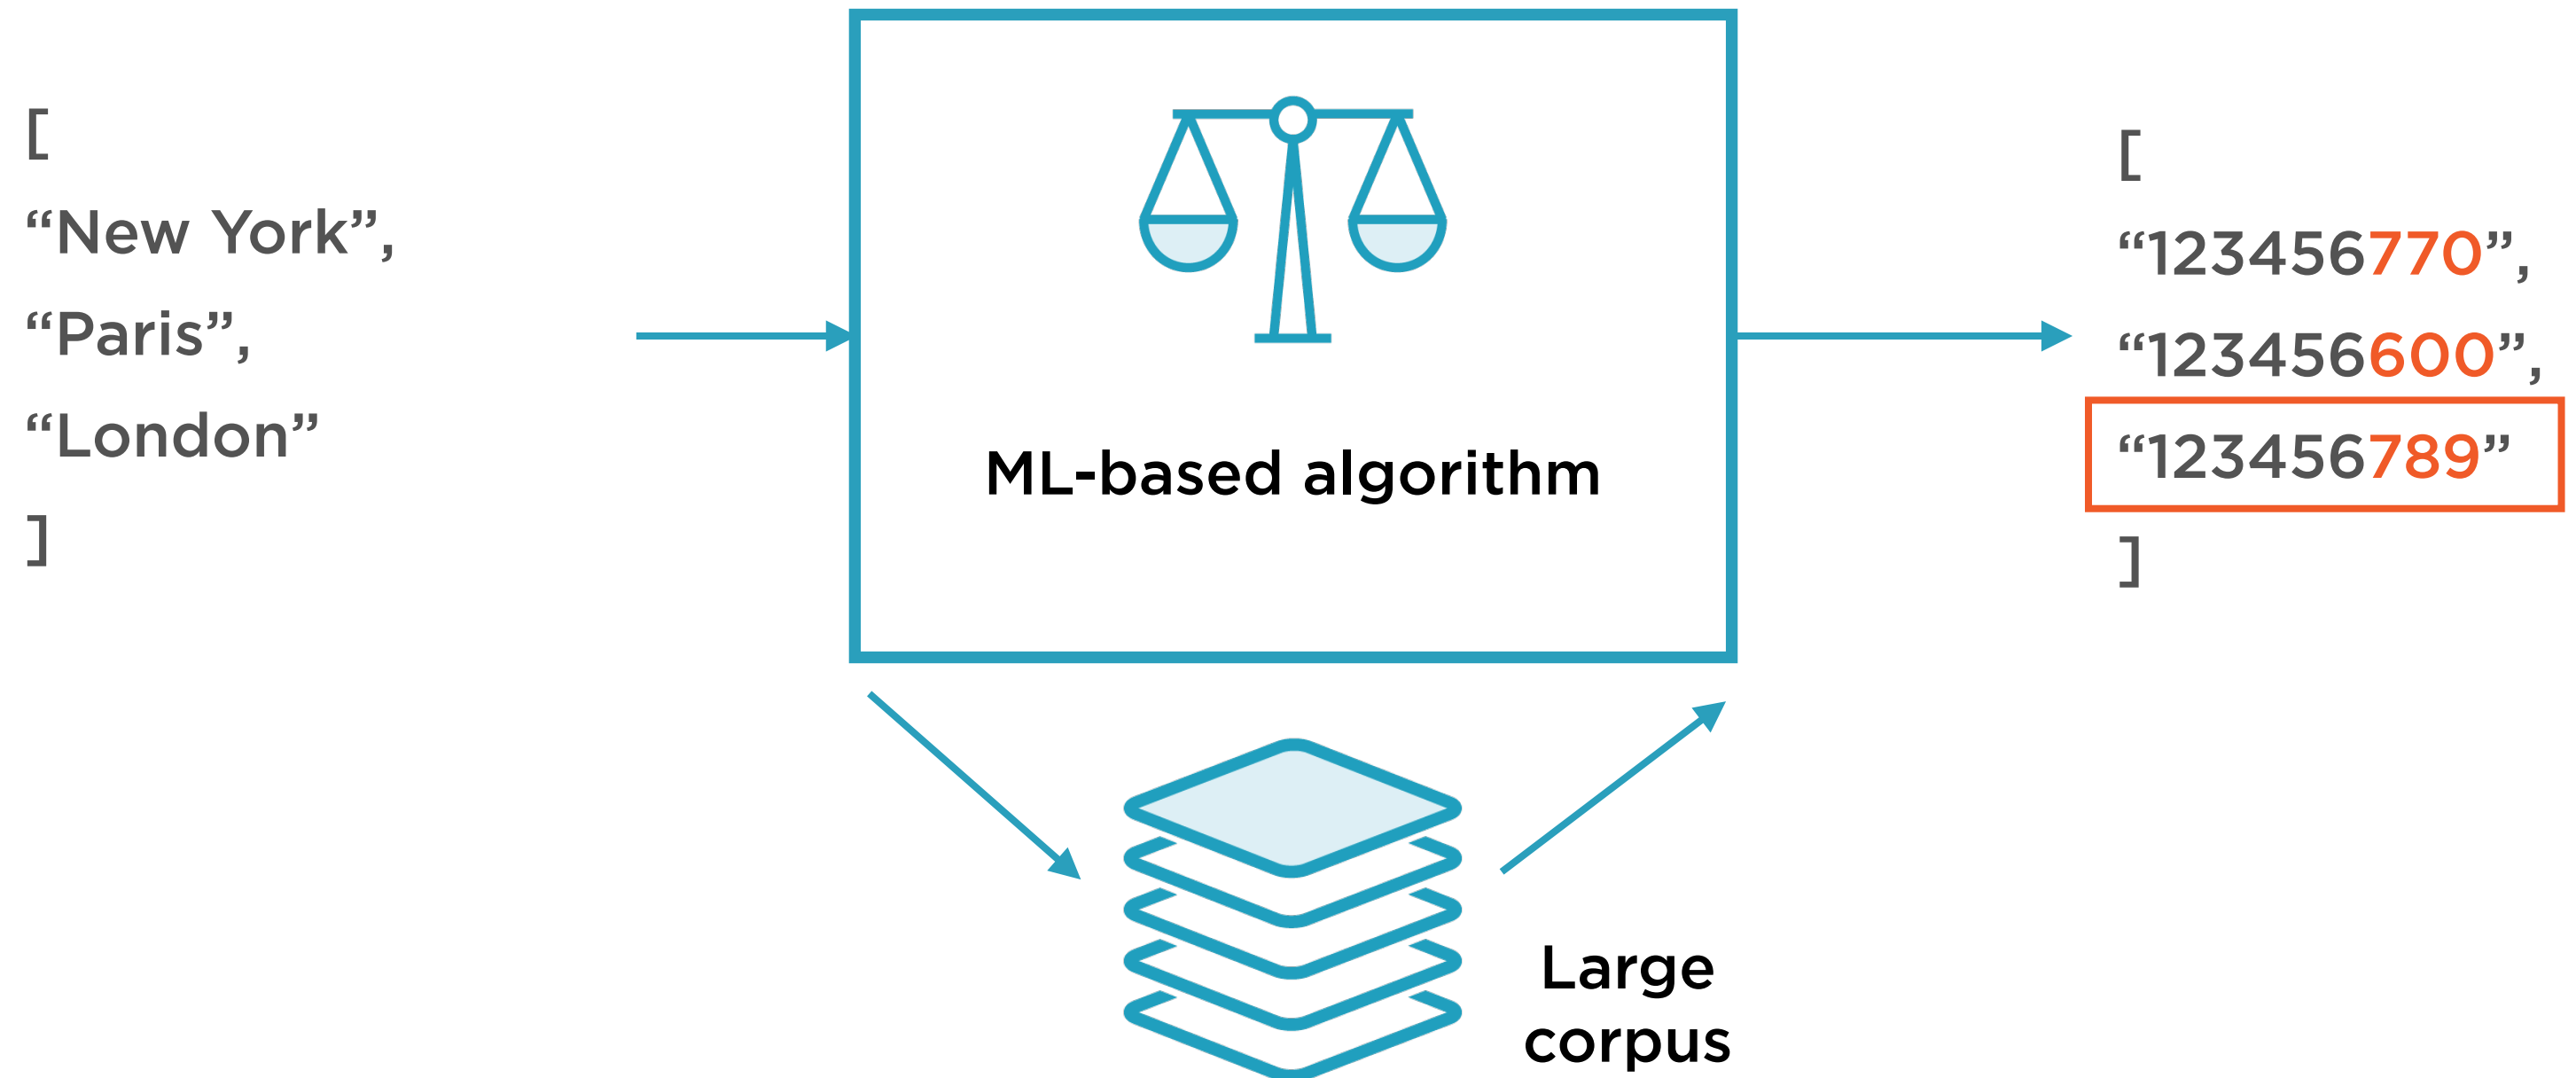
Capture meanings and semantic
relationships, generated using ML
models

~~“Birds~~ Words of a feather flock
together”

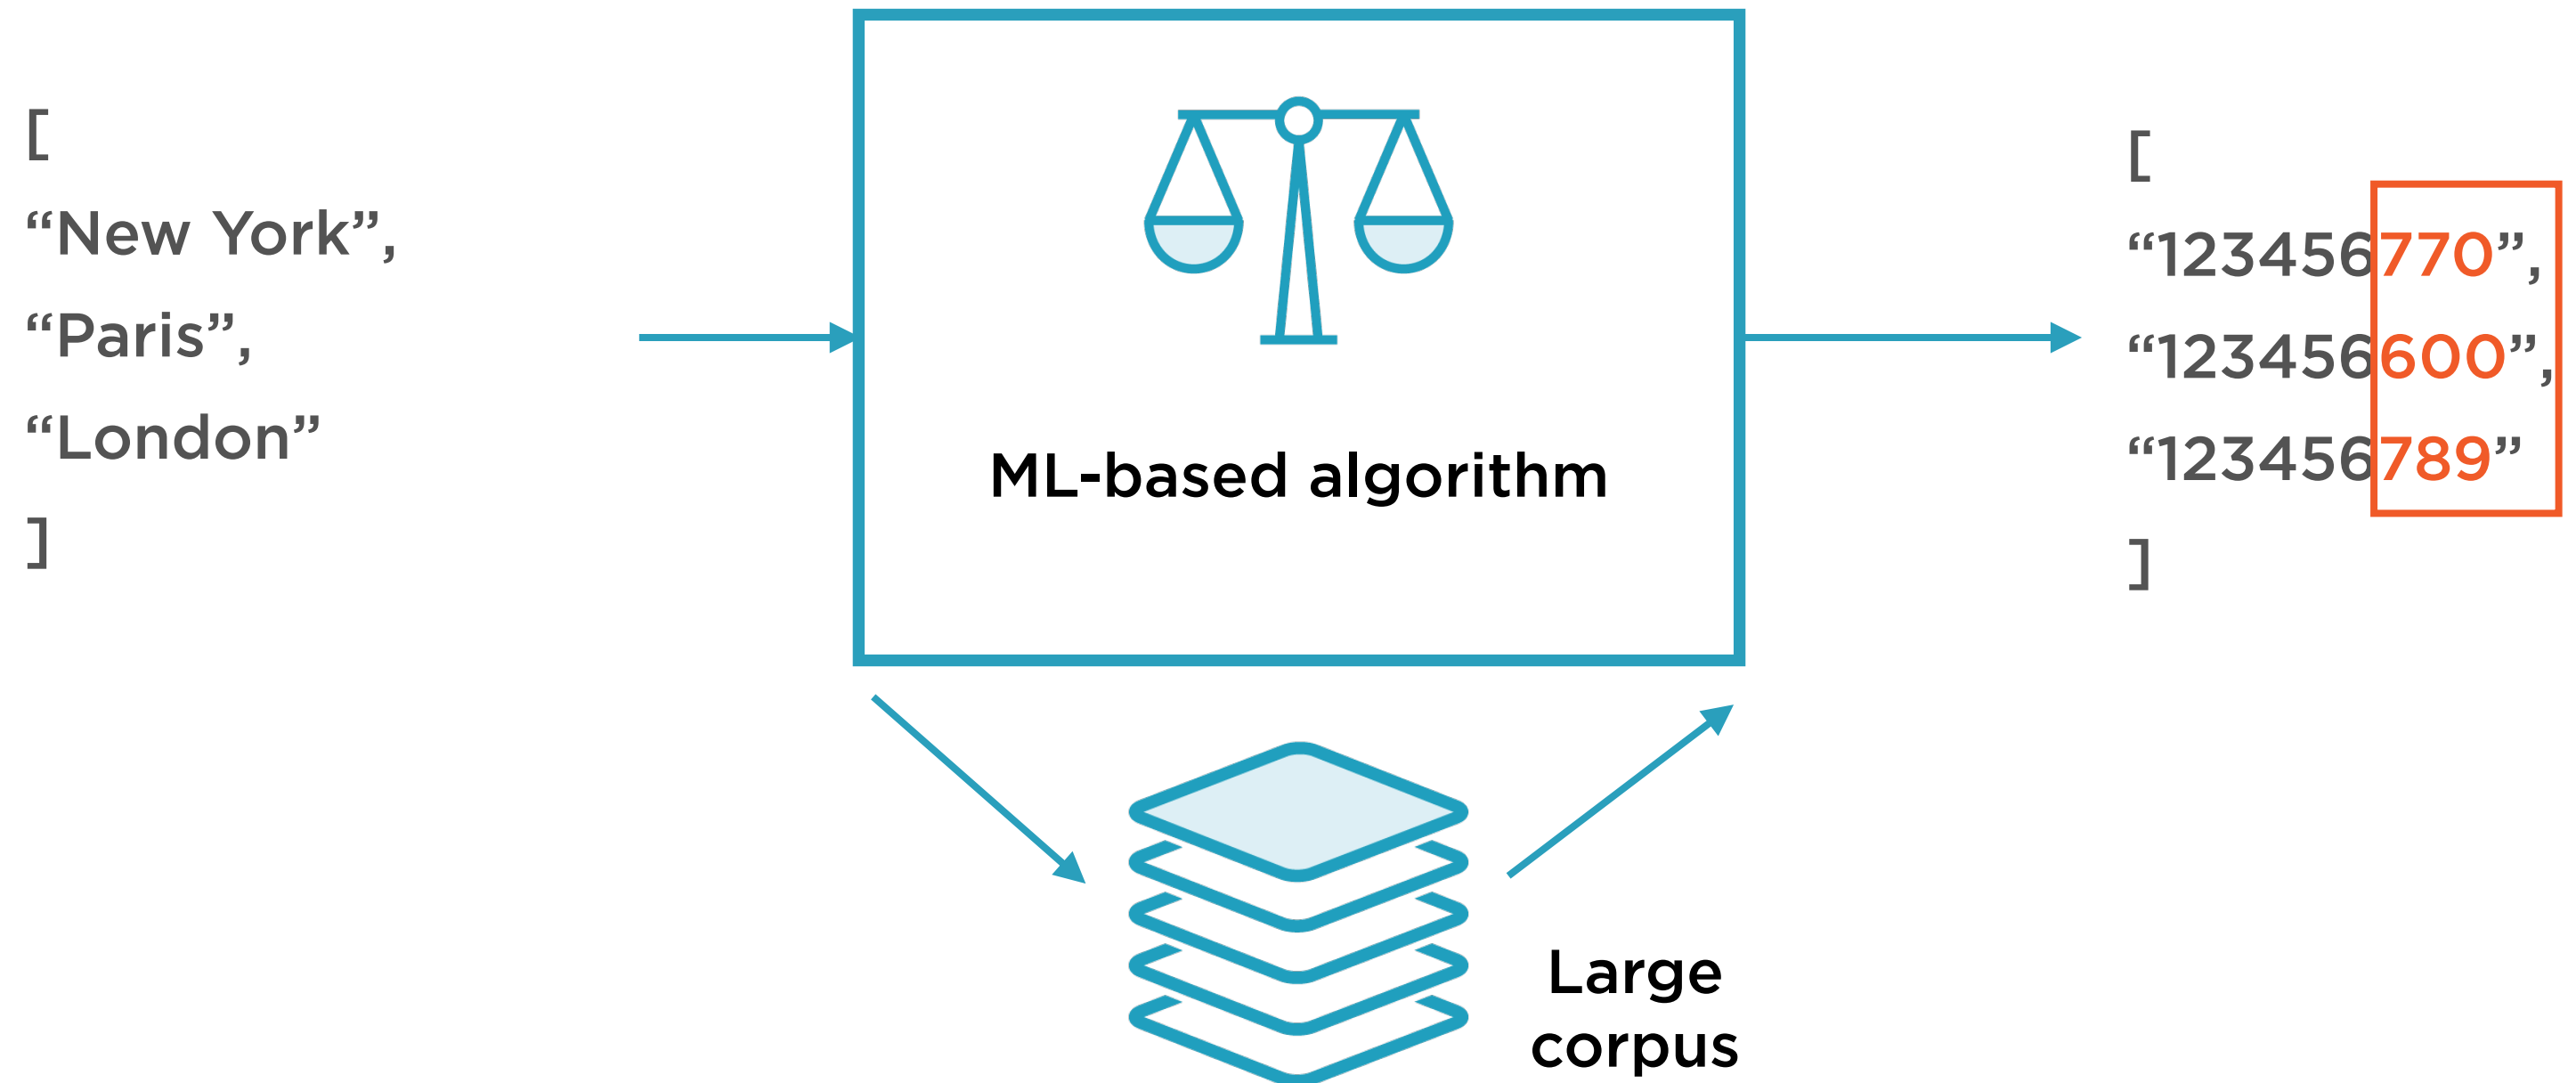
Prediction-based Word Embeddings




Prediction-based Word Embeddings



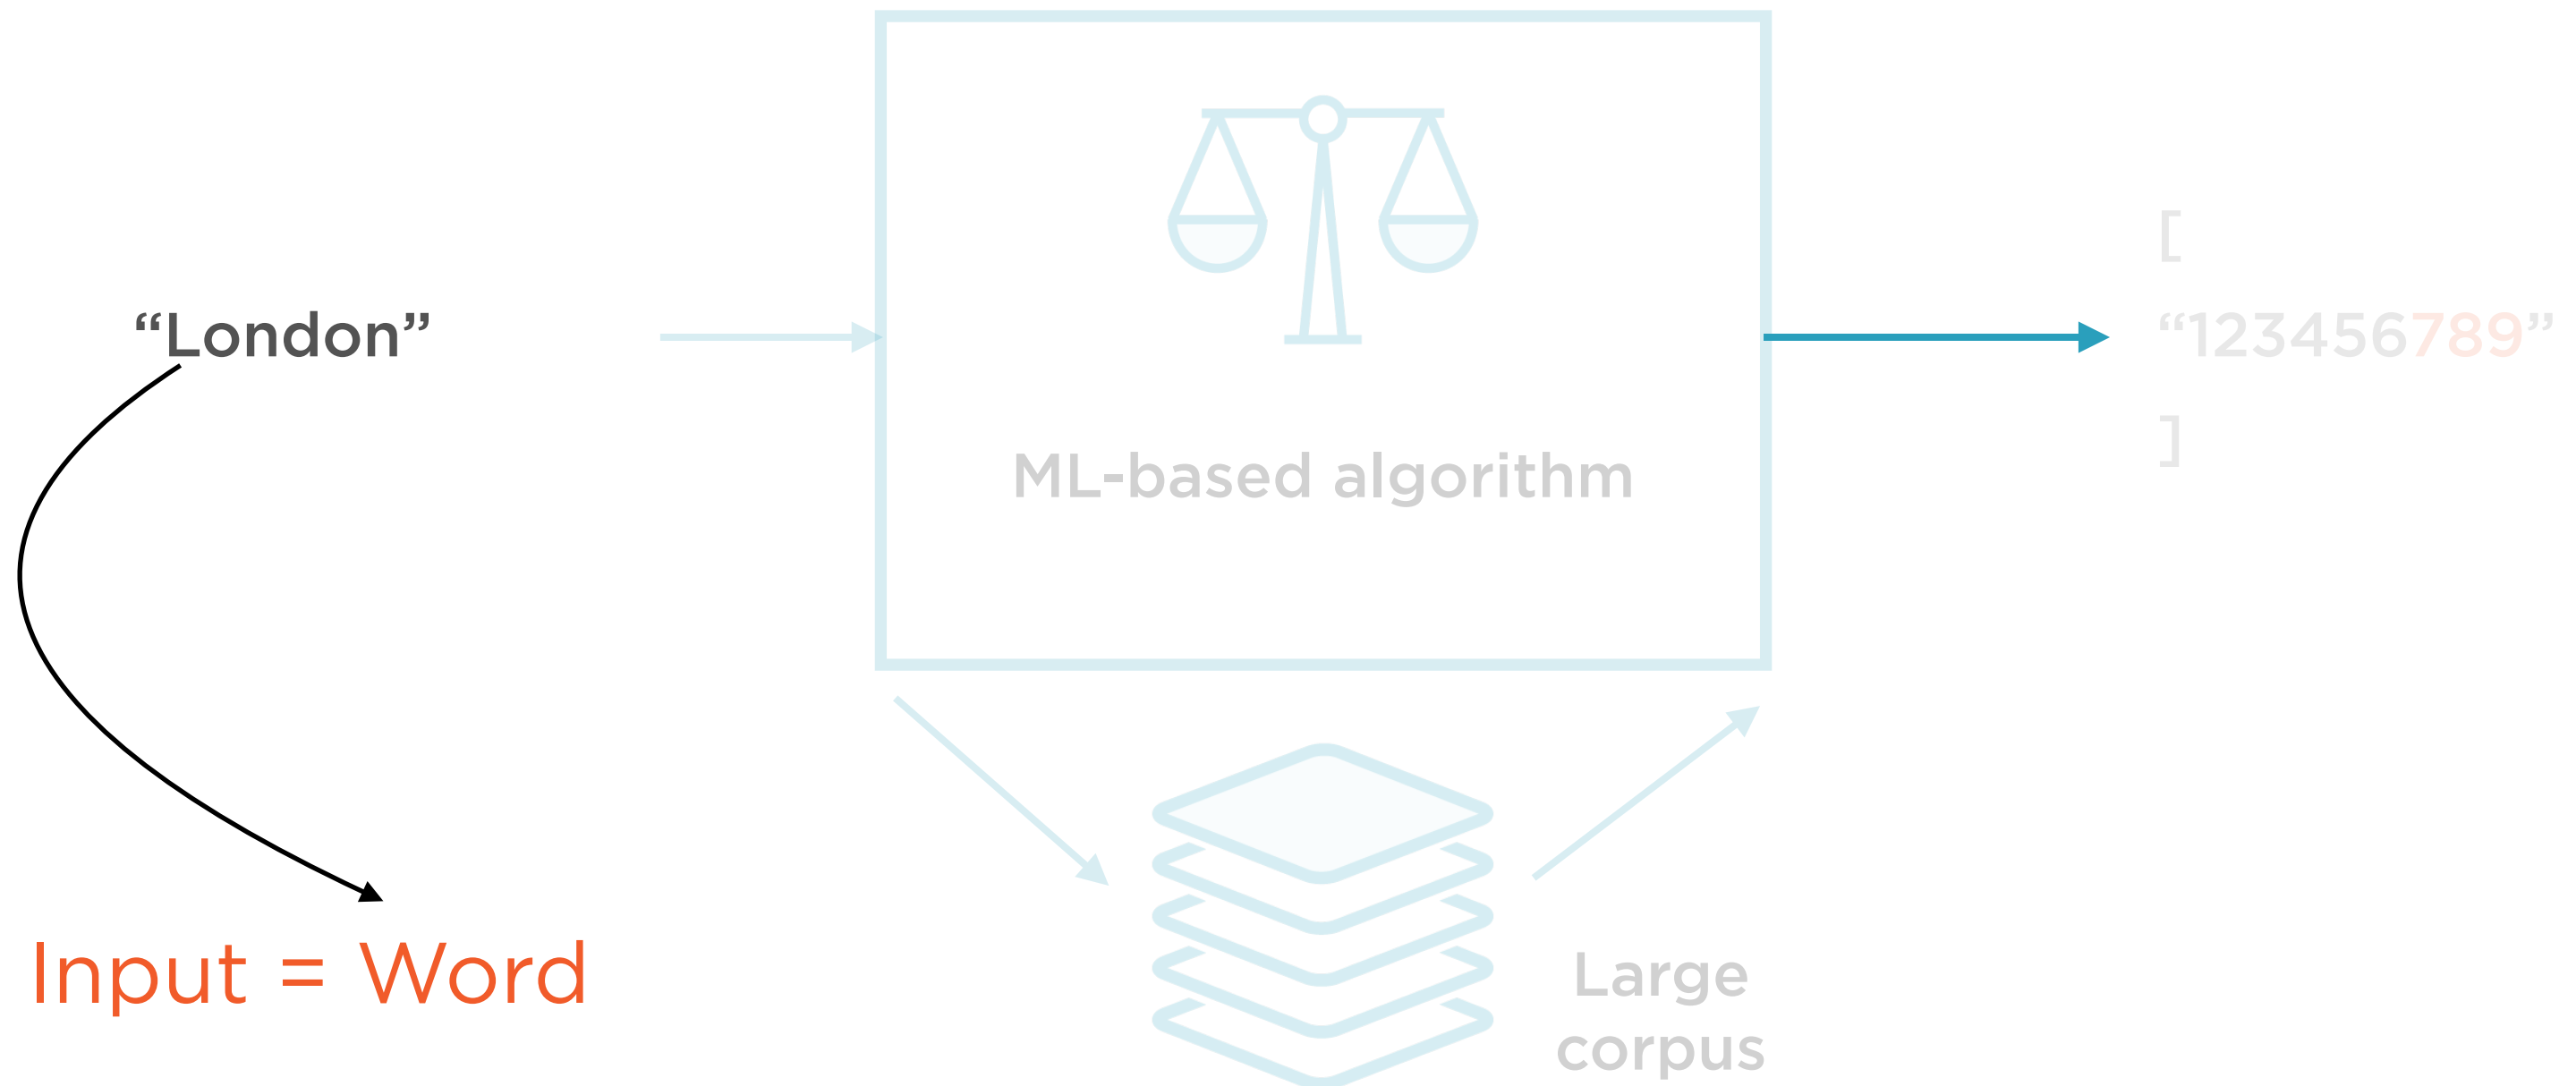
Prediction-based Word Embeddings




Prediction-based Word Embeddings



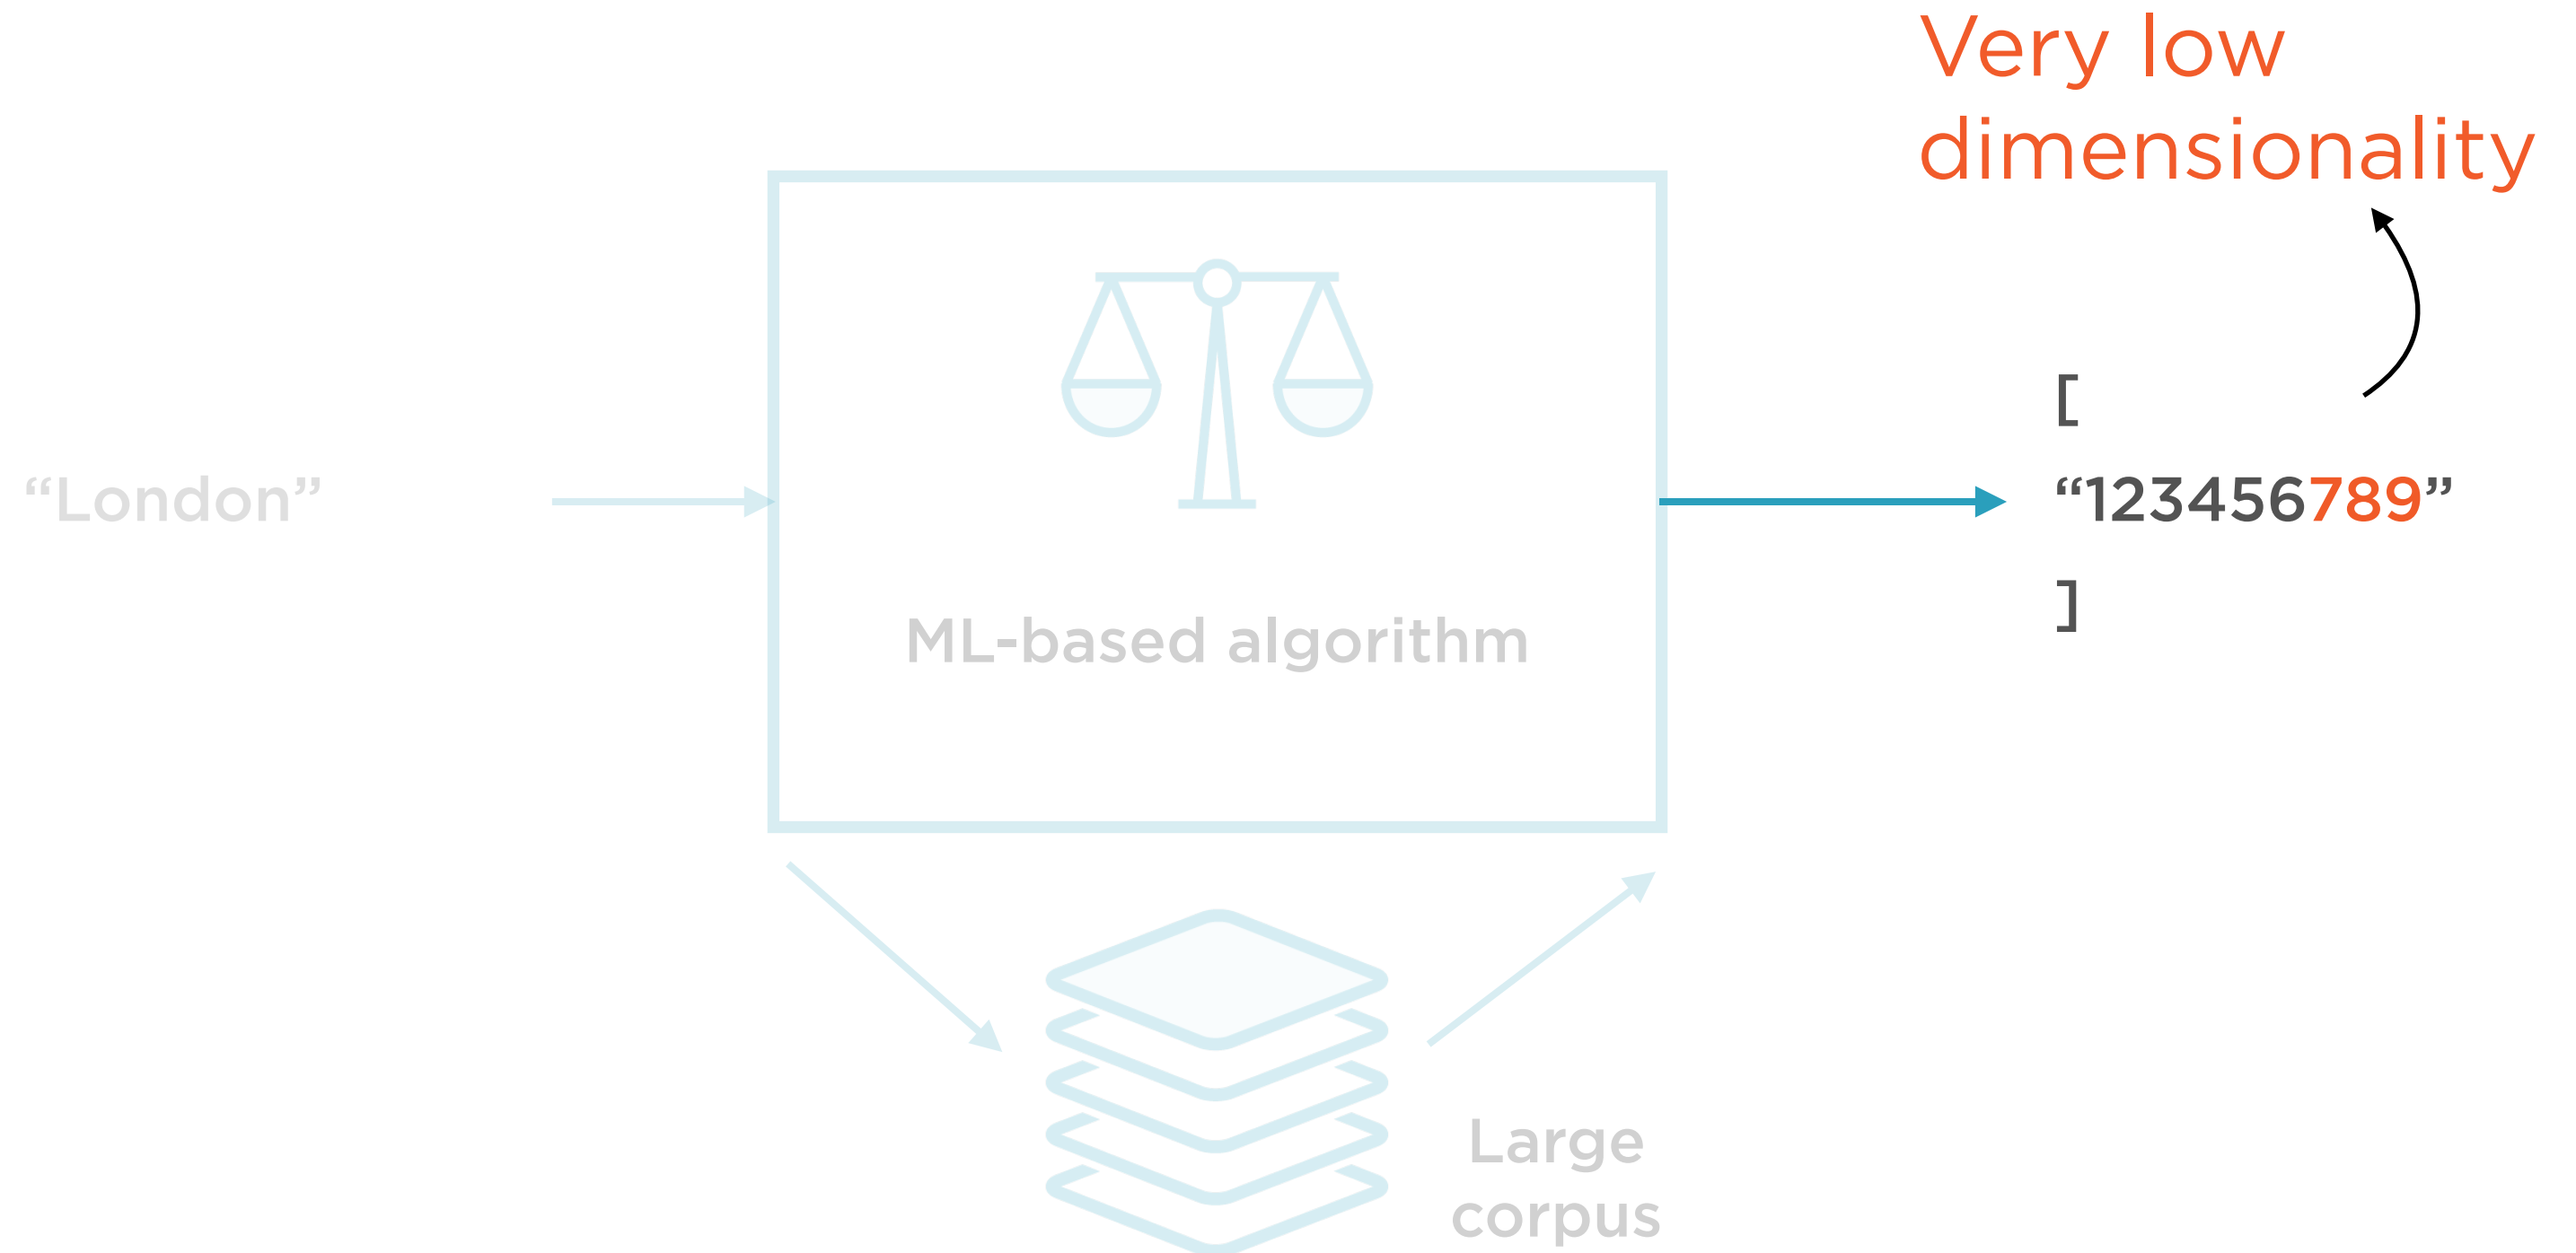
Prediction-based Word Embeddings



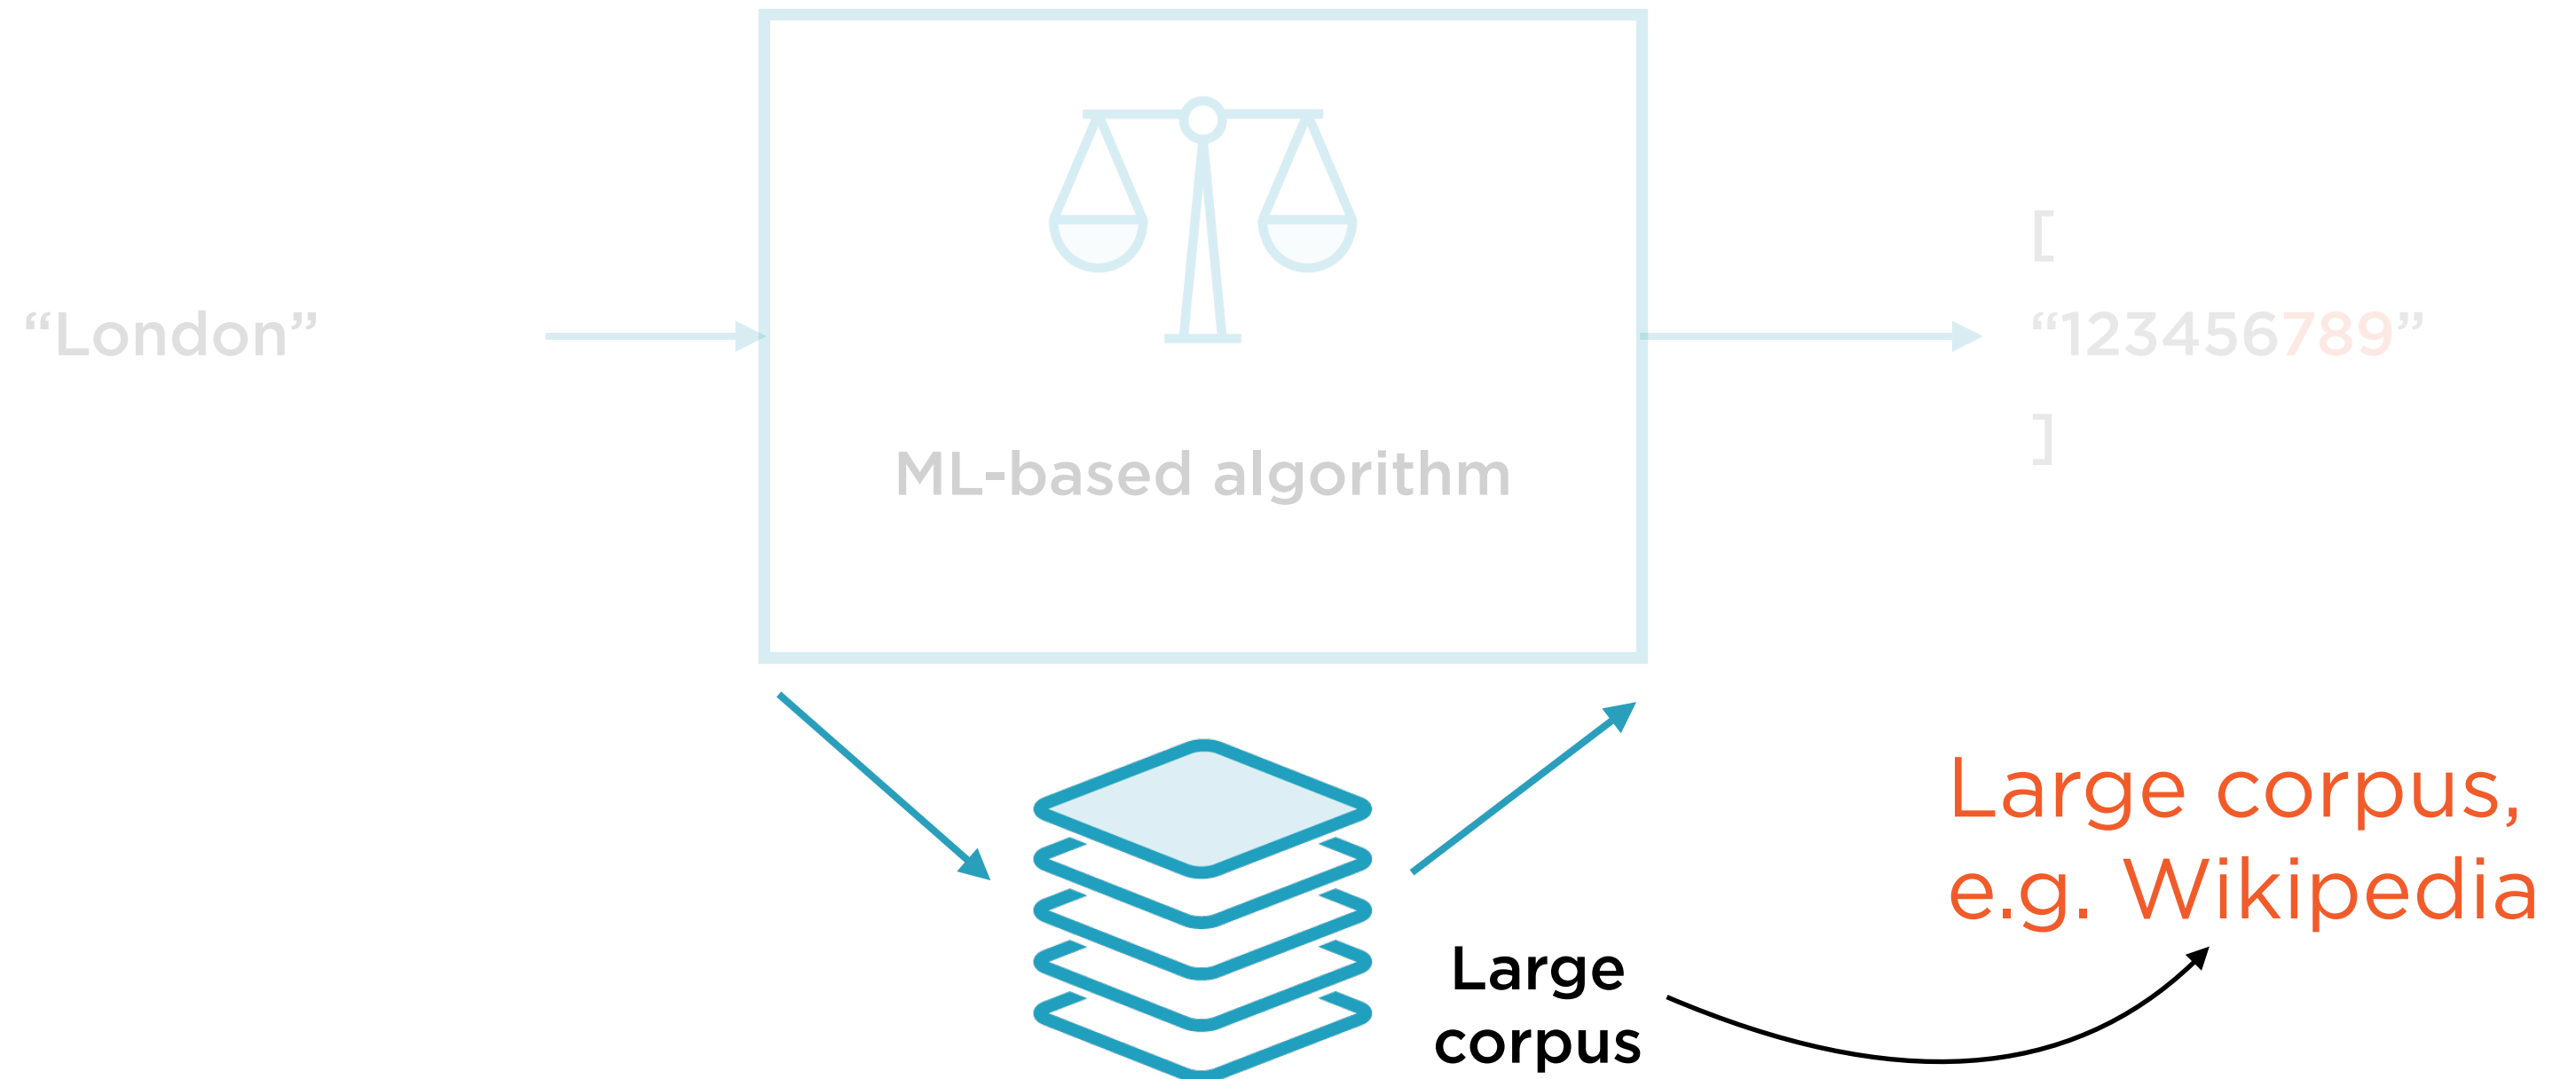
Prediction-based Word Embeddings



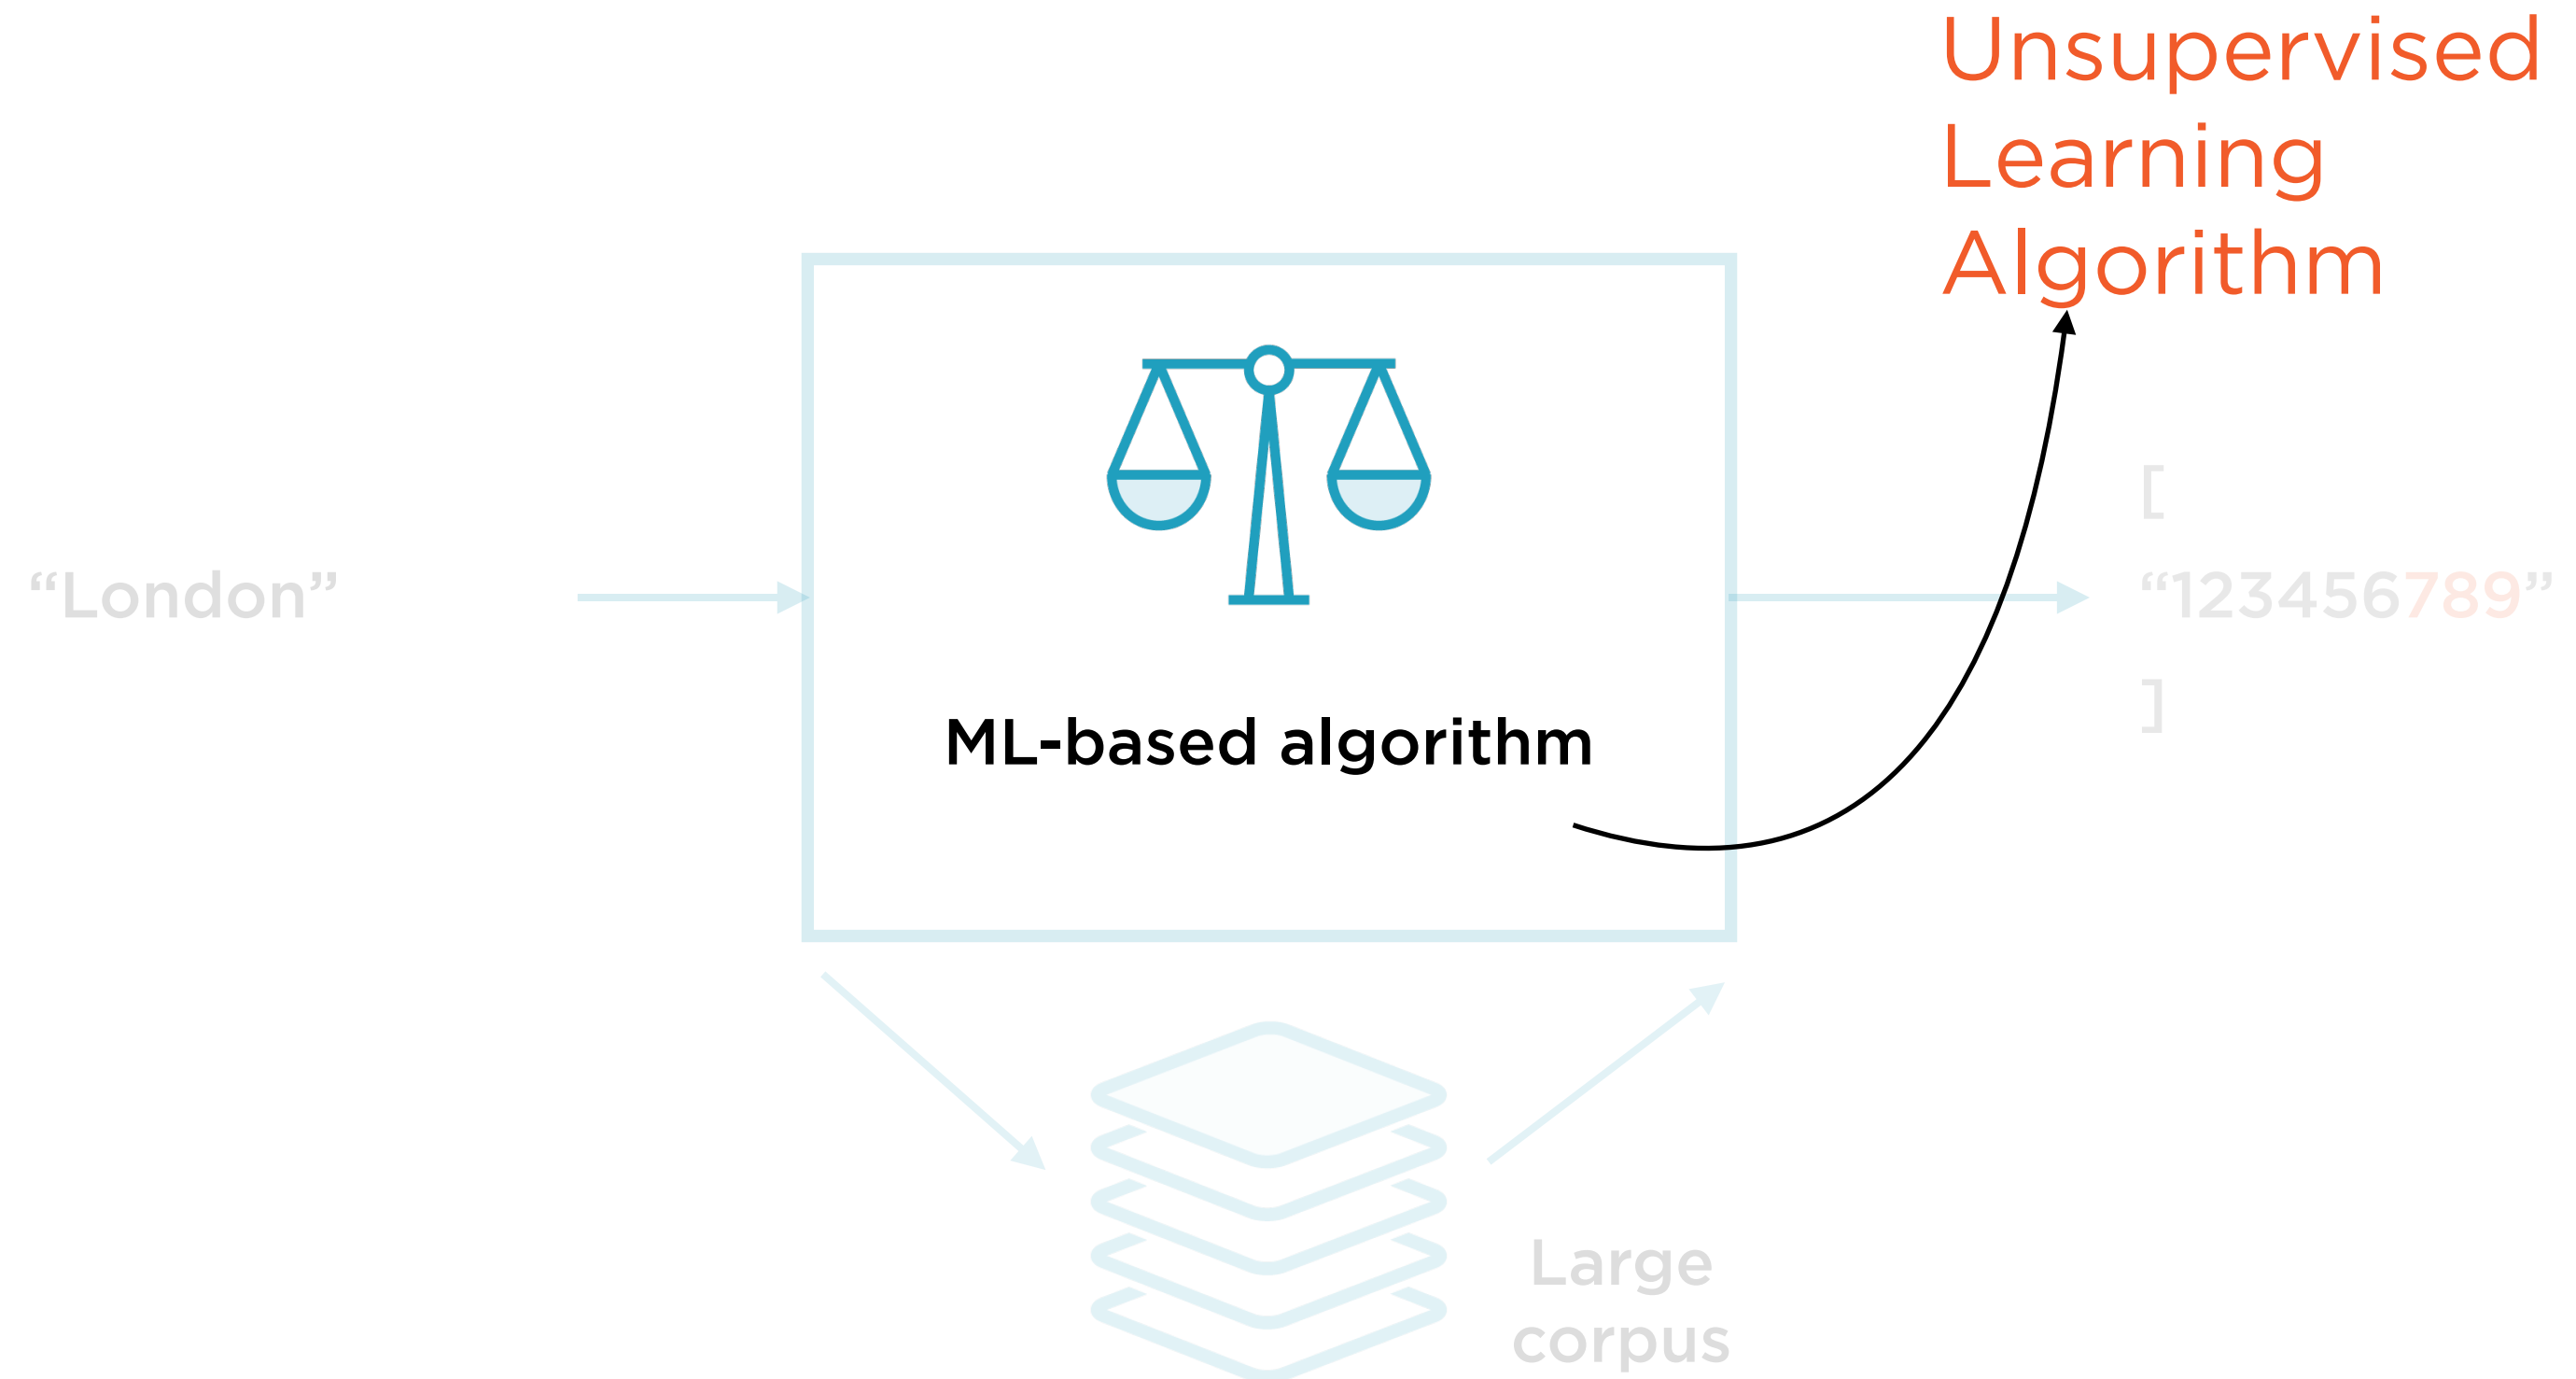
Prediction-based Word Embeddings



Prediction-based Word Embeddings



Prediction-based Word Embeddings



Pre-trained Word Embeddings

Word2Vec

GLoVe

Word2Vec



Most popular word embedding model

Mikolov (Google), 2013

Use simple NN (not deep) to learn embeddings

GLoVe



Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher,
Christopher D. Manning, (Stanford) 2014

Uses word-word co-occurrence matrix,
nearest-neighbors for word relationships

Given words from its context,
predict the word

Given a word, **predict the
words in its context**

Word Embeddings



Encode each word as a vector of other words

What words? A few related ones

Embedding("true") = ["false"]

Embedding("London") = ["New York", "Paris"]

Magic



Word embeddings capture meaning

“Queen” ~ “King” == “Woman” ~ “Man”

“Paris” ~ “France” == “London” ~ “England”

Dramatic dimensionality reduction


Embeddings are a way to
encode words capturing the
context around them

Demo

Explore pre-trained GloVe word embeddings

Multilayer RNNs

Conventional RNNs




The cell unrolled through time form the layers of the
neural network


Multilayer RNNs

Add additional RNNs on top of the original RNN,
where each RNN added is another layer


Multilayer RNNs




Multilayer RNNs




Multilayer RNNs




Multilayer RNNs



Multilayer RNNs



Multilayer RNNs



Multi-layer RNNs (also called deep RNNs) are another simple concept. The idea is that we add additional RNNs on top of the initial standard RNN, where each RNN added is another layer. The hidden state output by the first (bottom) RNN at time-step t will be the input to the RNN above it at time step t . The prediction is then made from the final hidden state of the final (highest) layer.

The image below shows a multi-layer unidirectional RNN, where the layer number is given as a superscript. Also note that each layer needs their own initial hidden state, h_0^l .


Bidirectional RNNs

`d = "The movie was terrible. Hated it."`

Document as Word Sequence

Model a document as an ordered sequence of words

Conventional RNN



Conventional RNN


The sky is cloudy, it looks like. _____



The blank slot will be filled in using
preceding information only

Bidirectional RNN

The sky is _____, it looks like rain



The diagram shows a horizontal orange arrow pointing from right to left, positioned above the word 'rain' in the sentence. This represents the backward pass of a bidirectional recurrent neural network, where information is processed from the end of the sentence towards the beginning.

Bidirectional RNNs will also use following
information

`d = ".it Hated .terrible was movie The"`

Reverse the Document Word Sequence

Bidirectional RNNs run words both forward and backward


Is the use of a Bidirectional RNN
cheating?

For streaming applications - yes; for
batch applications - **no!**

Bidirectional RNNs

Outputs

y_t




Forward
Layer

Backward
Layer

Inputs

x_t

Bidirectional RNNs




Speech recognition: Utterance = batch of words spoken together

Text auto-complete, auto-correct: Fragment = batch of characters input so far


Language Modeling: e.g. Cause-effect relationships

Conventional RNNs




The cell unrolled through time form the layers of the
neural network

Bidirectional RNNs




Bidirectional RNNs



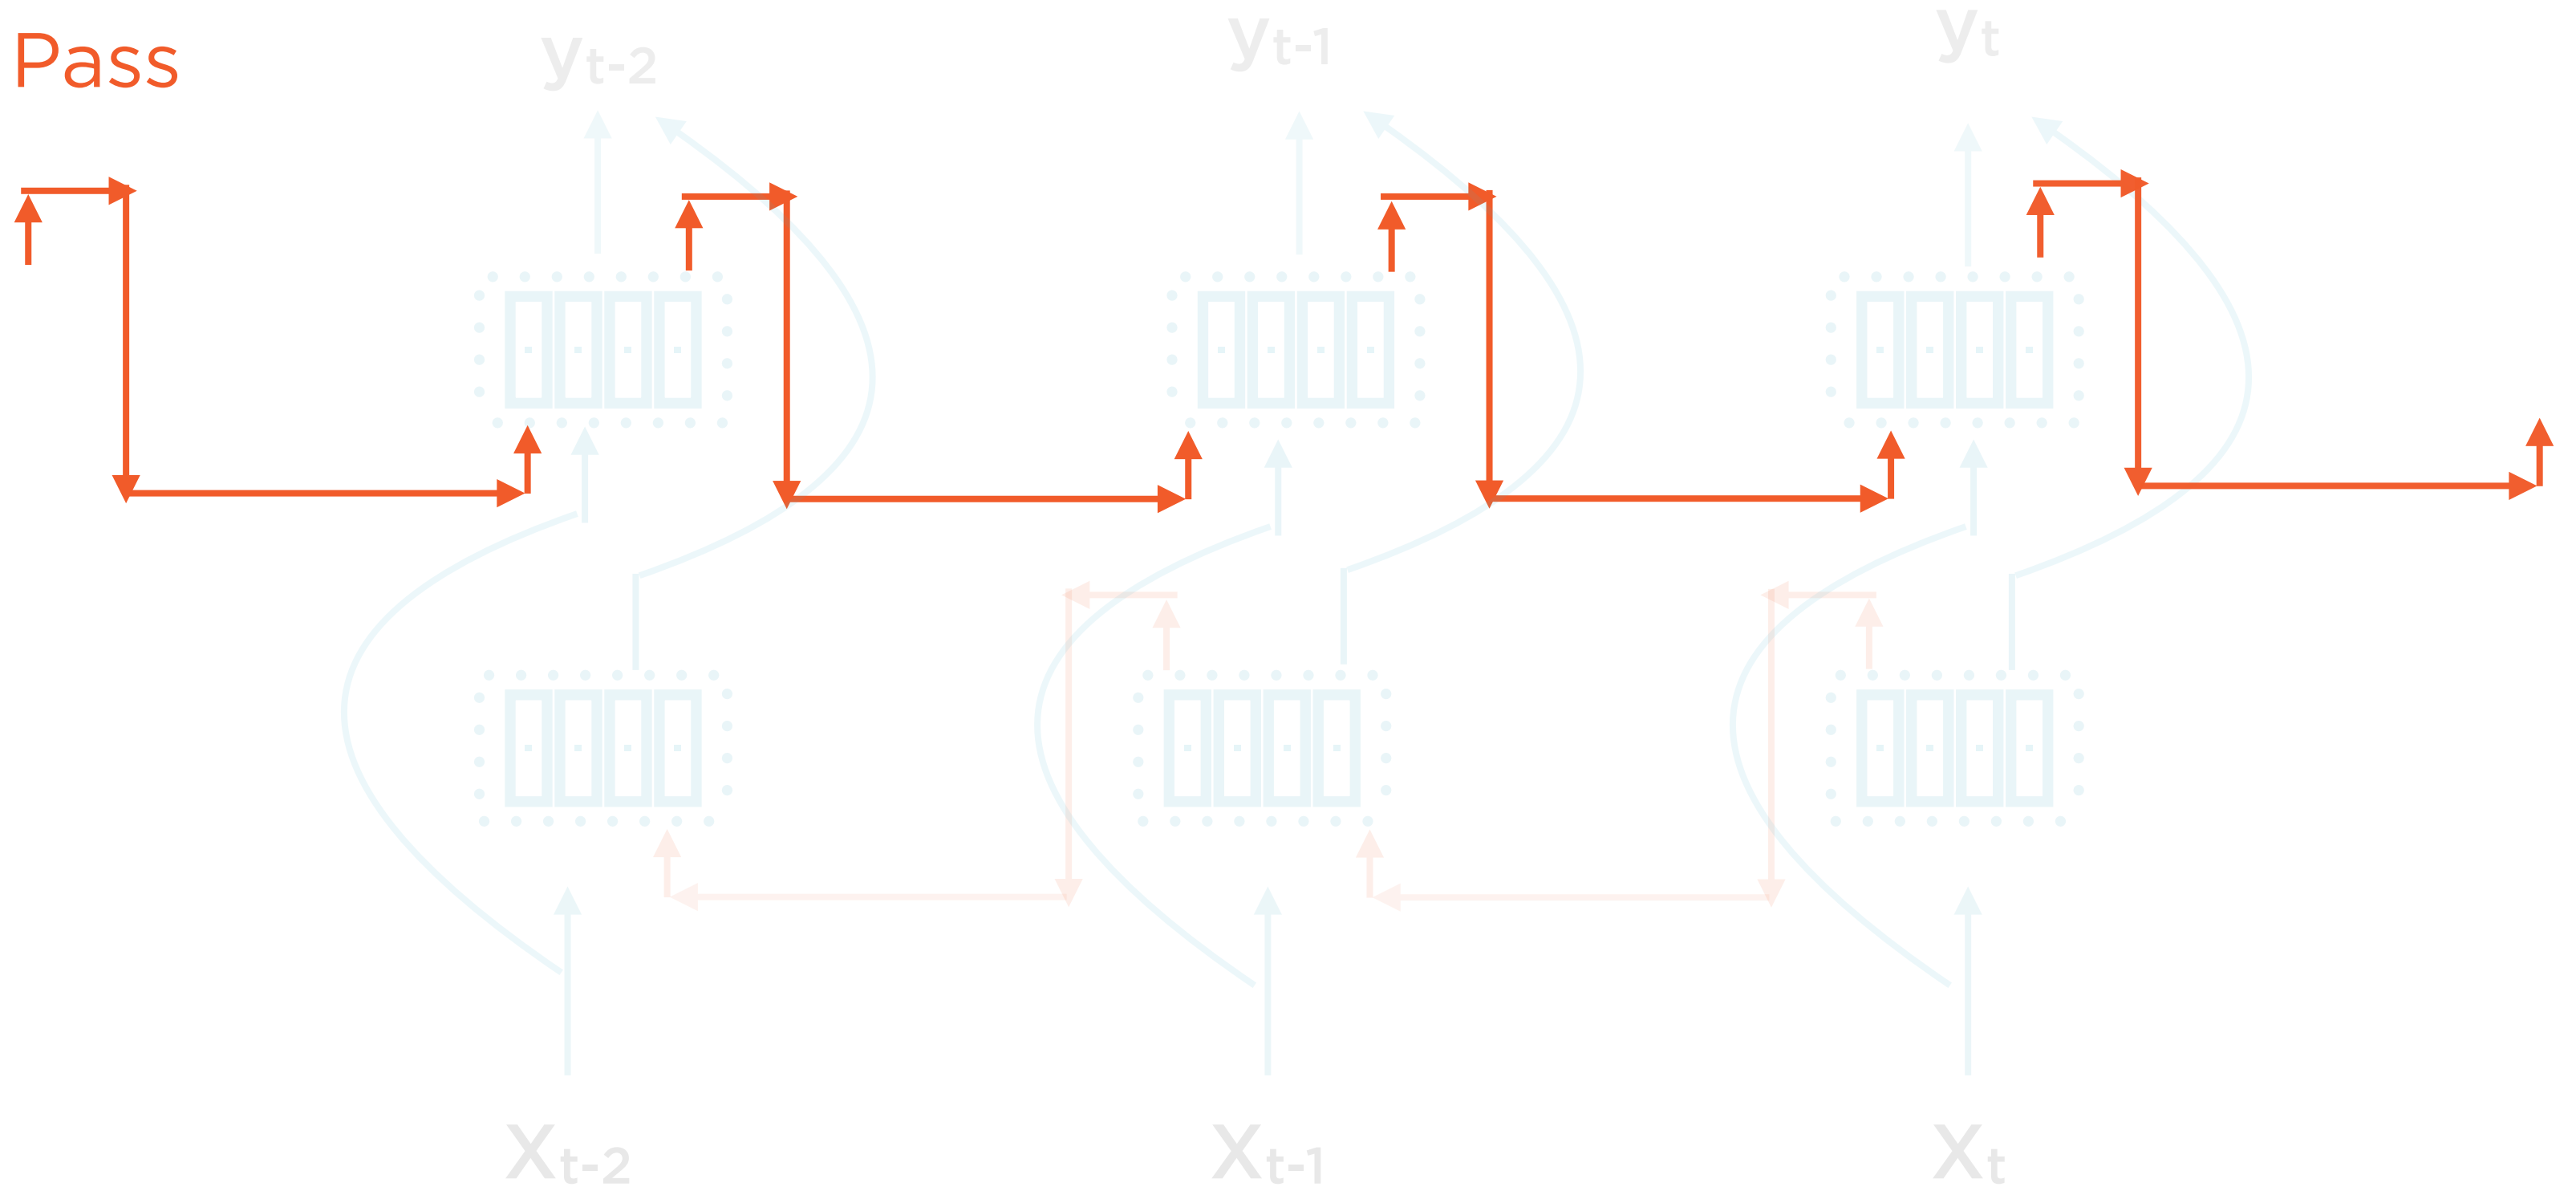
Bidirectional RNNs

Outputs




Bidirectional RNNs

Forward
Pass




Bidirectional RNNs

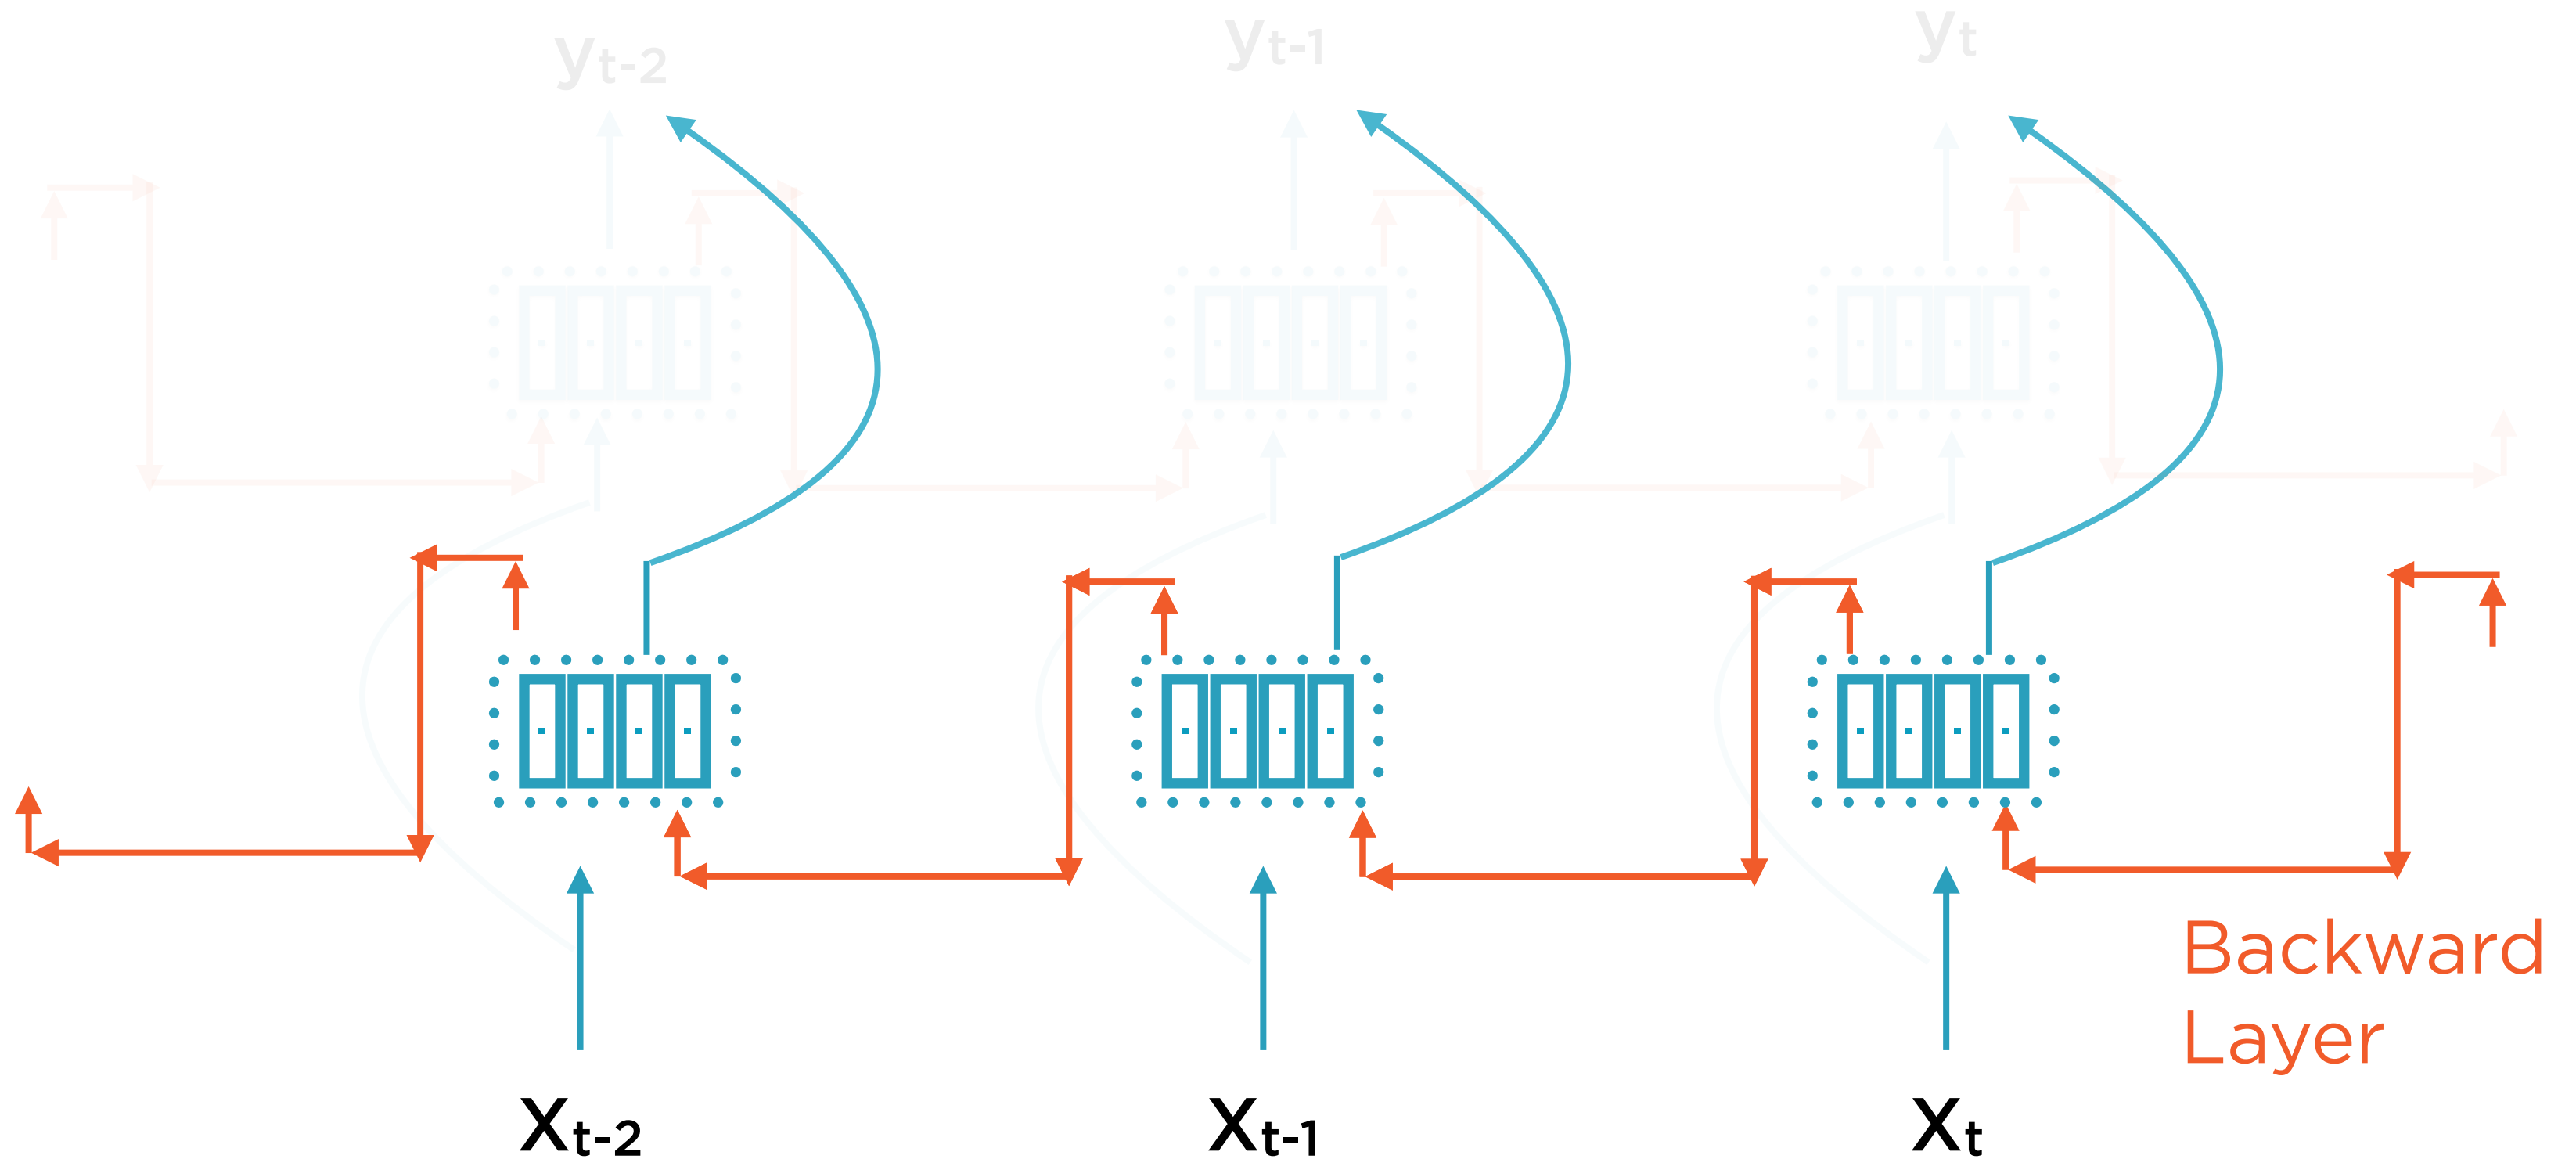
Forward
Layer



Bidirectional RNNs

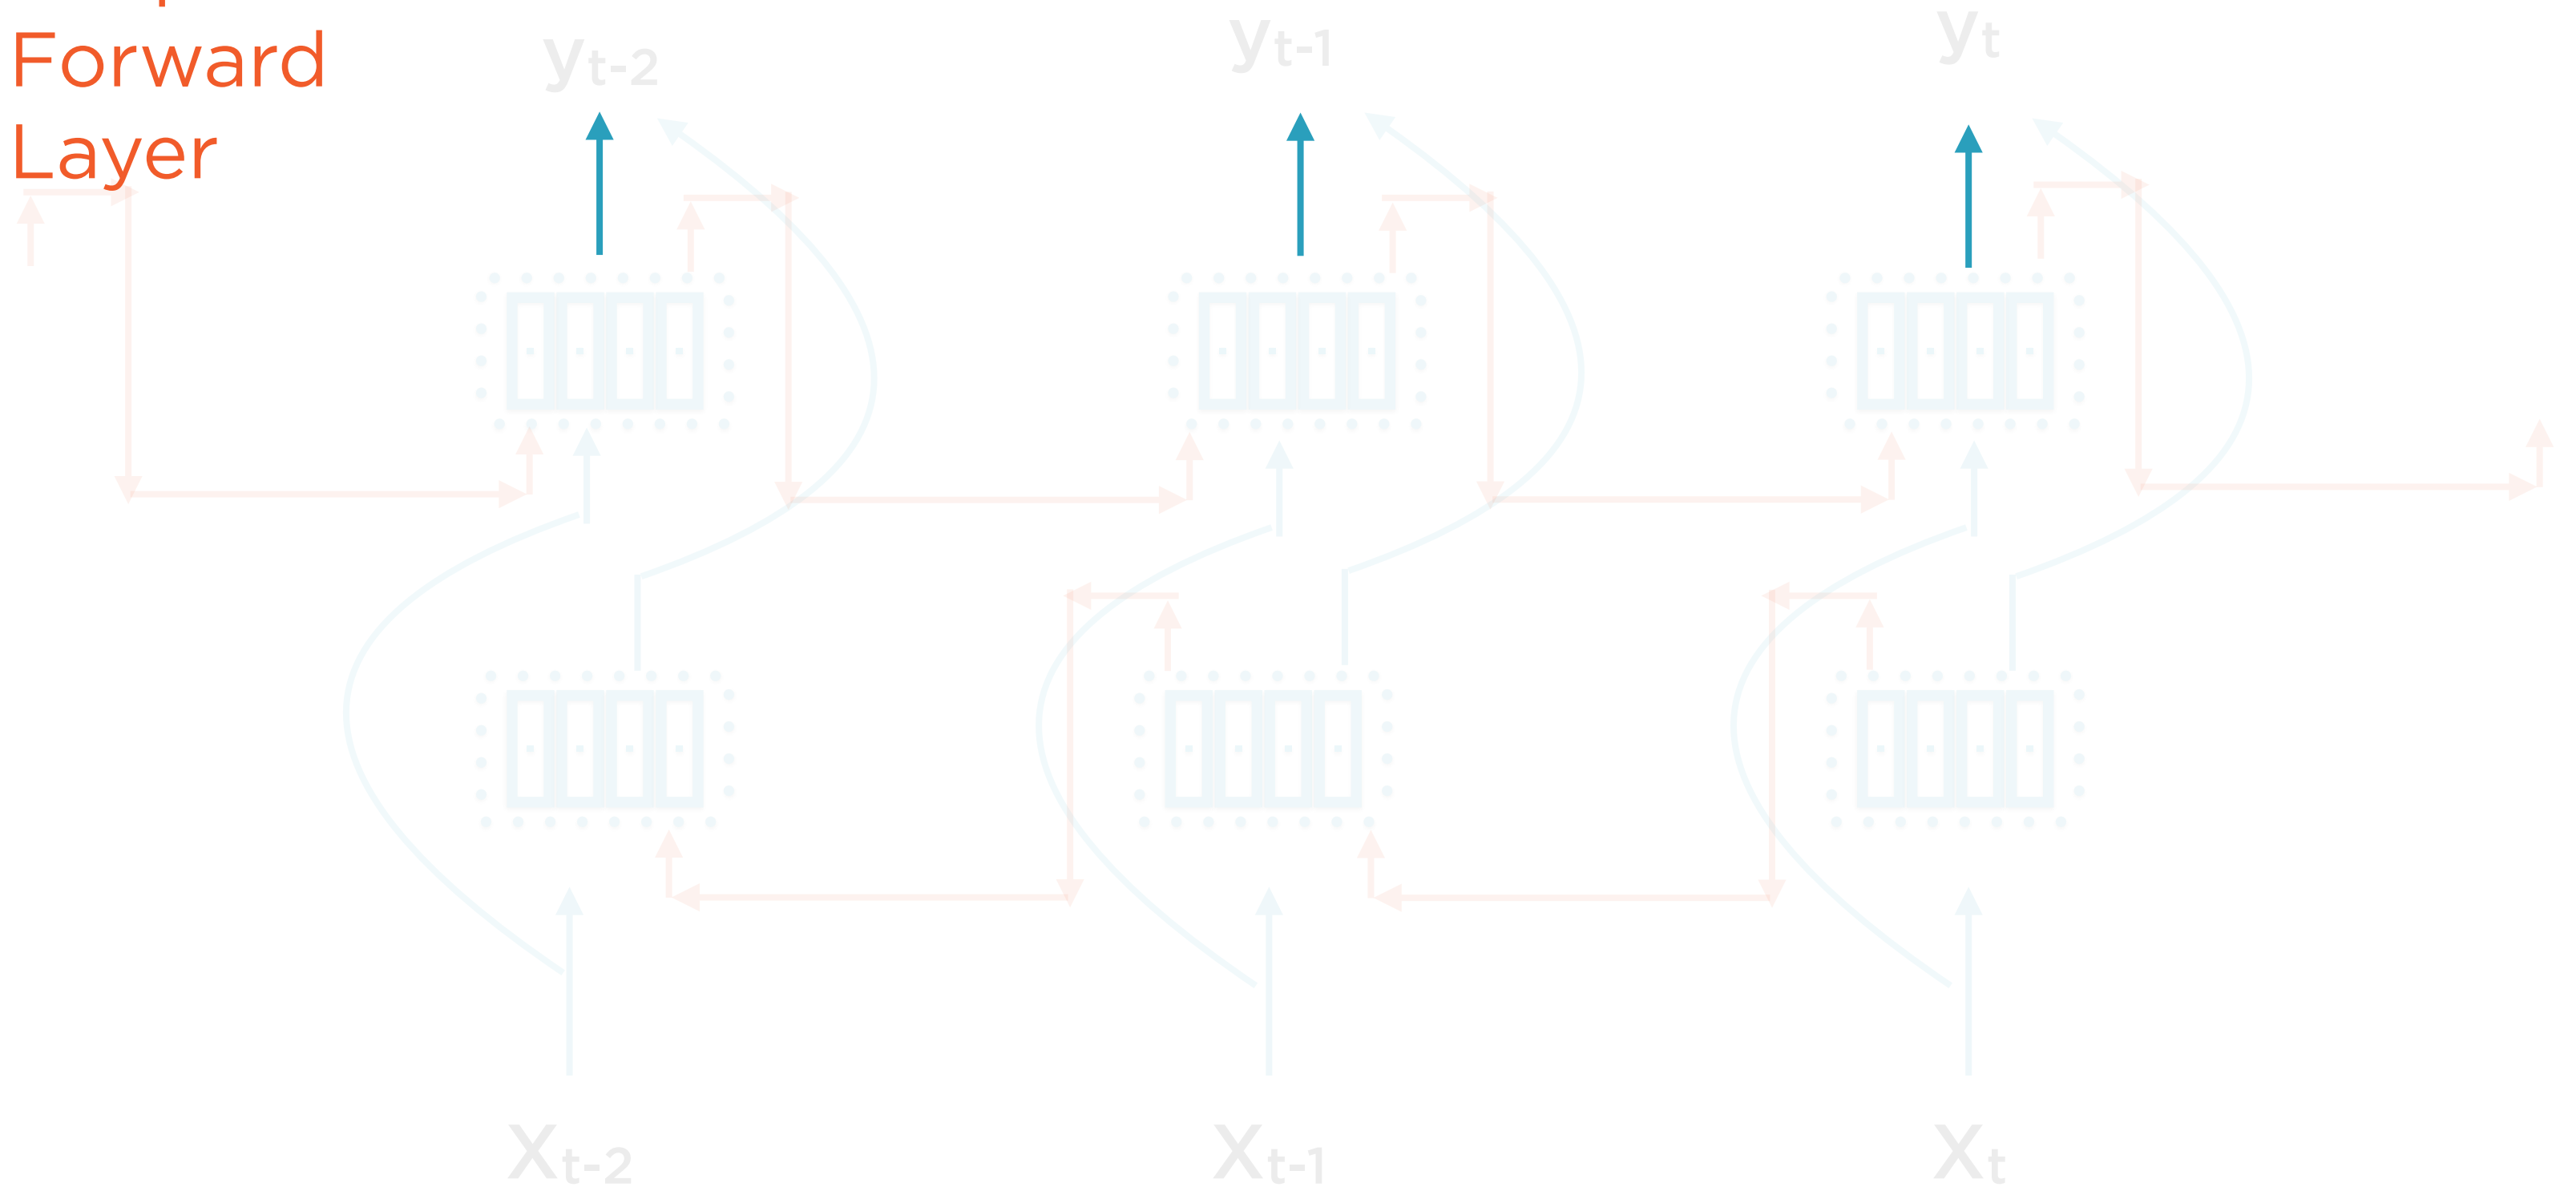


Bidirectional RNNs




Bidirectional RNNs

Outputs of
Forward
Layer




Bidirectional RNNs


Outputs of
Backward
Layer



Bidirectional RNNs




Bidirectional RNNs



Bidirectional RNNs

Outputs

y_t



Forward
Layer

Backward
Layer

Inputs

x_t

Forward and backward outputs can
be **combined** together in some
manner

Demo

Sentiment analysis using word embeddings

Summary

Representing textual data as features for machine learning

Count, frequency, and probability-based embeddings

GloVe and word2vec for pre-trained word embeddings

Performing sentiment analysis using pre-trained word embeddings

Bi-directional RNNs for sentiment analysis